

2017 PSI4 ORDNANCE SURVEY

LORI A. BURNS

PSI WWDC

3 NOV 2017

PSI4

INPUT MODES

Psithon

```
molecule dimer {
    0
    H 1 1.0
    H 1 1.0 2 90.0
    --
    H 1 R 2 135.0 3 0.0
    O 4 1.0 1 135.0 3 0.0
    H 5 1.0 4 90.0 90.0
    --
    R = 3.0
}

set basis aug-cc-pvdz

energy("sapt2+dmp2")
```

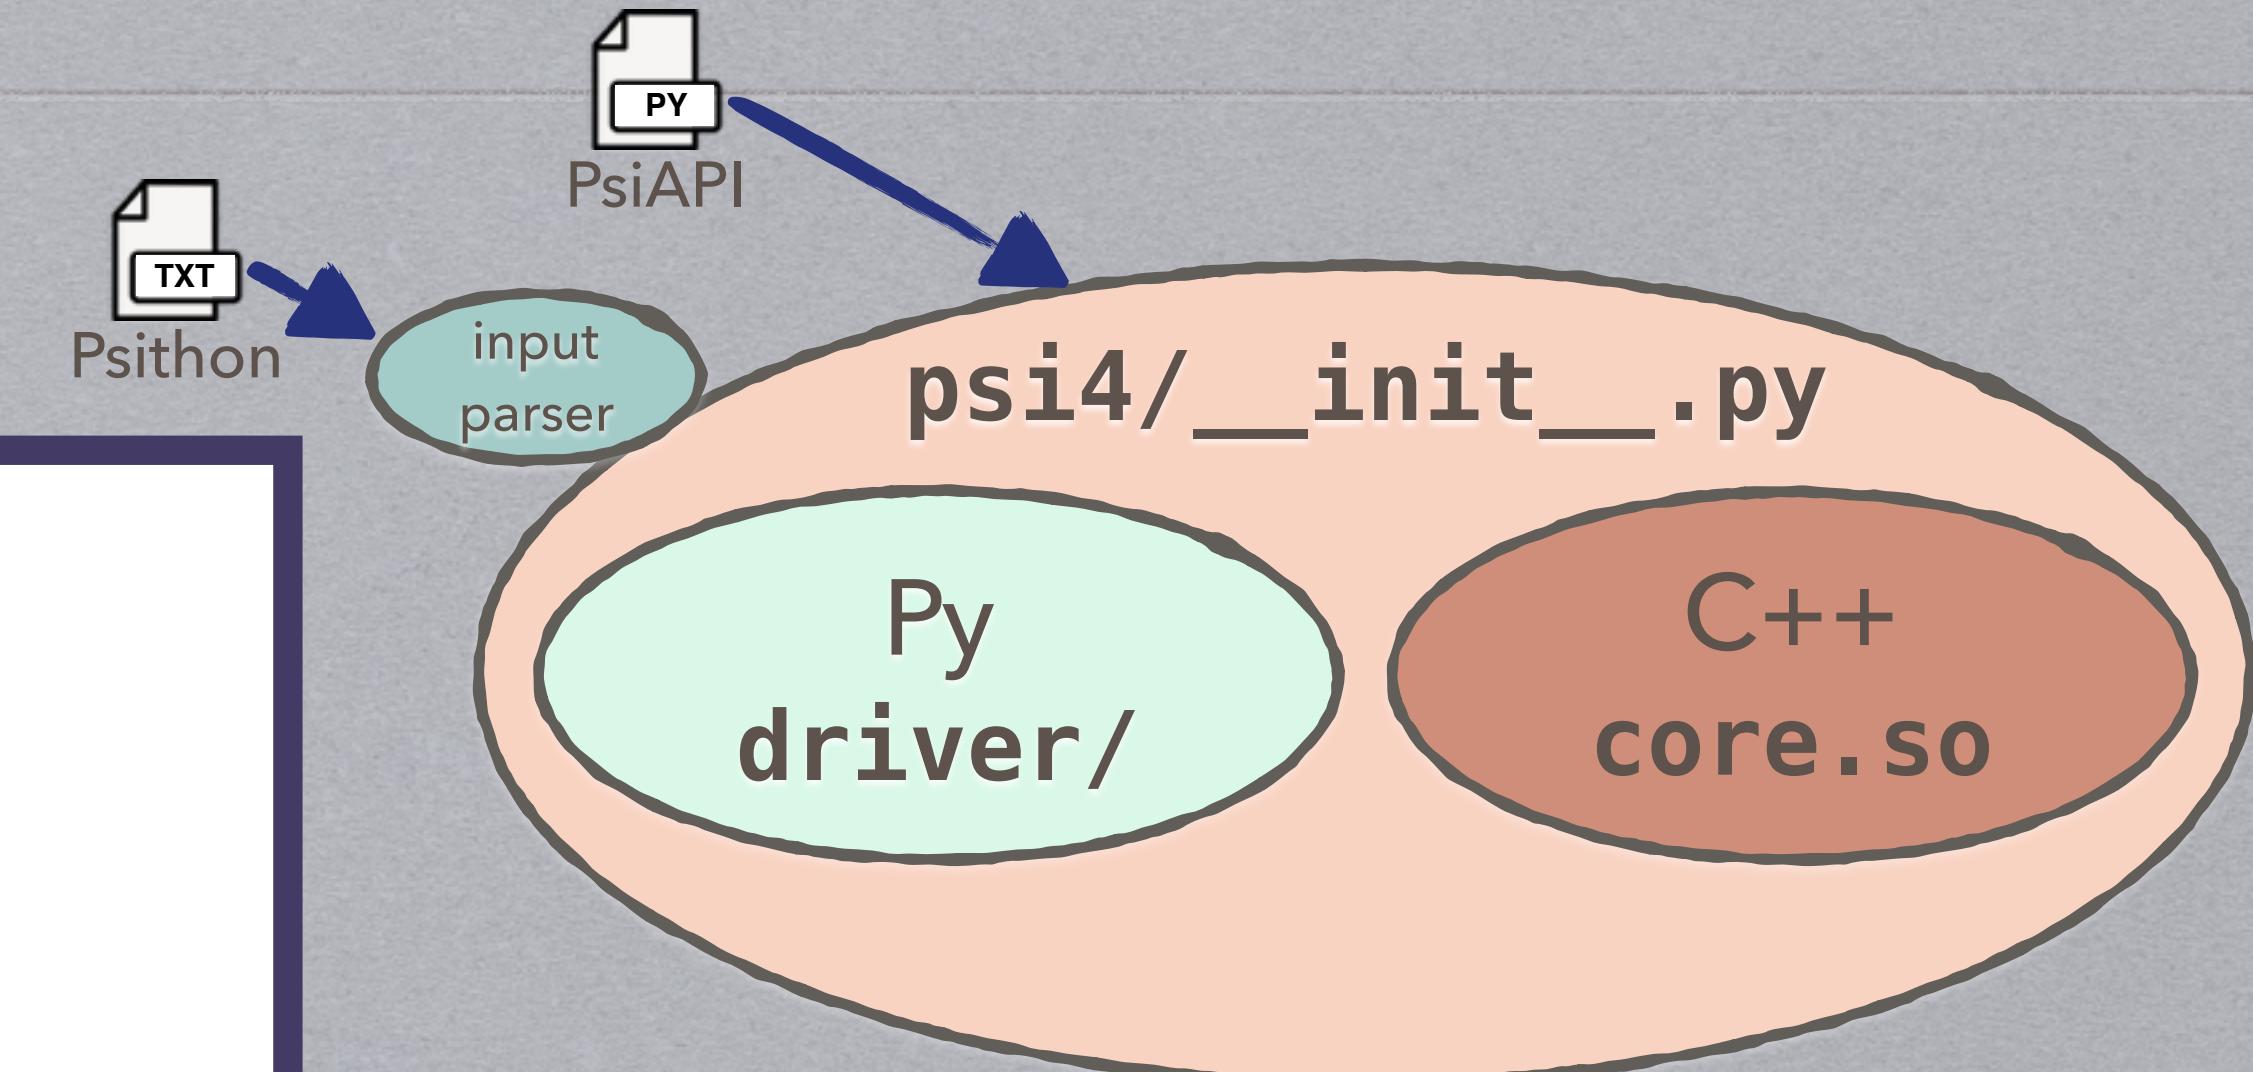
PsiAPI

```
import psi4

psi4.geometry("""
0
H 1 1.0
H 1 1.0 2 90.0
--
H 1 R 2 135.0 3 0.0
O 4 1.0 1 135.0 3 0.0
H 5 1.0 4 90.0 90.0
R = 3.0
""")

psi4.set_options({
    "basis": "aug-cc-pvdz"})

psi4.energy("sapt2+dmp2")
```

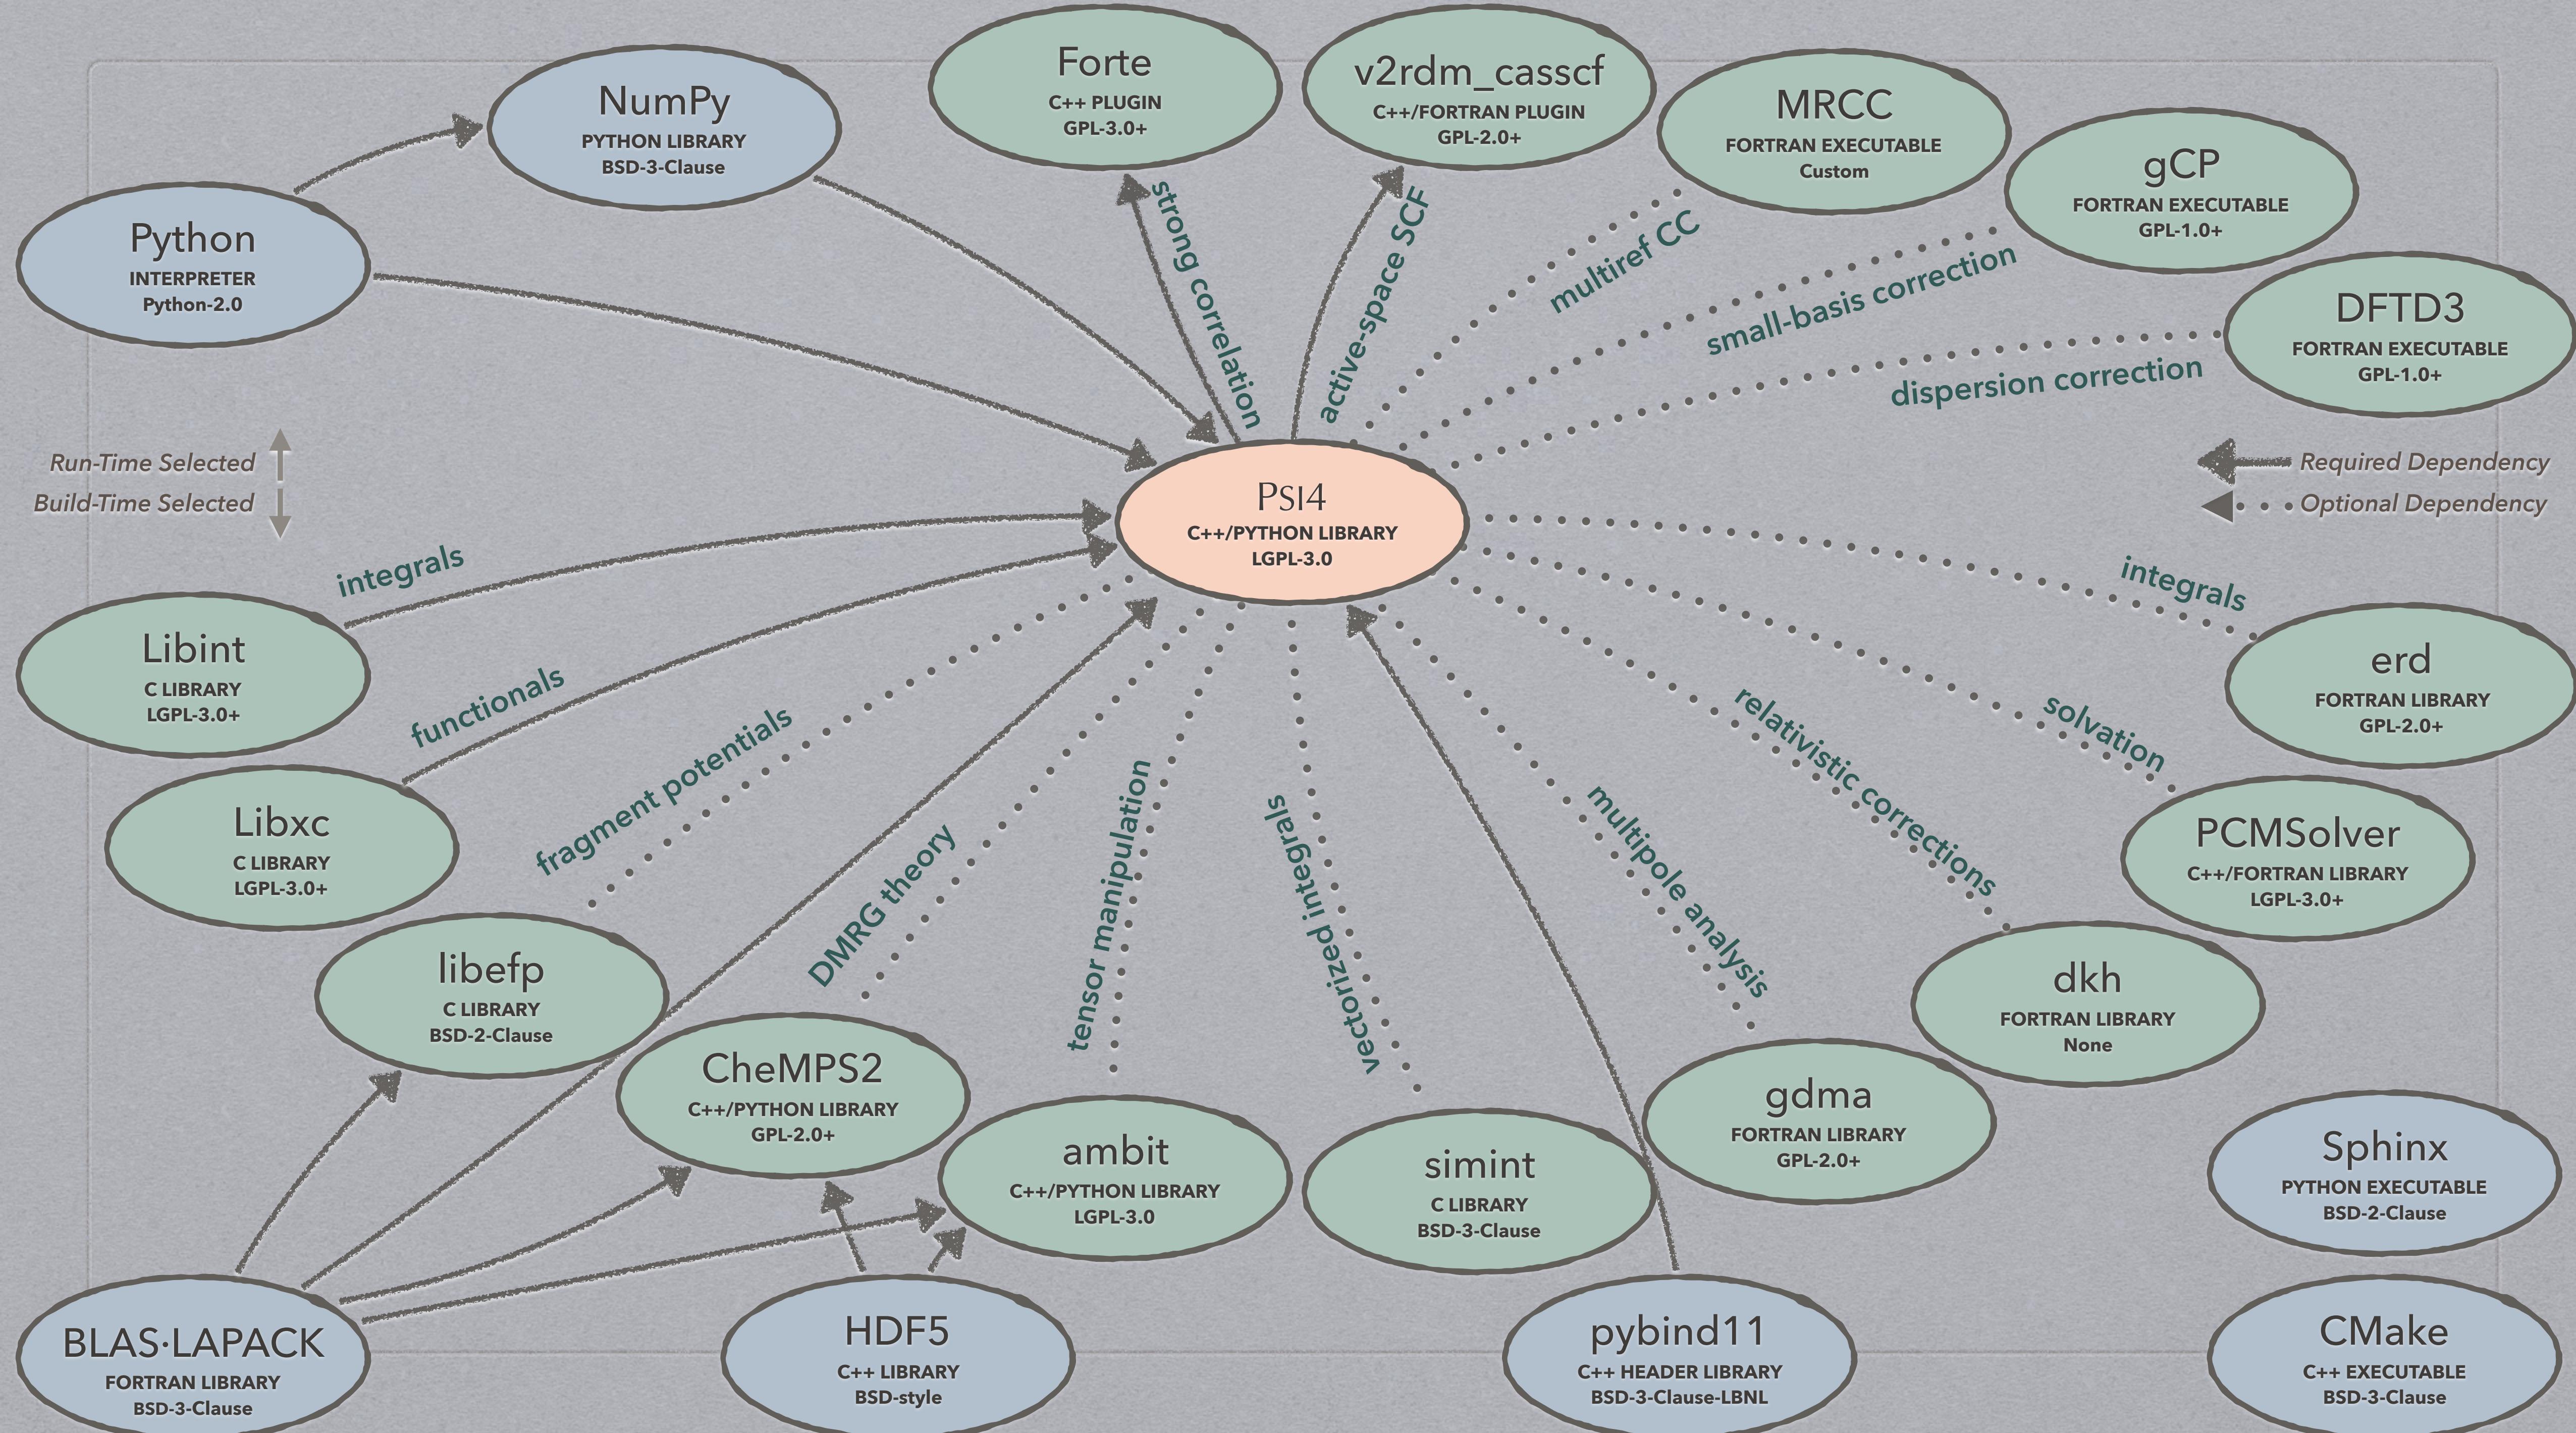


PYTHON MODULE W/C++ EXTENSION

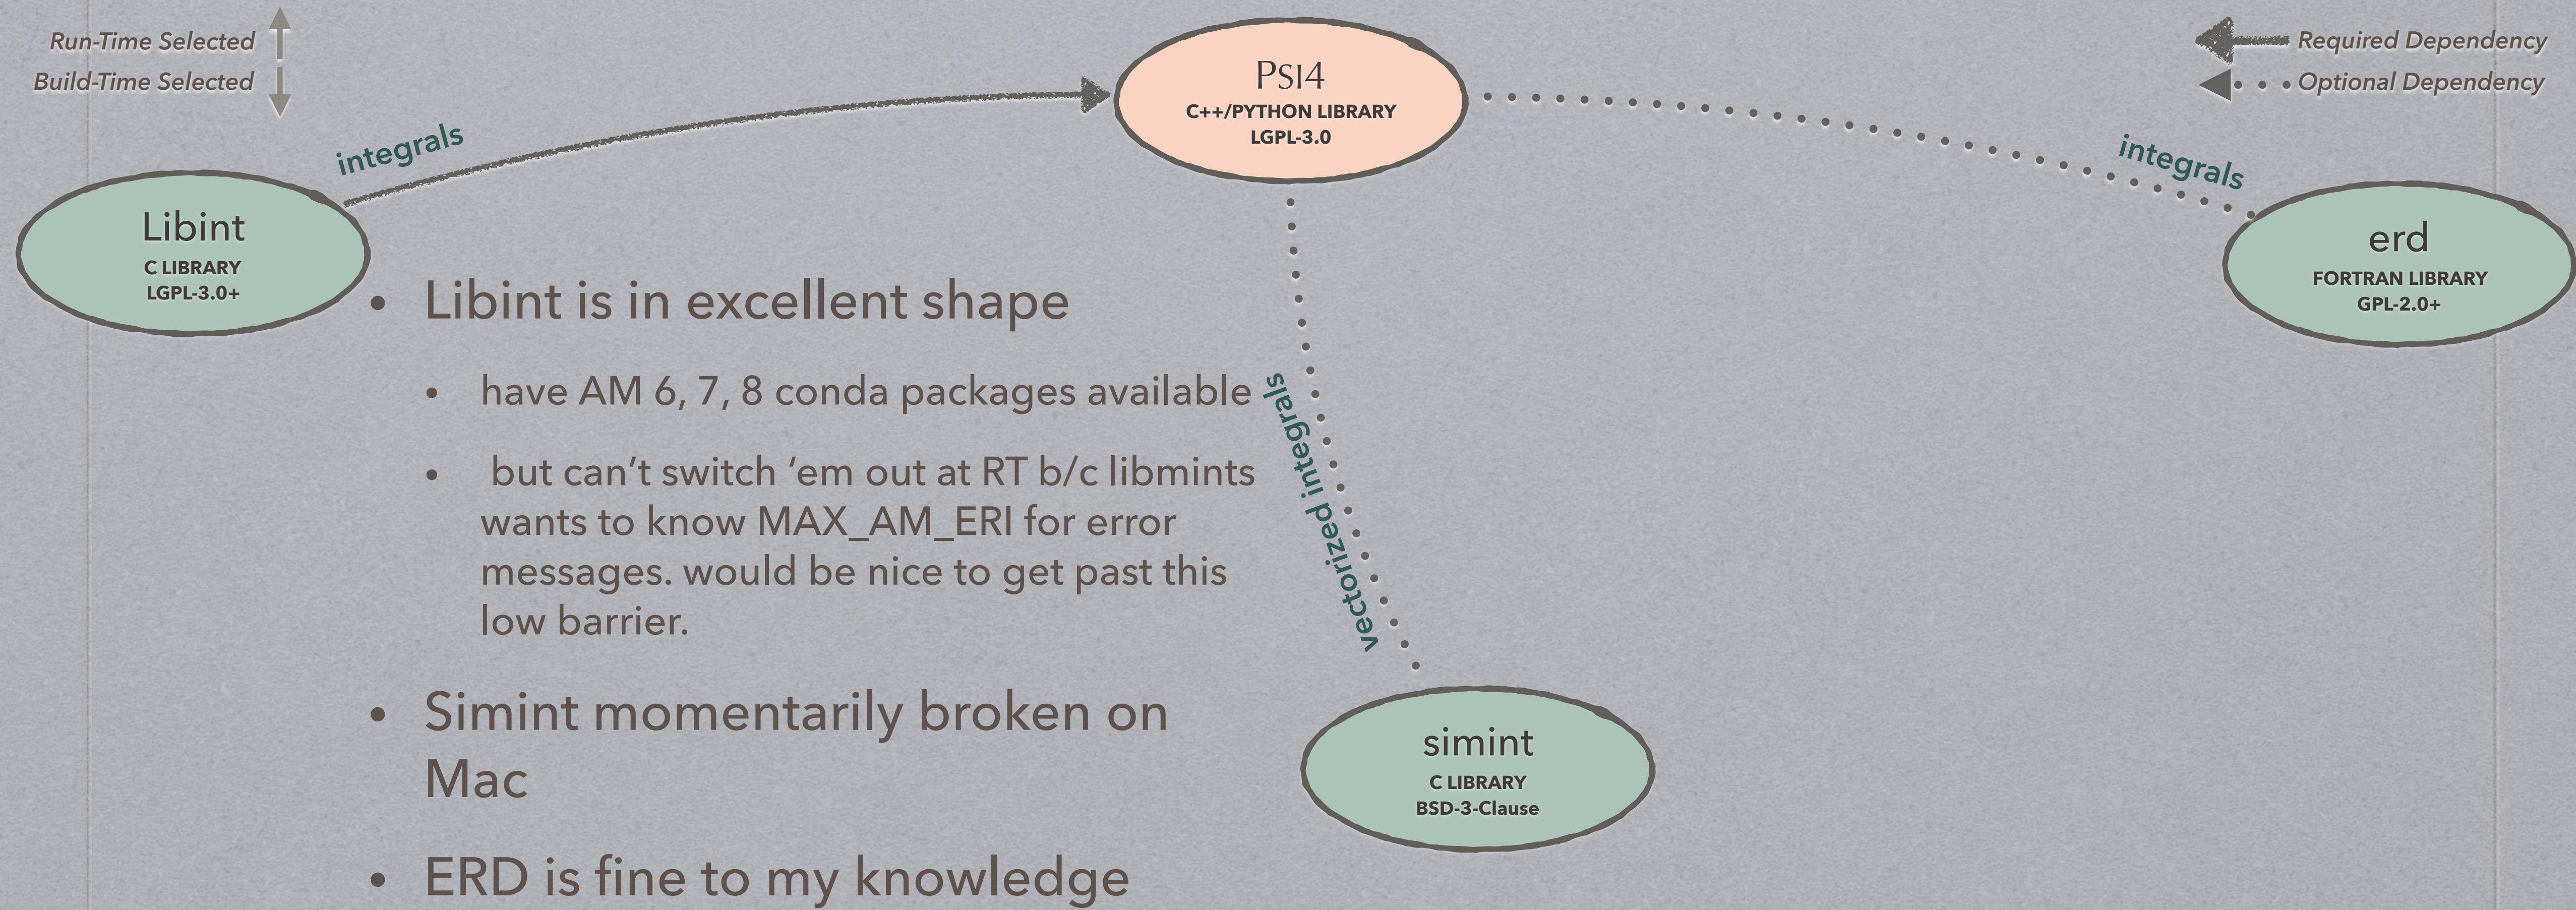
PSI4 c. 2017, 1.1

- **Psithon (normal)**
 - been around forever
 - would deprecate the term as domain-specific languages
not come il faut but it's useful to name contrast to psiapi
 - sloppier namespace
 - set PATH, run with **psi4 input.in**
- **PsiAPI (new)**
 - new mode, just as powerful
 - don't care what we call it, but a name is handy
 - set PYTHONPATH, run with **python input.py**

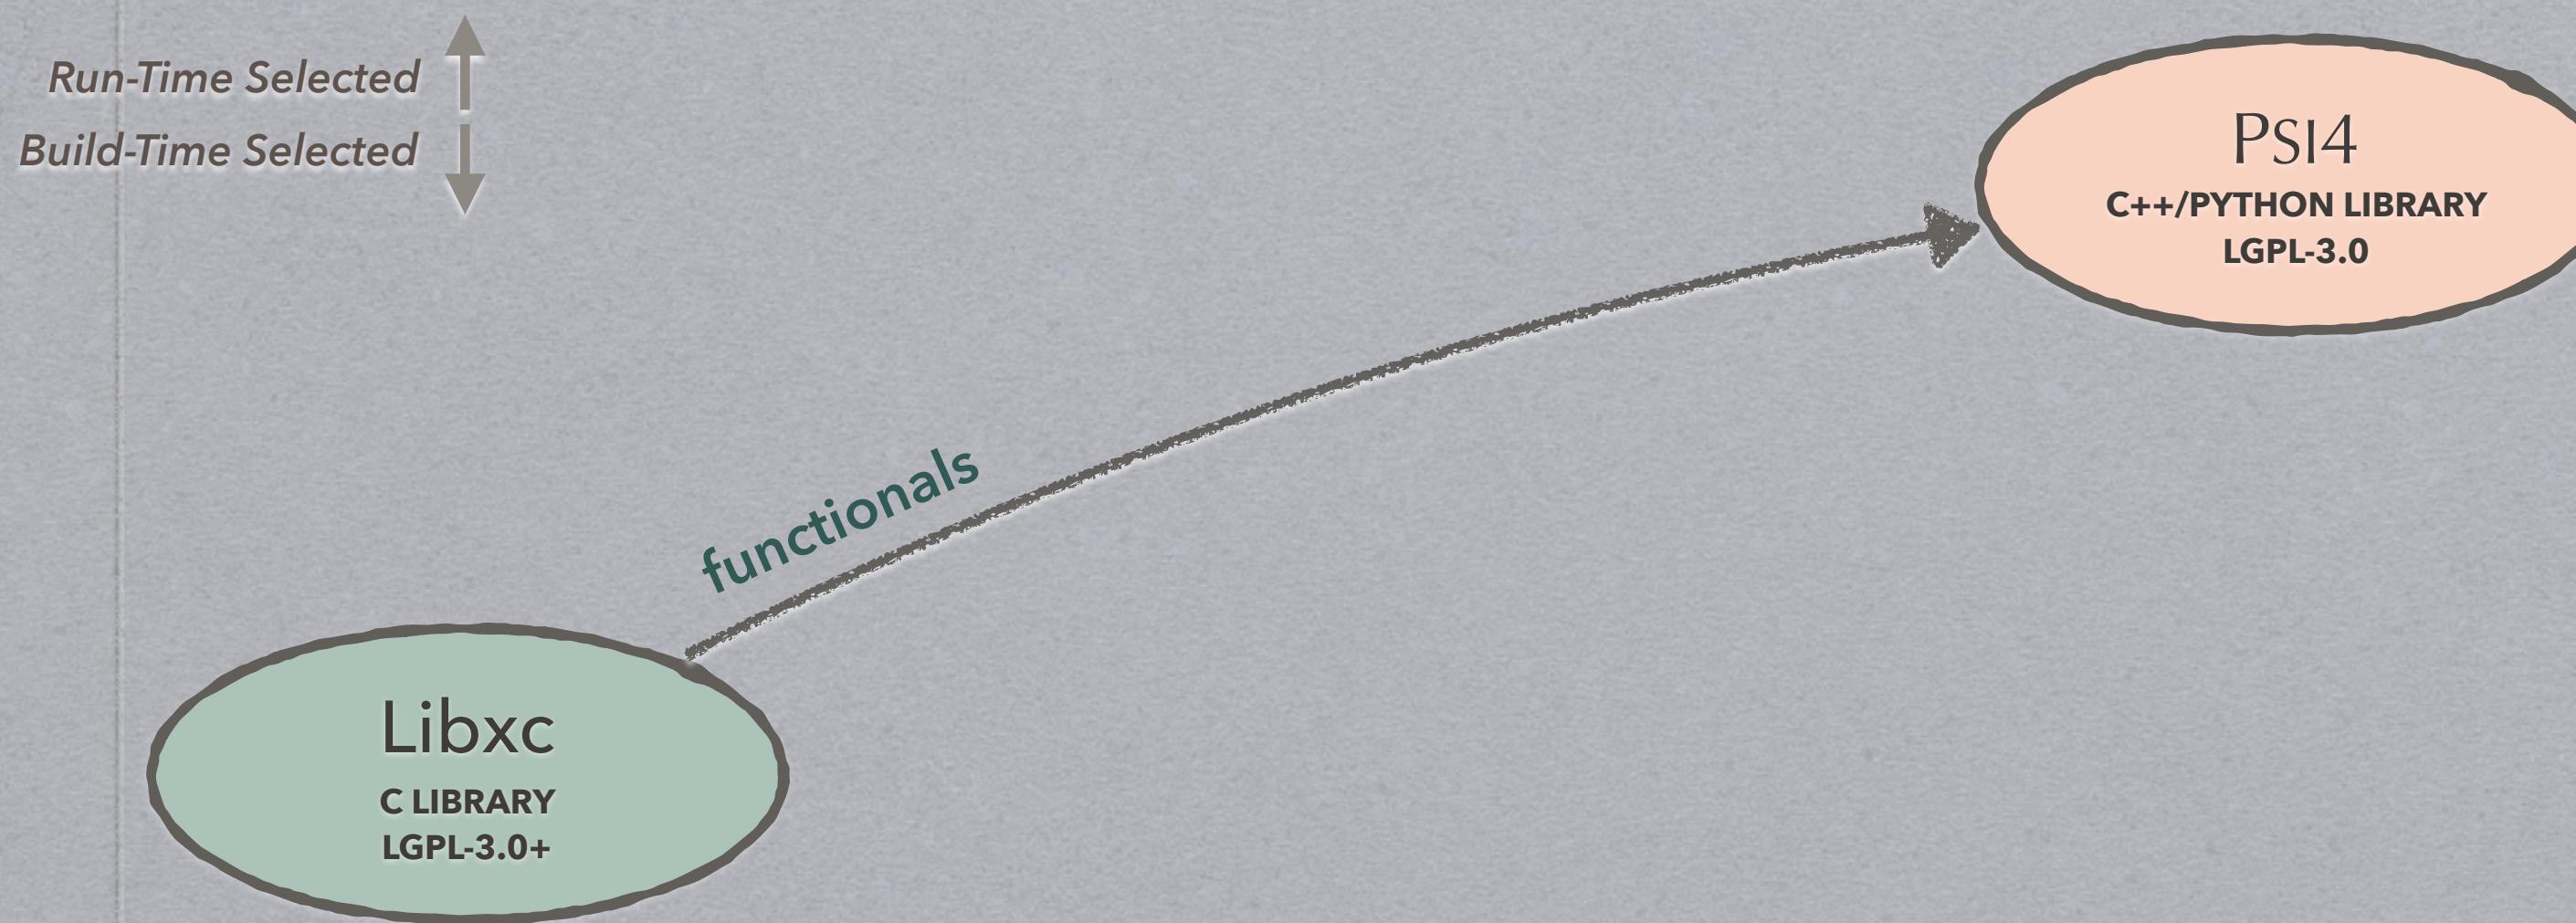
```
>>> psi4 --psiapi-path
export PATH=/Users/username/miniconda/envs/p4dev/bin:$PATH
export PYTHONPATH=/Users/username/gits/objdir/stage/usr/local/psi4/lib:$PYTHONPATH
>>> python -c "import psi4; print(psi4.__version__)"
```



INTEGRALS

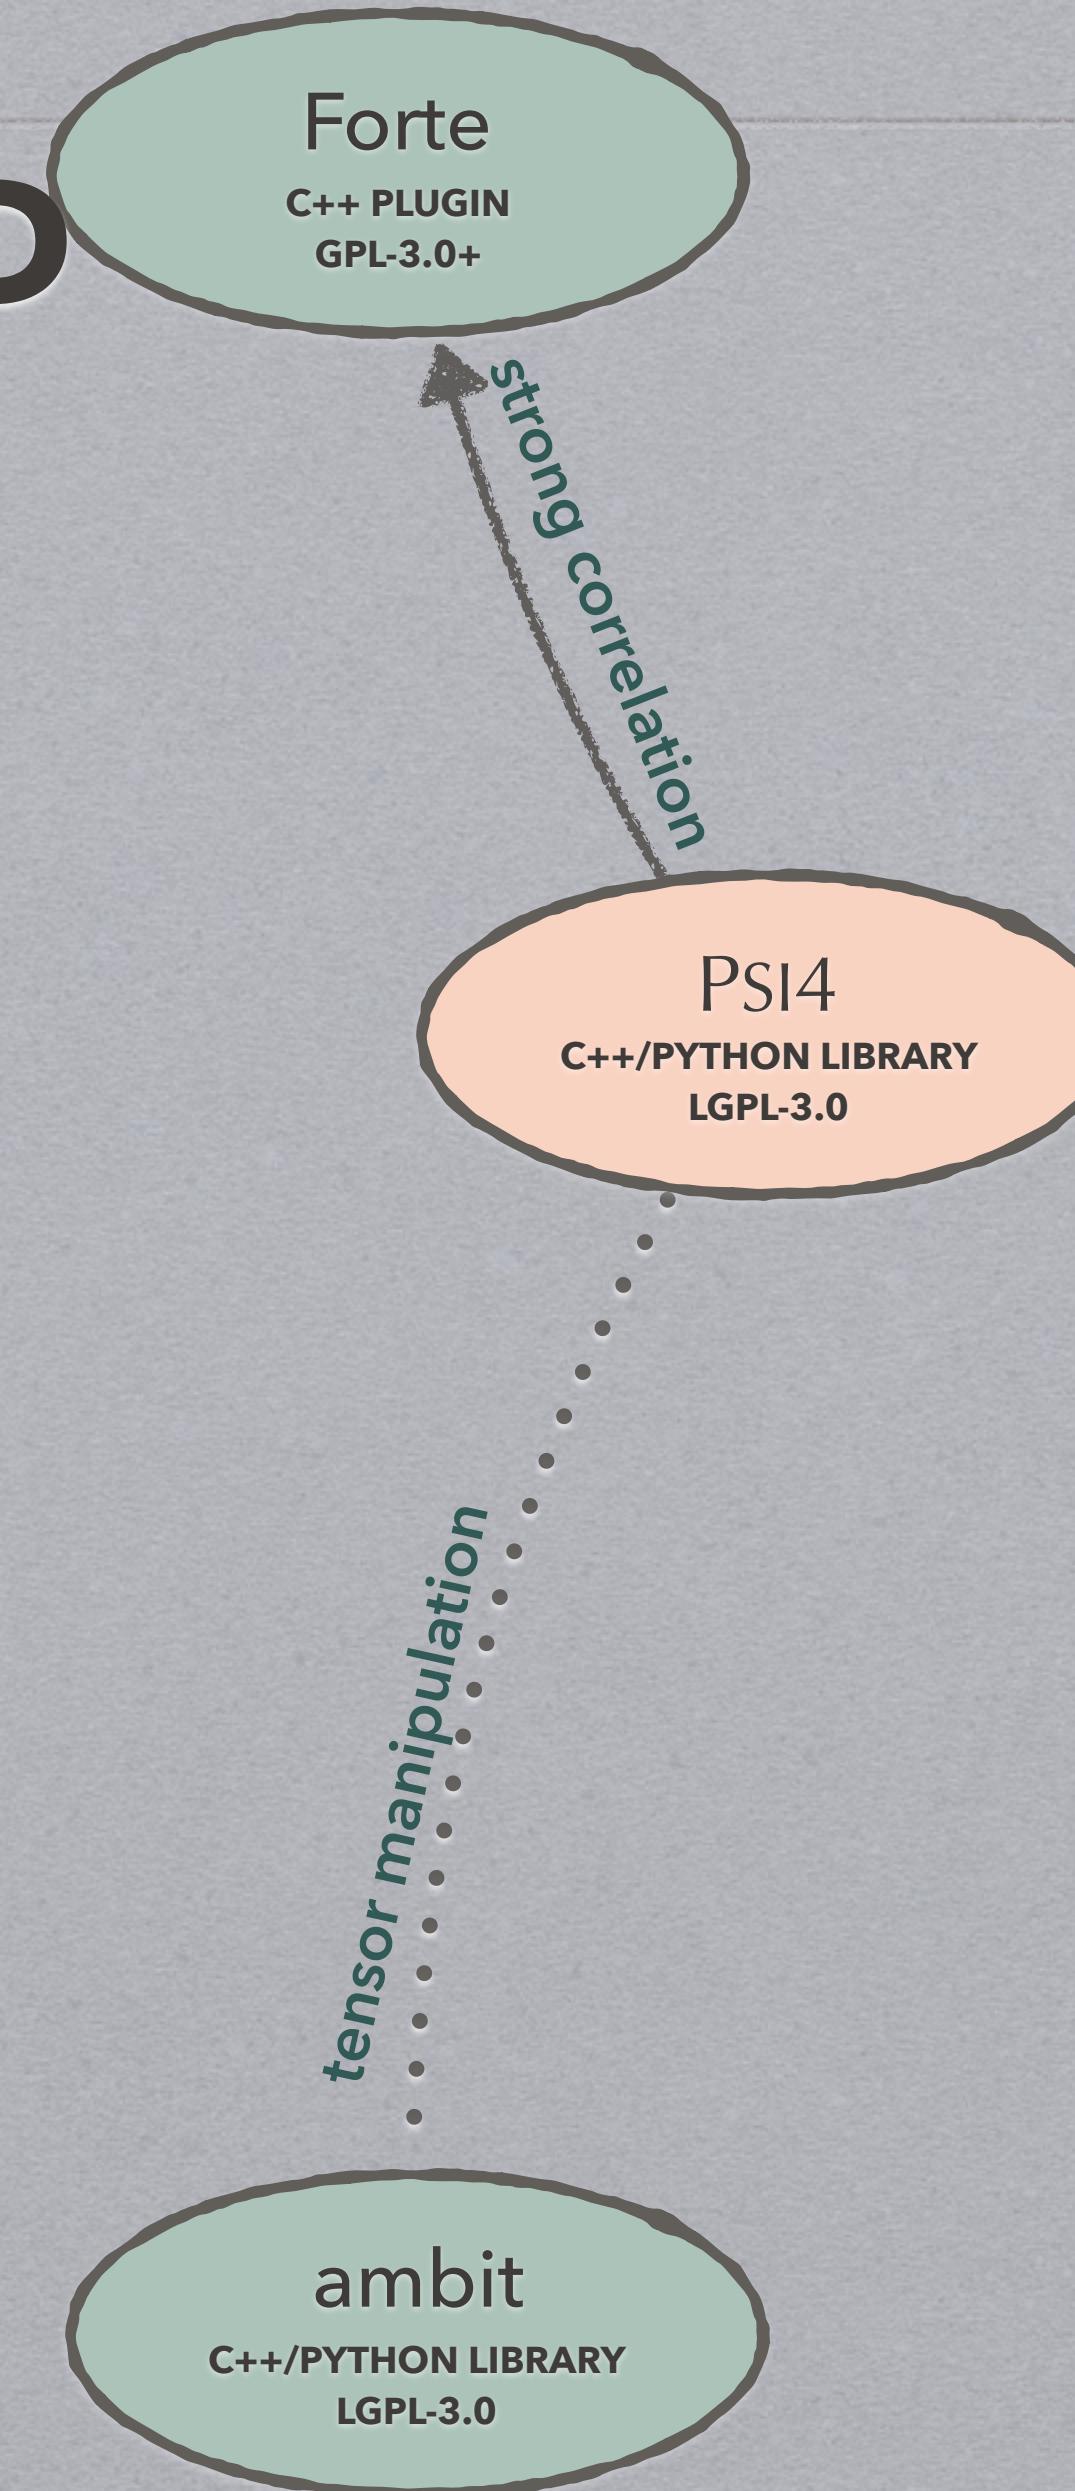


FUNCTIONALS



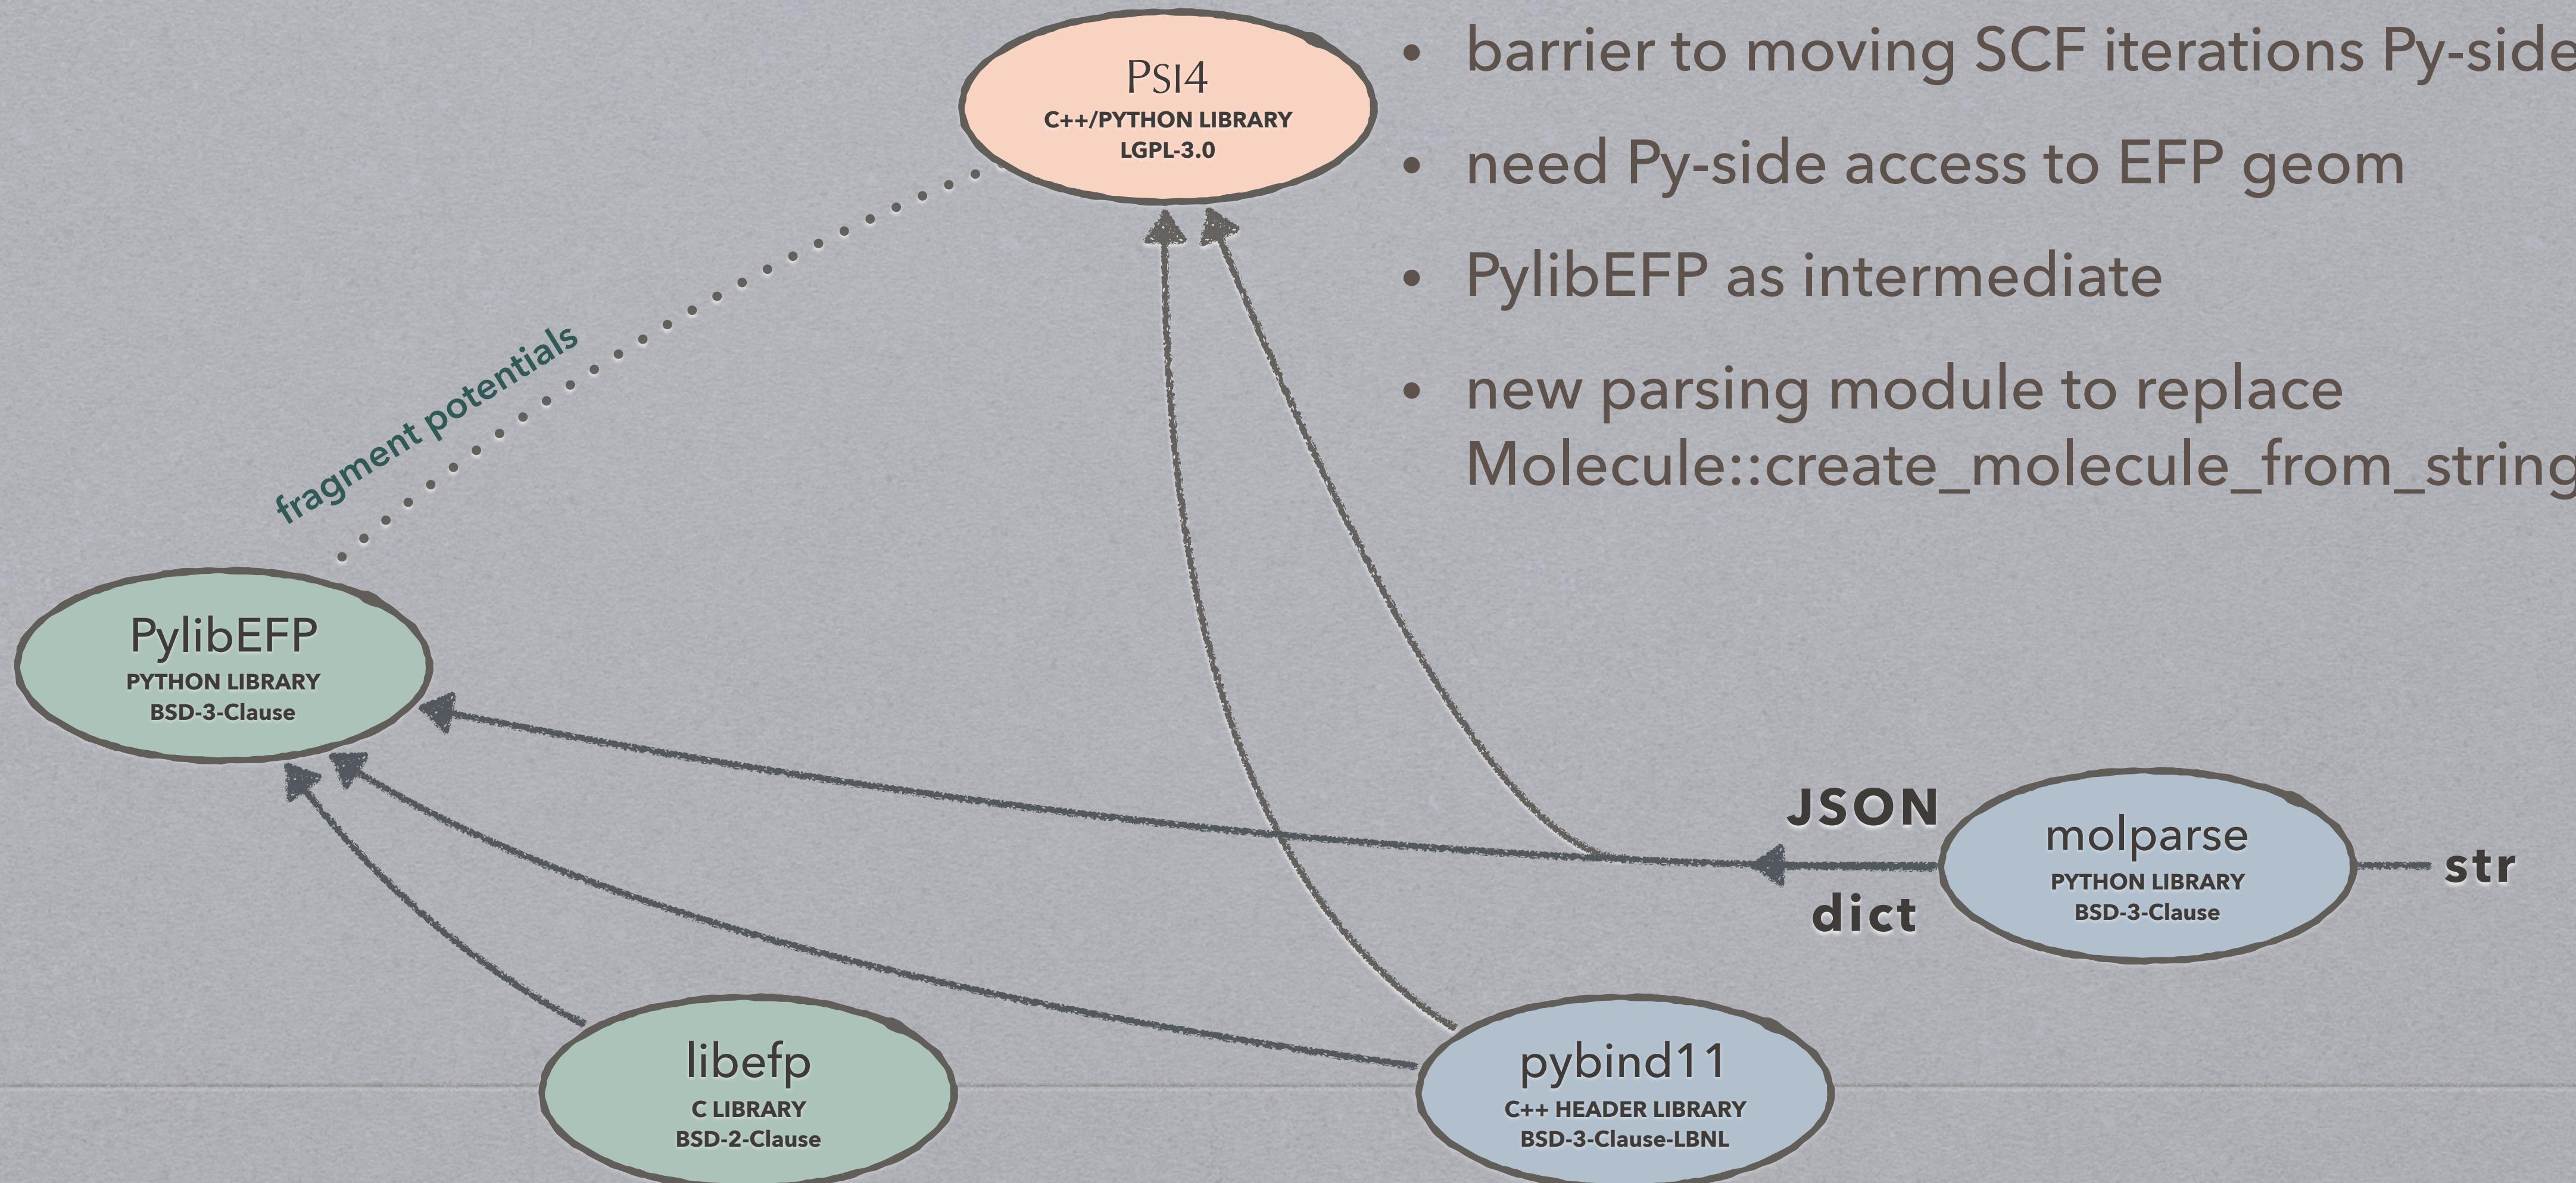
- Libxc was doing fine; we were prepared to switch to upstream when they indicated stability (e.g., RC1)
- Unfortunately, they simply declared 4.0
- New interface separates internal vs. external functional parameter setters in such a way as to greatly restrict RT-manipulation of fctls. Probably can't reproduce current Psi capabilities in 4.0.
- So currently still stuck at 3.0, pulling from GH:psi4/libxc, rather than GL:libxc/libxc.
 - joining PySCF & Horton stuck at Libxc 2.0.
 - we should communicate release dates and provide migration paths when we break API, lest others get irked at us

TENSOR CROWD



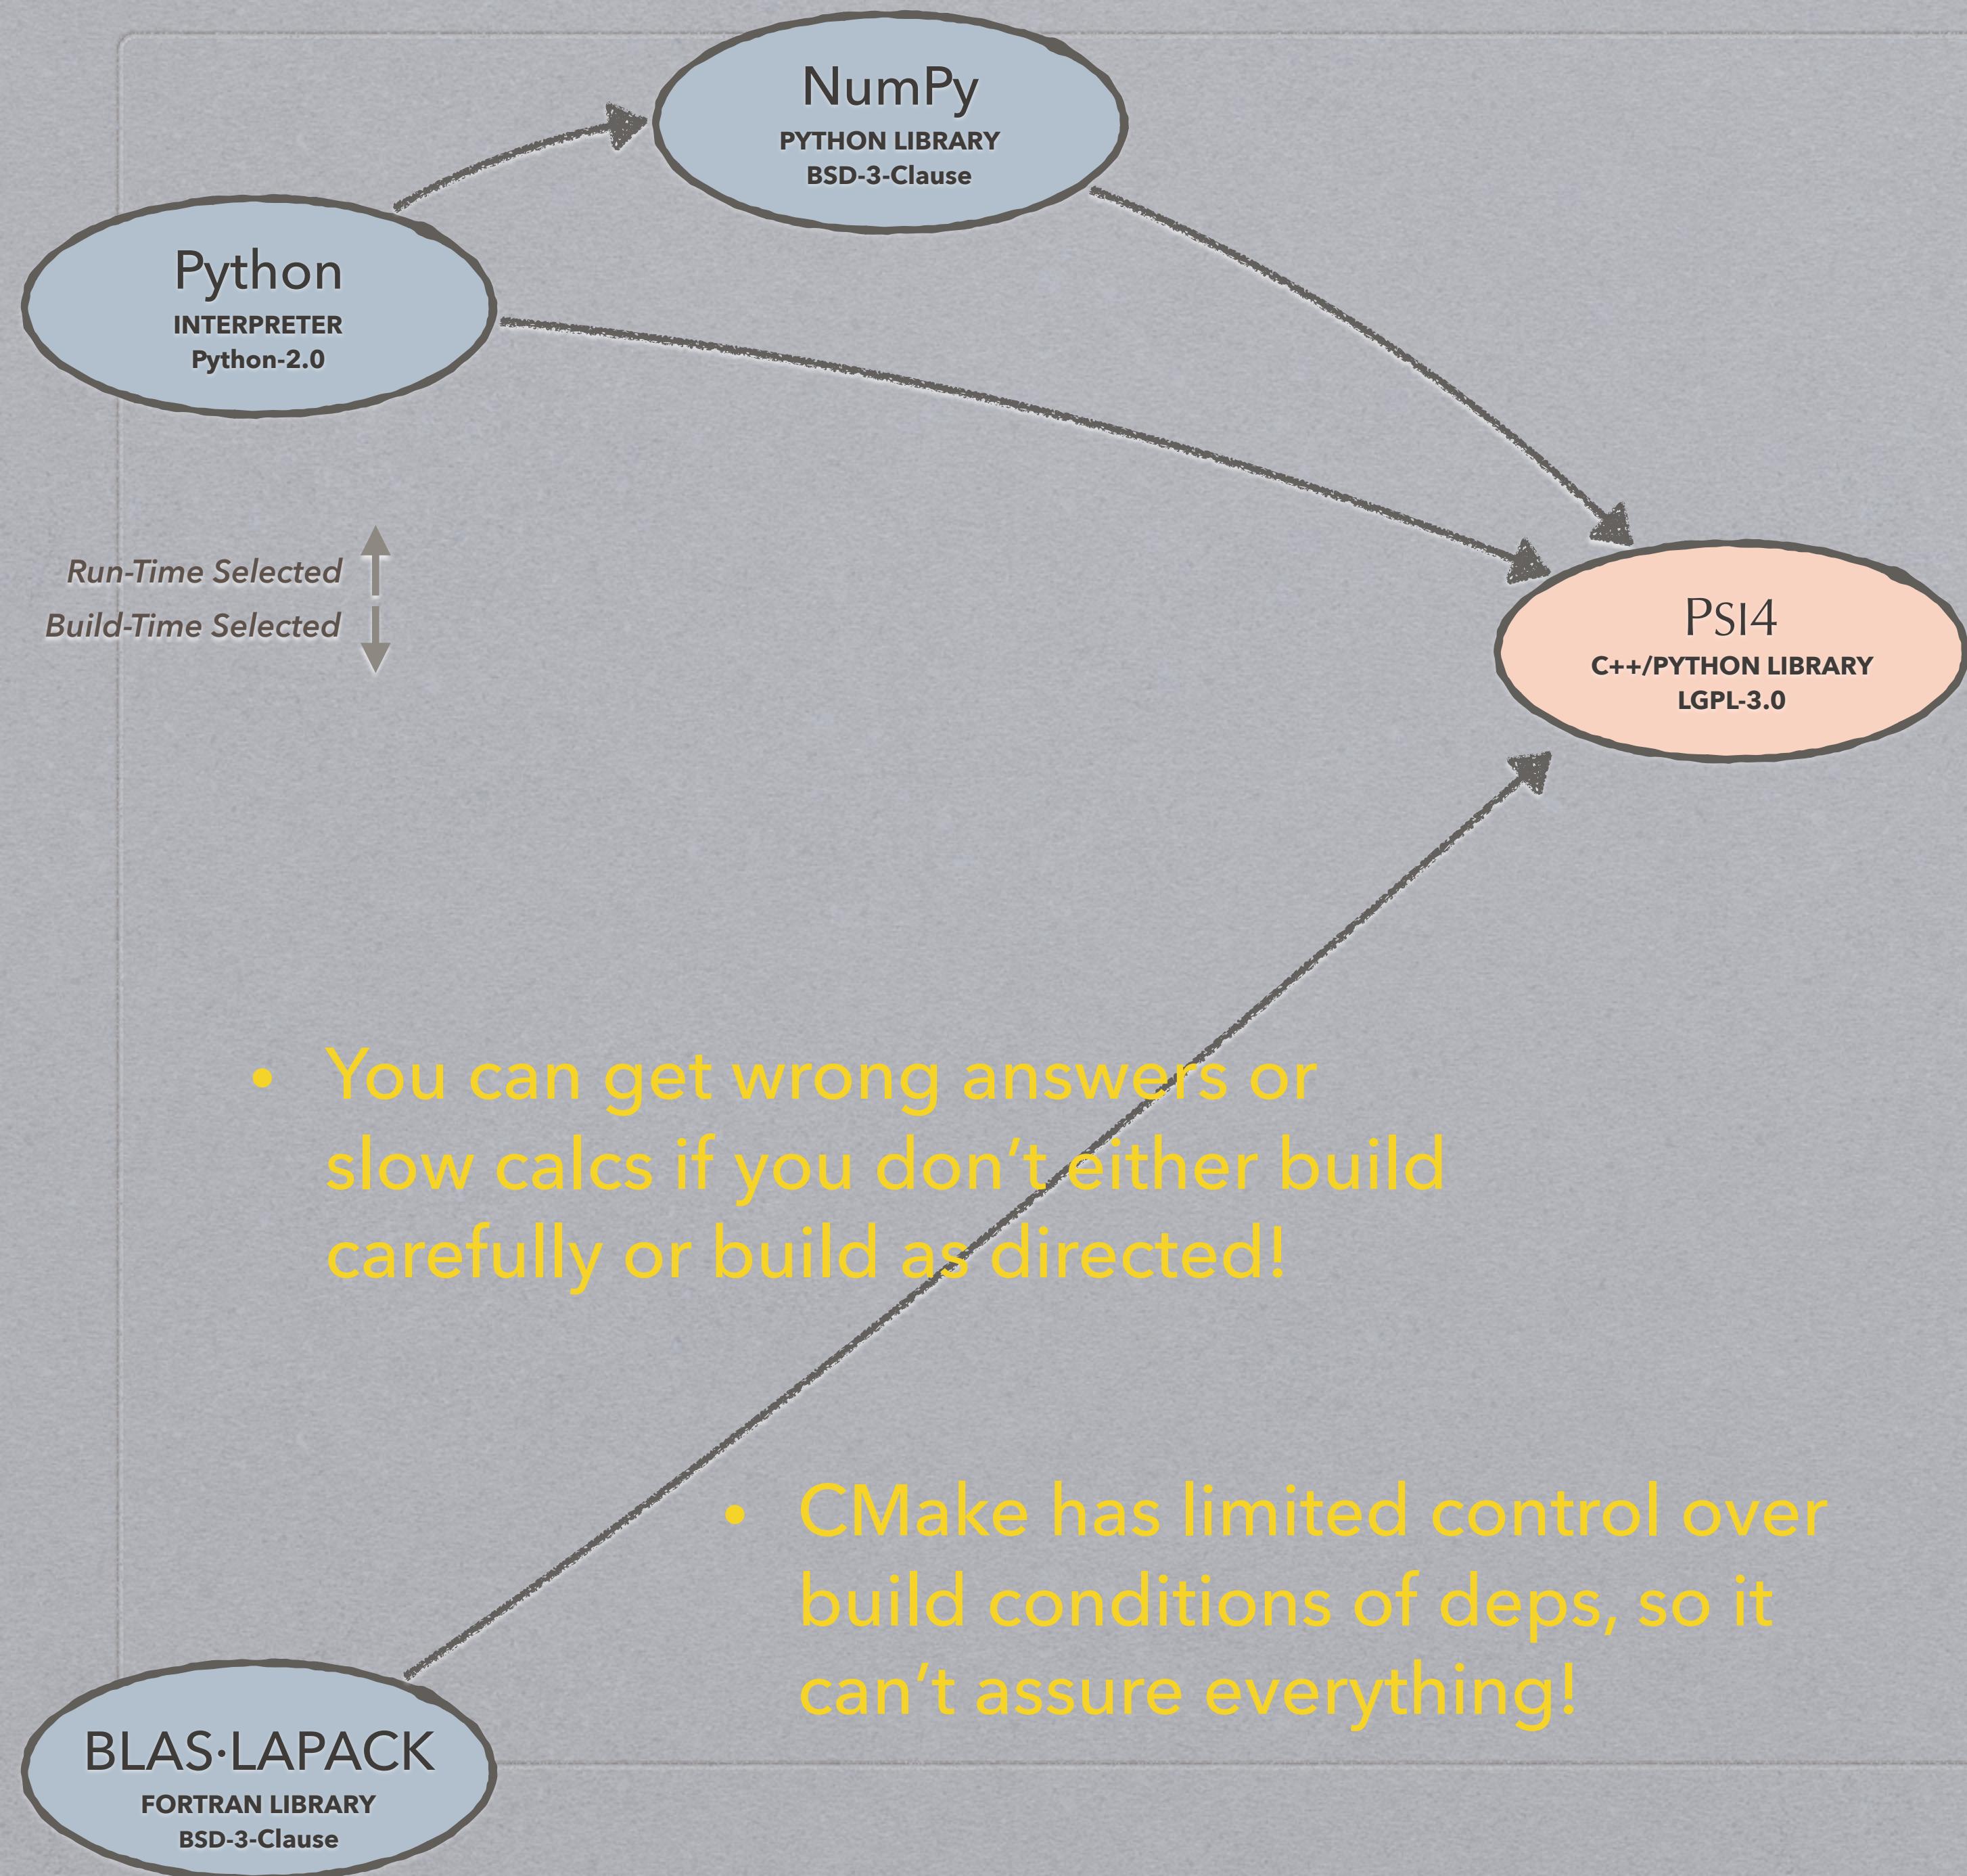
- Have ambit and Forte come to an agreement on the one true master branch?
- If we can add some internal versioning, are these ready to conda package?
- Forte, I believe, has pybound itself. More on potential complications later.
- Any other niftiness from Forte that should move back to Psi?

EFP COMPLICATIONS



- EFP presently octopus-like C-side
 - global EFP in P::e
 - parsing in libmints.Molecule
 - binary, library, if/else in SCF
- barrier to moving SCF iterations Py-side
- need Py-side access to EFP geom
- PylibEFP as intermediate
- new parsing module to replace Molecule::create_molecule_from_string

INSIDIOUS LAPACK



- Previously for conda packages (and you can still arrange your dev build this way), used Accelerate for Mac and statically linked MKL for Linux.
- Most ppl used NumPy from default channel linked to intel_thread/core.
- Initially noticed by RMcG, NumPy linalg errors show up with mixed LAPACK (and I mean wrong answers, not notification that env is wrong) and also threading performance can be suppressed if first lib loaded isn't multi architecture capable.
- Conda pkgs reworked to link to mkl_rt. **bin/psi4-path-advisor** sets this up for you, too.
- Numpy from default conda channel links to intel_thread/core, not rt, so import order starts to become important.
- Numpy from intel conda channel links to mkl_rt, so more robust wrt threading. Recommended. Now available for py27, 35, 36.
- Could avoid (but not eliminate) many user problems by importing numpy first in psithon, at cost of losing "how to get numpy" error msg.

OPENMP

BLAS·LAPACK
FORTRAN LIBRARY
BSD-3-Clause

```
## [Linux|Mac]/compiler_fam/[dynamic|single_dynamic_lib]/[native|no_choice|intel]_OpenMP_threading
+
## (a) = iomp5 only choice with icpc and that's self-contained but want the extra flags so g++-based plugins run correctly
## (b) = clang-based so no intervention req'd
## (c) = to satisfy OpenMP symbols and suppress libgomp
+
## L-gnu/Dyn/N      -lmkl_intel_lp64 -lmkl_gnu_thread   -lmkl_core -lgomp          -lpthread -lm -ldl
## L/pgi/Dyn/N      -lmkl_intel_lp64 -lmkl_pgi_thread   -lmkl_core -pgf90libs -mp -lpthread -lm -ldl
+
## L/intel/Dyn/-    -lmkl_intel_lp64 -lmkl_intel_thread -lmkl_core -liomp5 -lpthread -lm -ldl           IN
## M/intel/Dyn/-    -lmkl_intel_lp64 -lmkl_intel_thread -lmkl_core -liomp5 -lpthread -lm -ldl           IN
## L-gnu/Dyn/I       -lmkl_intel_lp64 -lmkl_intel_thread -lmkl_core -liomp5 -lpthread -lm -ldl           I(c):[-fno-openmp] N(above)
## M-gnu/Dyn/-       -lmkl_intel_lp64 -lmkl_intel_thread -lmkl_core -liomp5 -lpthread -lm -ldl           I(c):[-fno-openmp] N(n/a):
## L/pgi/Dyn/I       -lmkl_intel_lp64 -lmkl_intel_thread -lmkl_core -liomp5 -lpthread -lm -ldl           I:? N(above)
## M/pgi/Dyn/-       -lmkl_intel_lp64 -lmkl_intel_thread -lmkl_core -liomp5 -lpthread -lm -ldl           I:? N(n/a):
## M/clang/Dyn/-     -lmkl_intel_lp64 -lmkl_intel_thread -lmkl_core -liomp5 -lpthread -lm -ldl           IN(b):[]
+
## L/intel/SDL/-    -lmkl_rt -lpthread -lm -ldl           IN(a):[-liomp5 -fno-openmp]
## M/intel/SDL/-    -lmkl_rt -lpthread -lm -ldl           IN(b):[]
## L-gnu/SDL/-       -lmkl_rt -lpthread -lm -ldl           I(c):[-liomp5 -fno-openmp] N: []
## M-gnu/SDL/-       -lmkl_rt -lpthread -lm -ldl           I(c):[-liomp5 -fno-openmp] N(n/a):
## L/pgi/SDL/-       -lmkl_rt -lpthread -lm -ldl           ?
## M/pgi/SDL/-       -lmkl_rt -lpthread -lm -ldl           ?
## M/clang/SDL/-     -lmkl_rt -lpthread -lm -ldl           IN(b):[]
```

C++/OpenMP

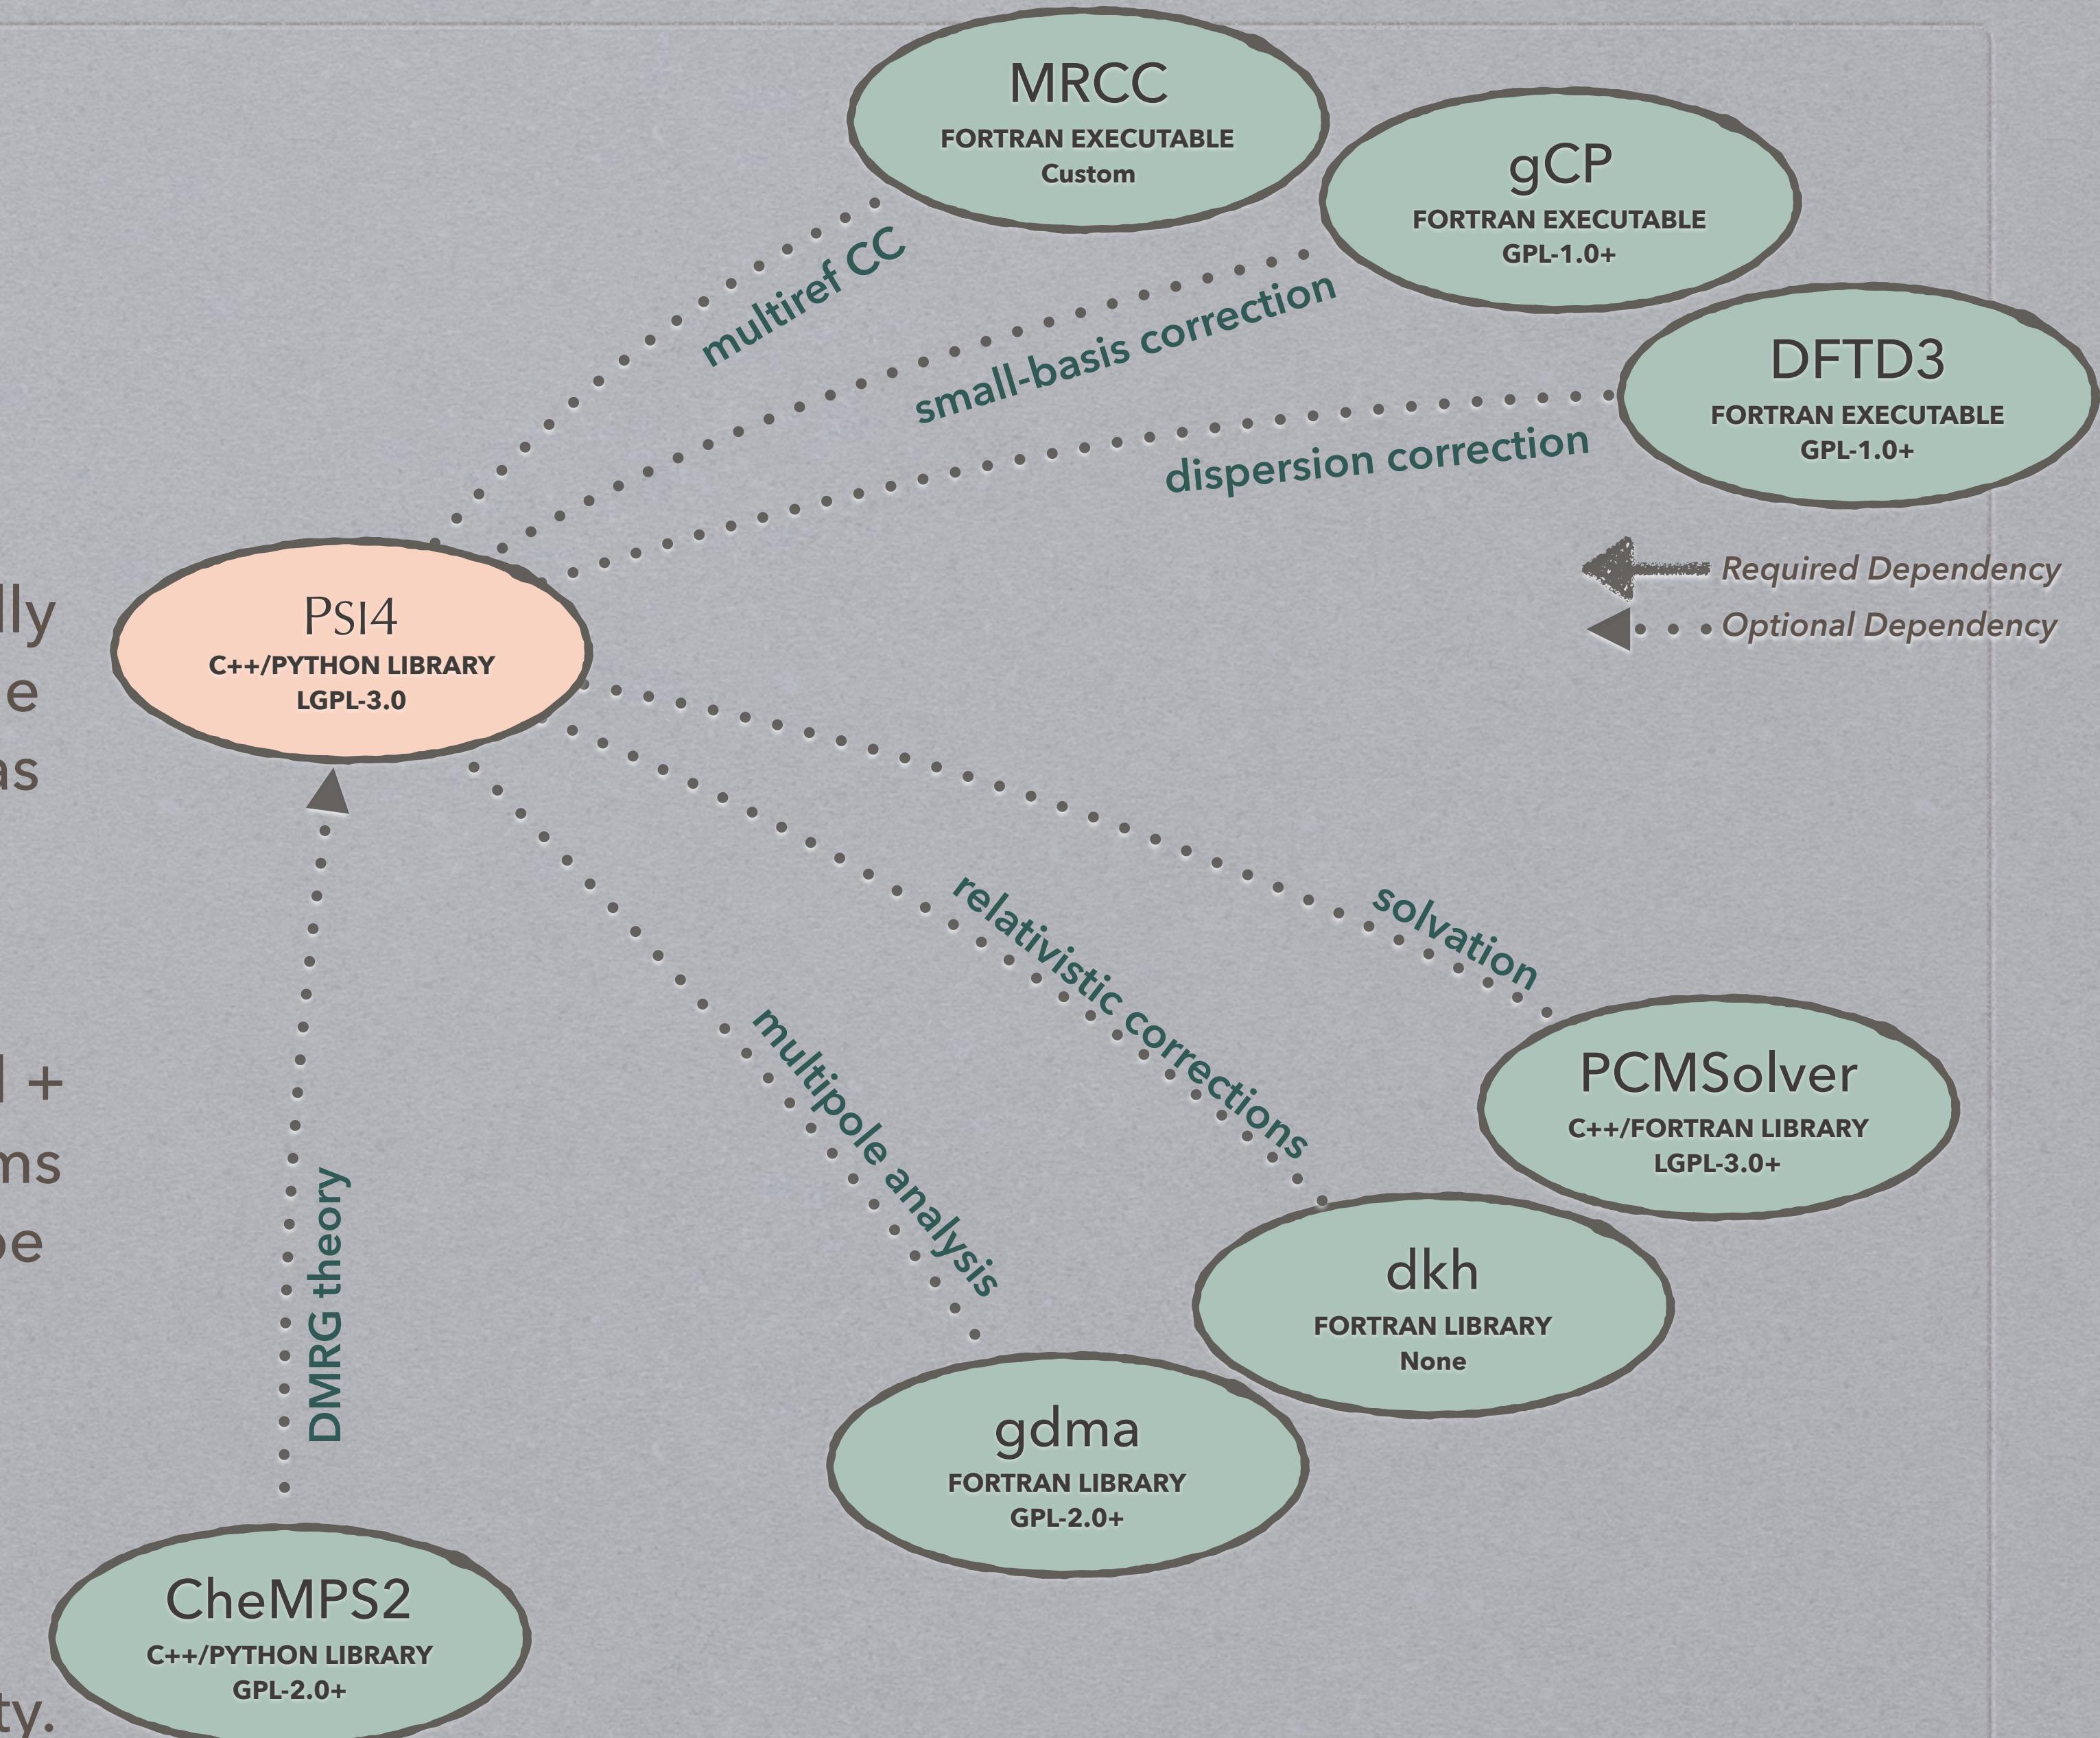
PSI4
C++/PYTHON LIBRARY
LGPL-3.0

- On Mac, AppleClang has no OMP. Can build with GCC w/ OMP, but it's slow and painful to invest in meagerly supported software. For conda have to make separate sse41 packages for old hardware.
- On Linux, Intel compilers provide **-msse2 -axCORE- AVX2, AVX** flags that optimize code for multiple architectures. For conda, we presently rely on a old and custom gcc5.2 package.
- LAPACK & OMP are sadly intertwined, and current released FindOpenMP.cmake is not too helpful.
- Conda has a new compiler toolchain with very recent GCC 7 on Linux and Clang 4.0.1 on Mac (so new Intel 2018 can't handle its headers). Note that that's real Clang with OMP support (it works). All these are conda-installable, so no more Xcode Clang or system GCC.
- Once Intel Mac gets with the program, plan is to
 - drop Mac GCC and use Intel atop Clang with OMP & multiarch, thus building all-in-one pkgs like Linux
 - need legitimate Mac Intel license once path is validated
- Once I get time, Linux plan is to
 - drop current gcc5.2 and move to Intel atop GCC 7, continuing multiarch
- OpenMP detection in CMake is improving at >3.9. If it improves enough, will rework Math/OMP detection to use it and bump CMake min from 3.3 to at least 3.6

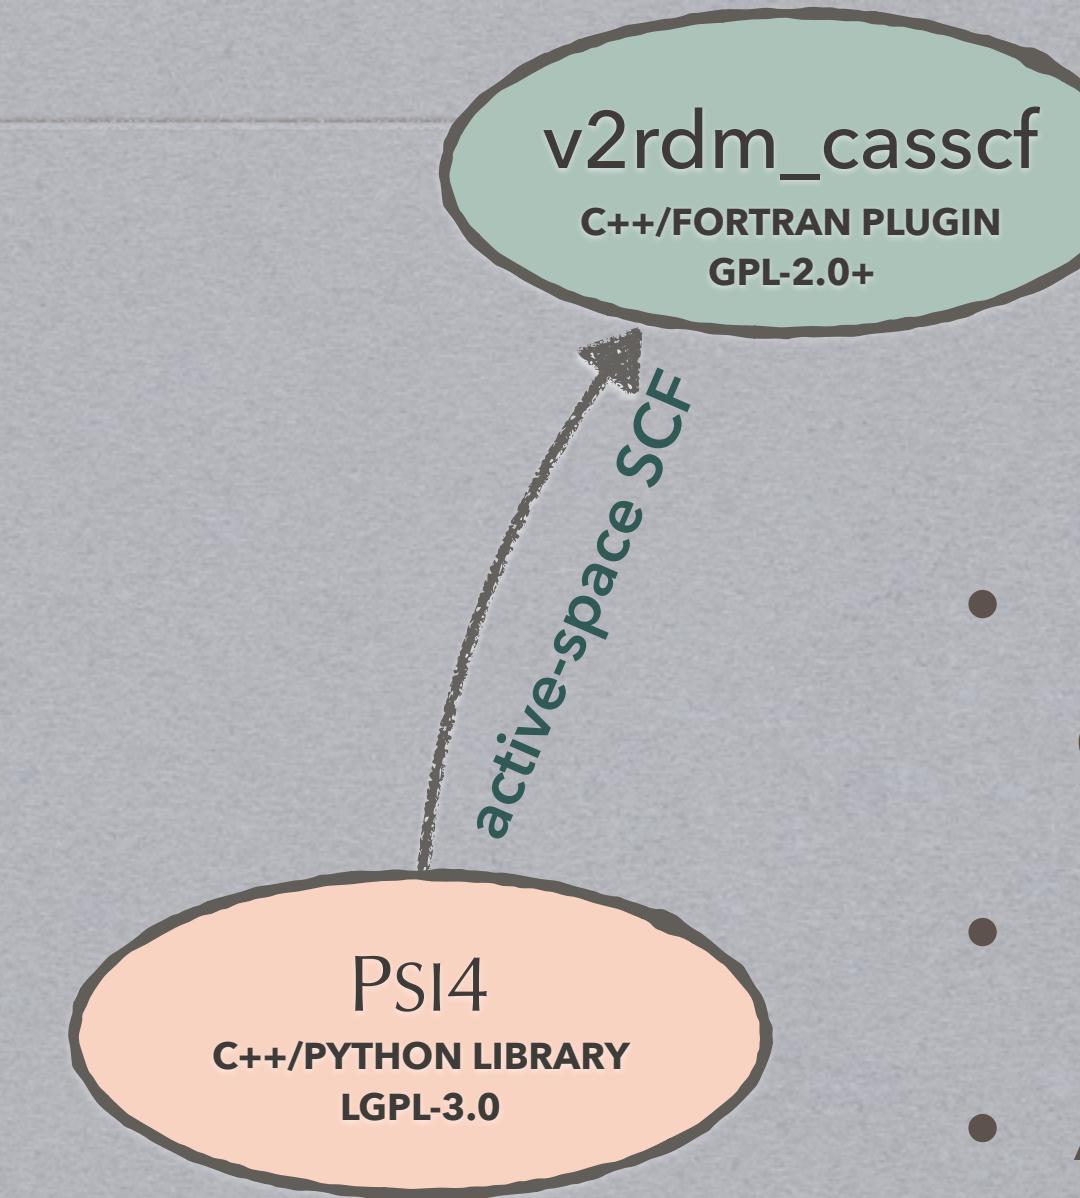
CMake
C++ EXECUTABLE
BSD-3-Clause

ADDONS

- CheMPS2 continues gamely, though I doubt it gets much use through Psi b/c parallelism.
- GDMA is great (except for that string length thing) but can't be updated trivially b/c bin → lib conversion. Persuade Stone to separate into library and driver? Call as exe through Py instead?
- DKH, no complaints
- PCMSolver I suspect could get a piebald + C++ lib treatment, but other QC programs probably constraining it. Should its I/O be JSON-ified? Can we call it from Py-side when SCF iterations move thither?
- DFTD3 good for now, but the DFTD4 paper is out (haven't seen any code).
- gCP could use mods to improve flexibility.
- MRCC, there when you're desperate.

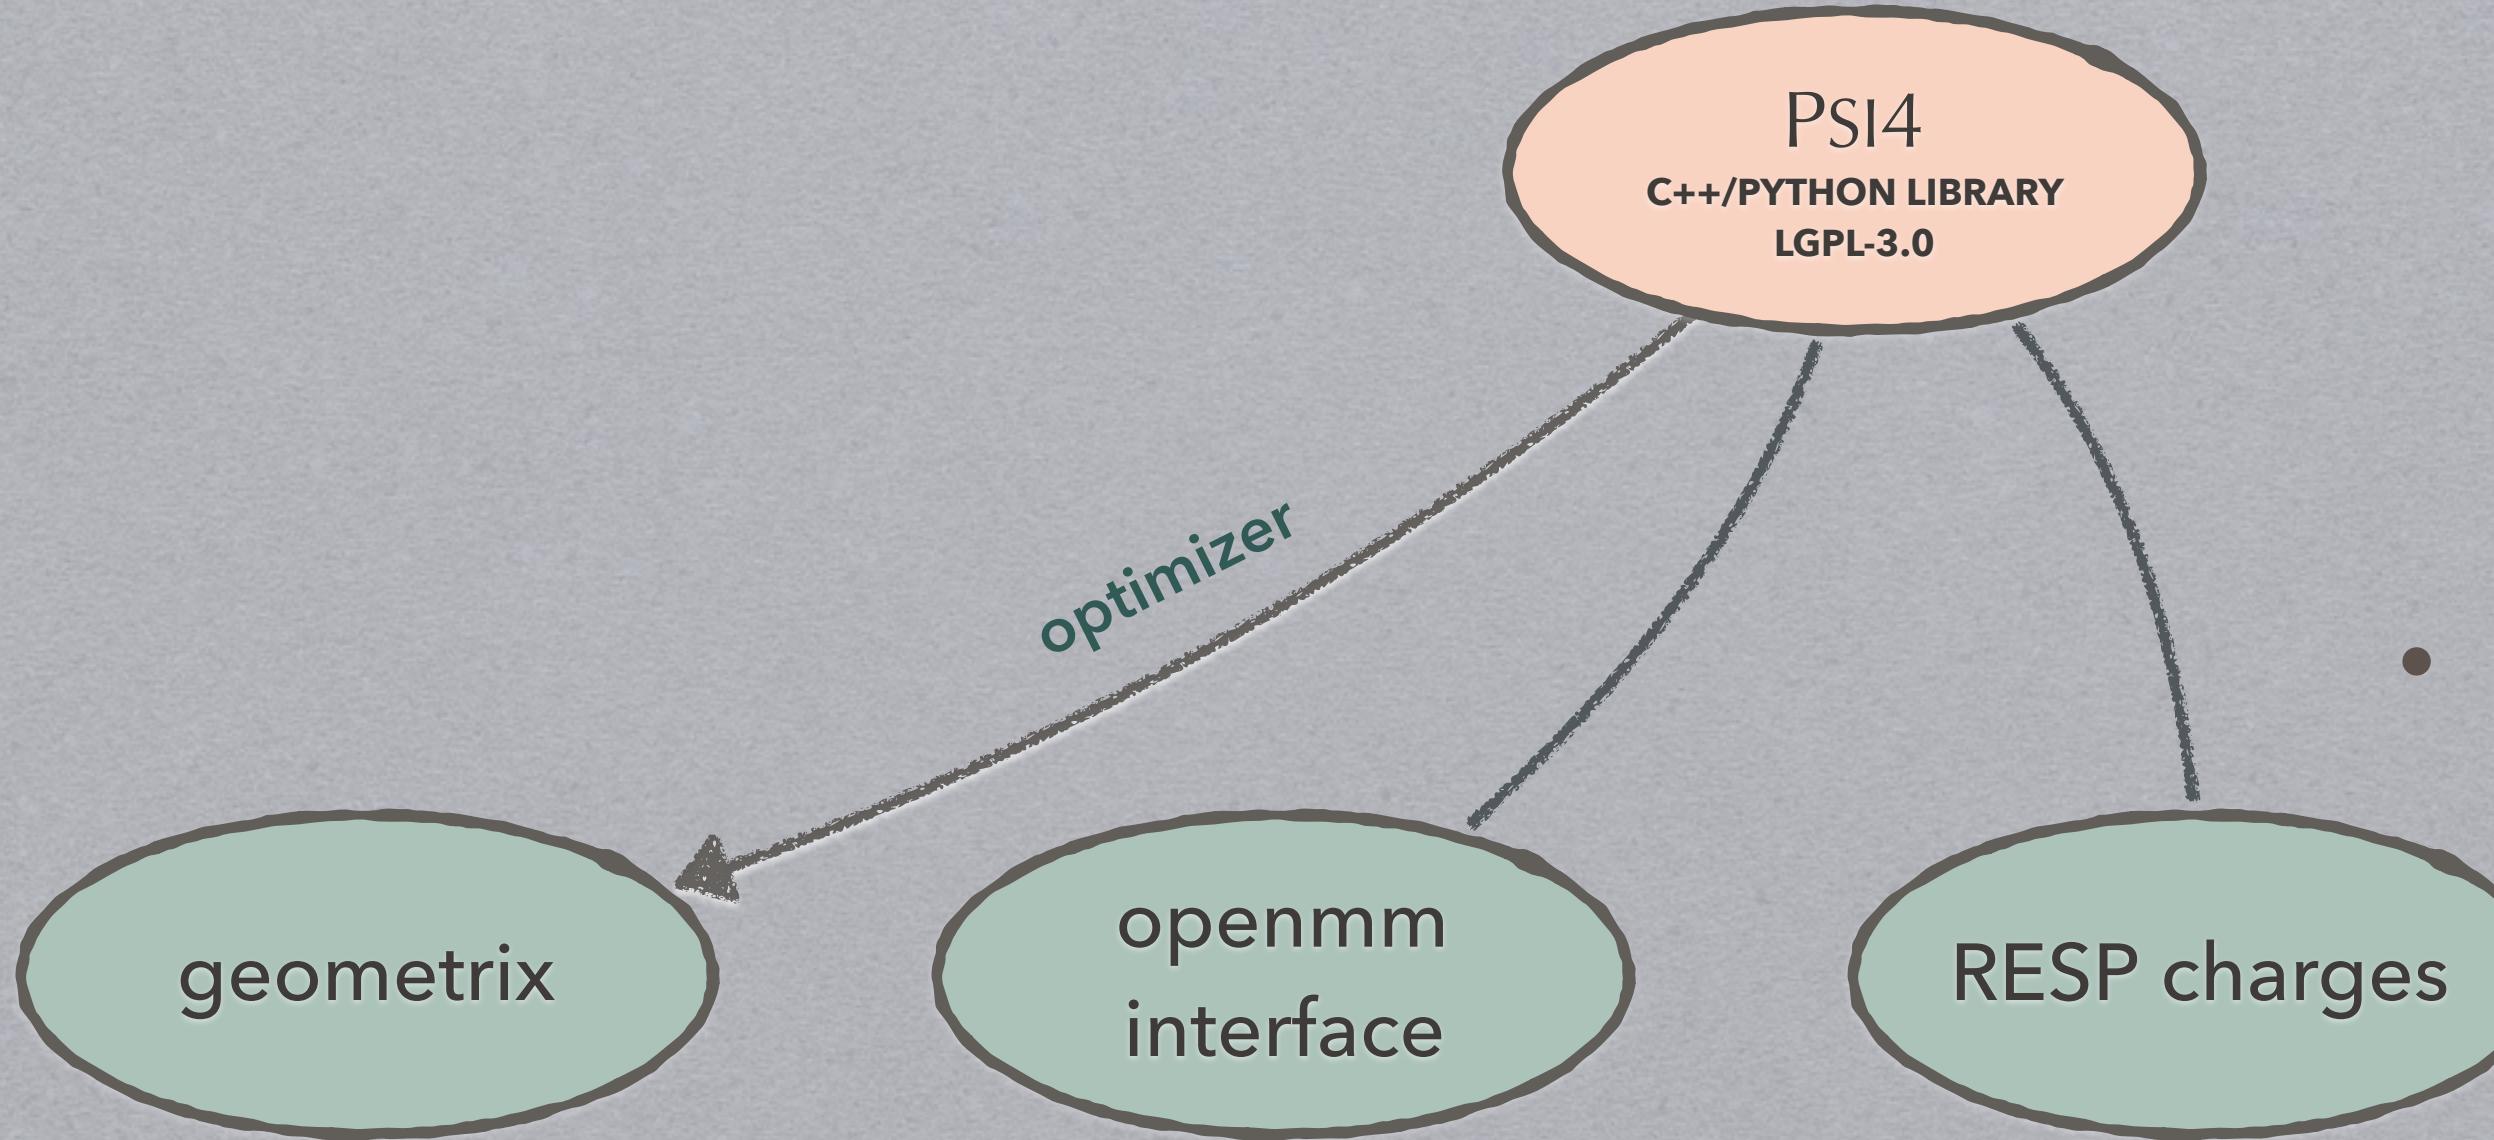


PLUGINS



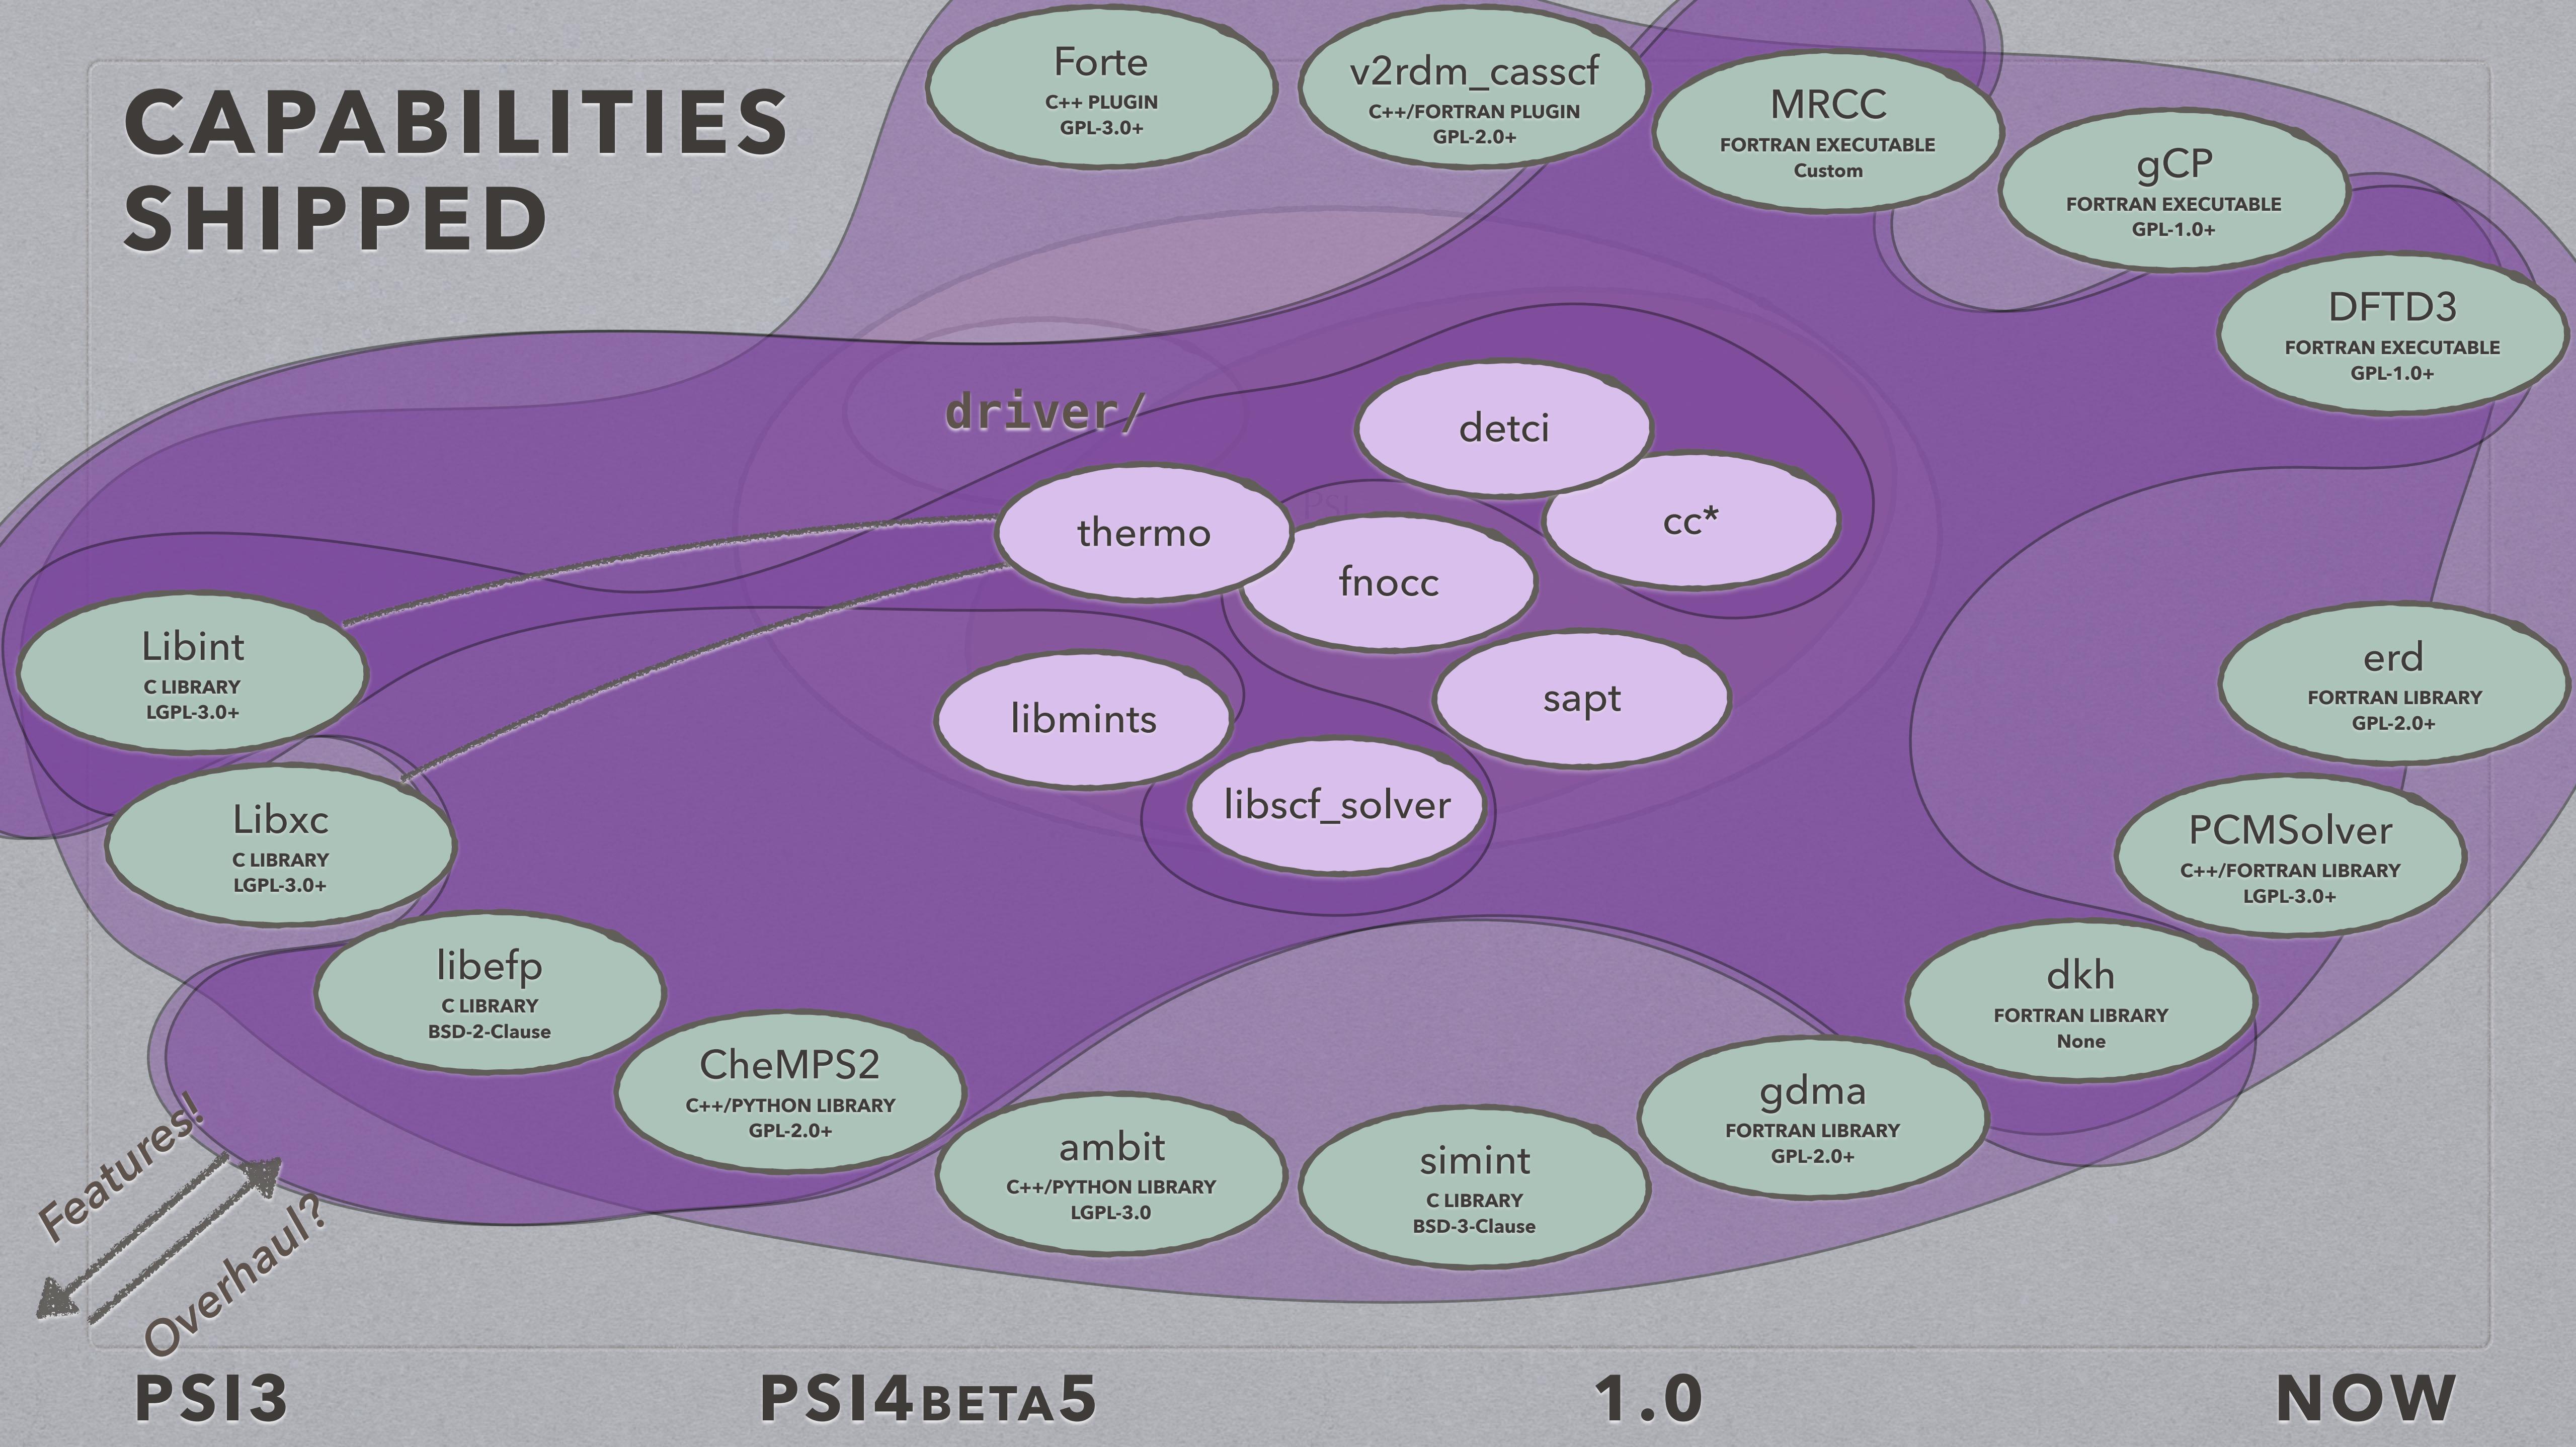
- Eugene newly updated with gradients
- I haven't reintegrated into conda
- Among those suffering from ABI changes w/o version changes
- Pain that our plugins feel wrt ABI/API breakage is akin to what any upstream package feels

DESERVING PROJECTS



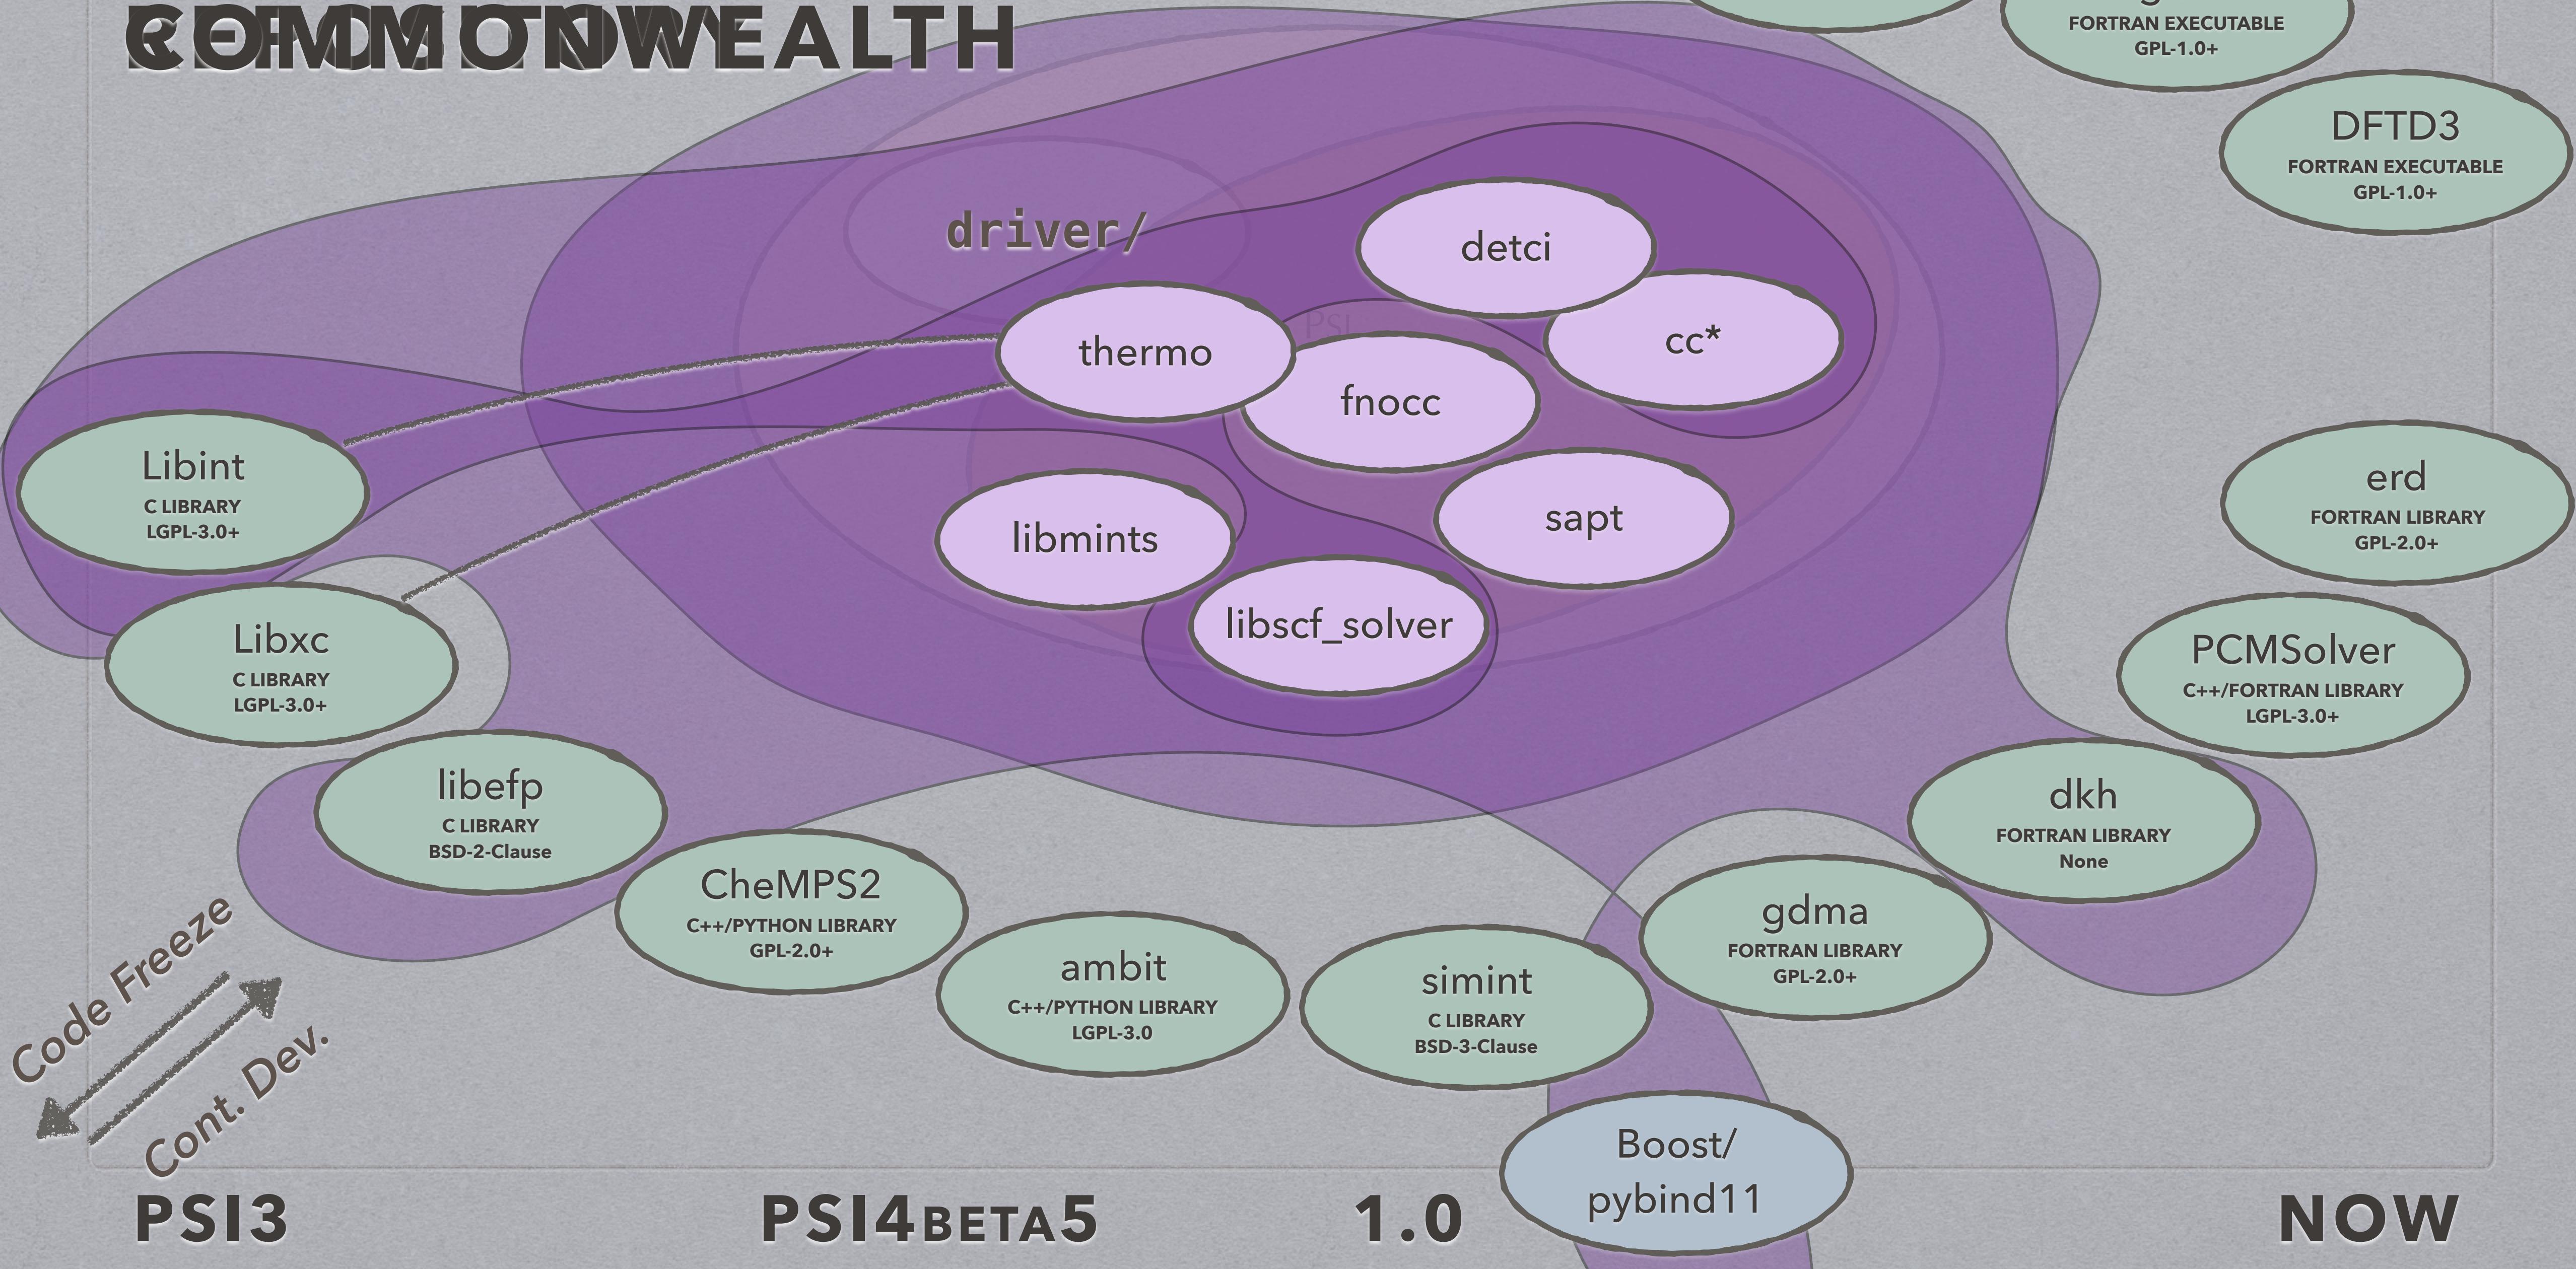
- I'd like to make even pure-Py projects build-on-demand (**cmake -DENABLE_v2rdm_casscf**), even though they're only interfaced at RT
- others, I'm sure

CAPABILITIES SHIPPED



EMPERENTOSI4

REDFM DOWEALTH



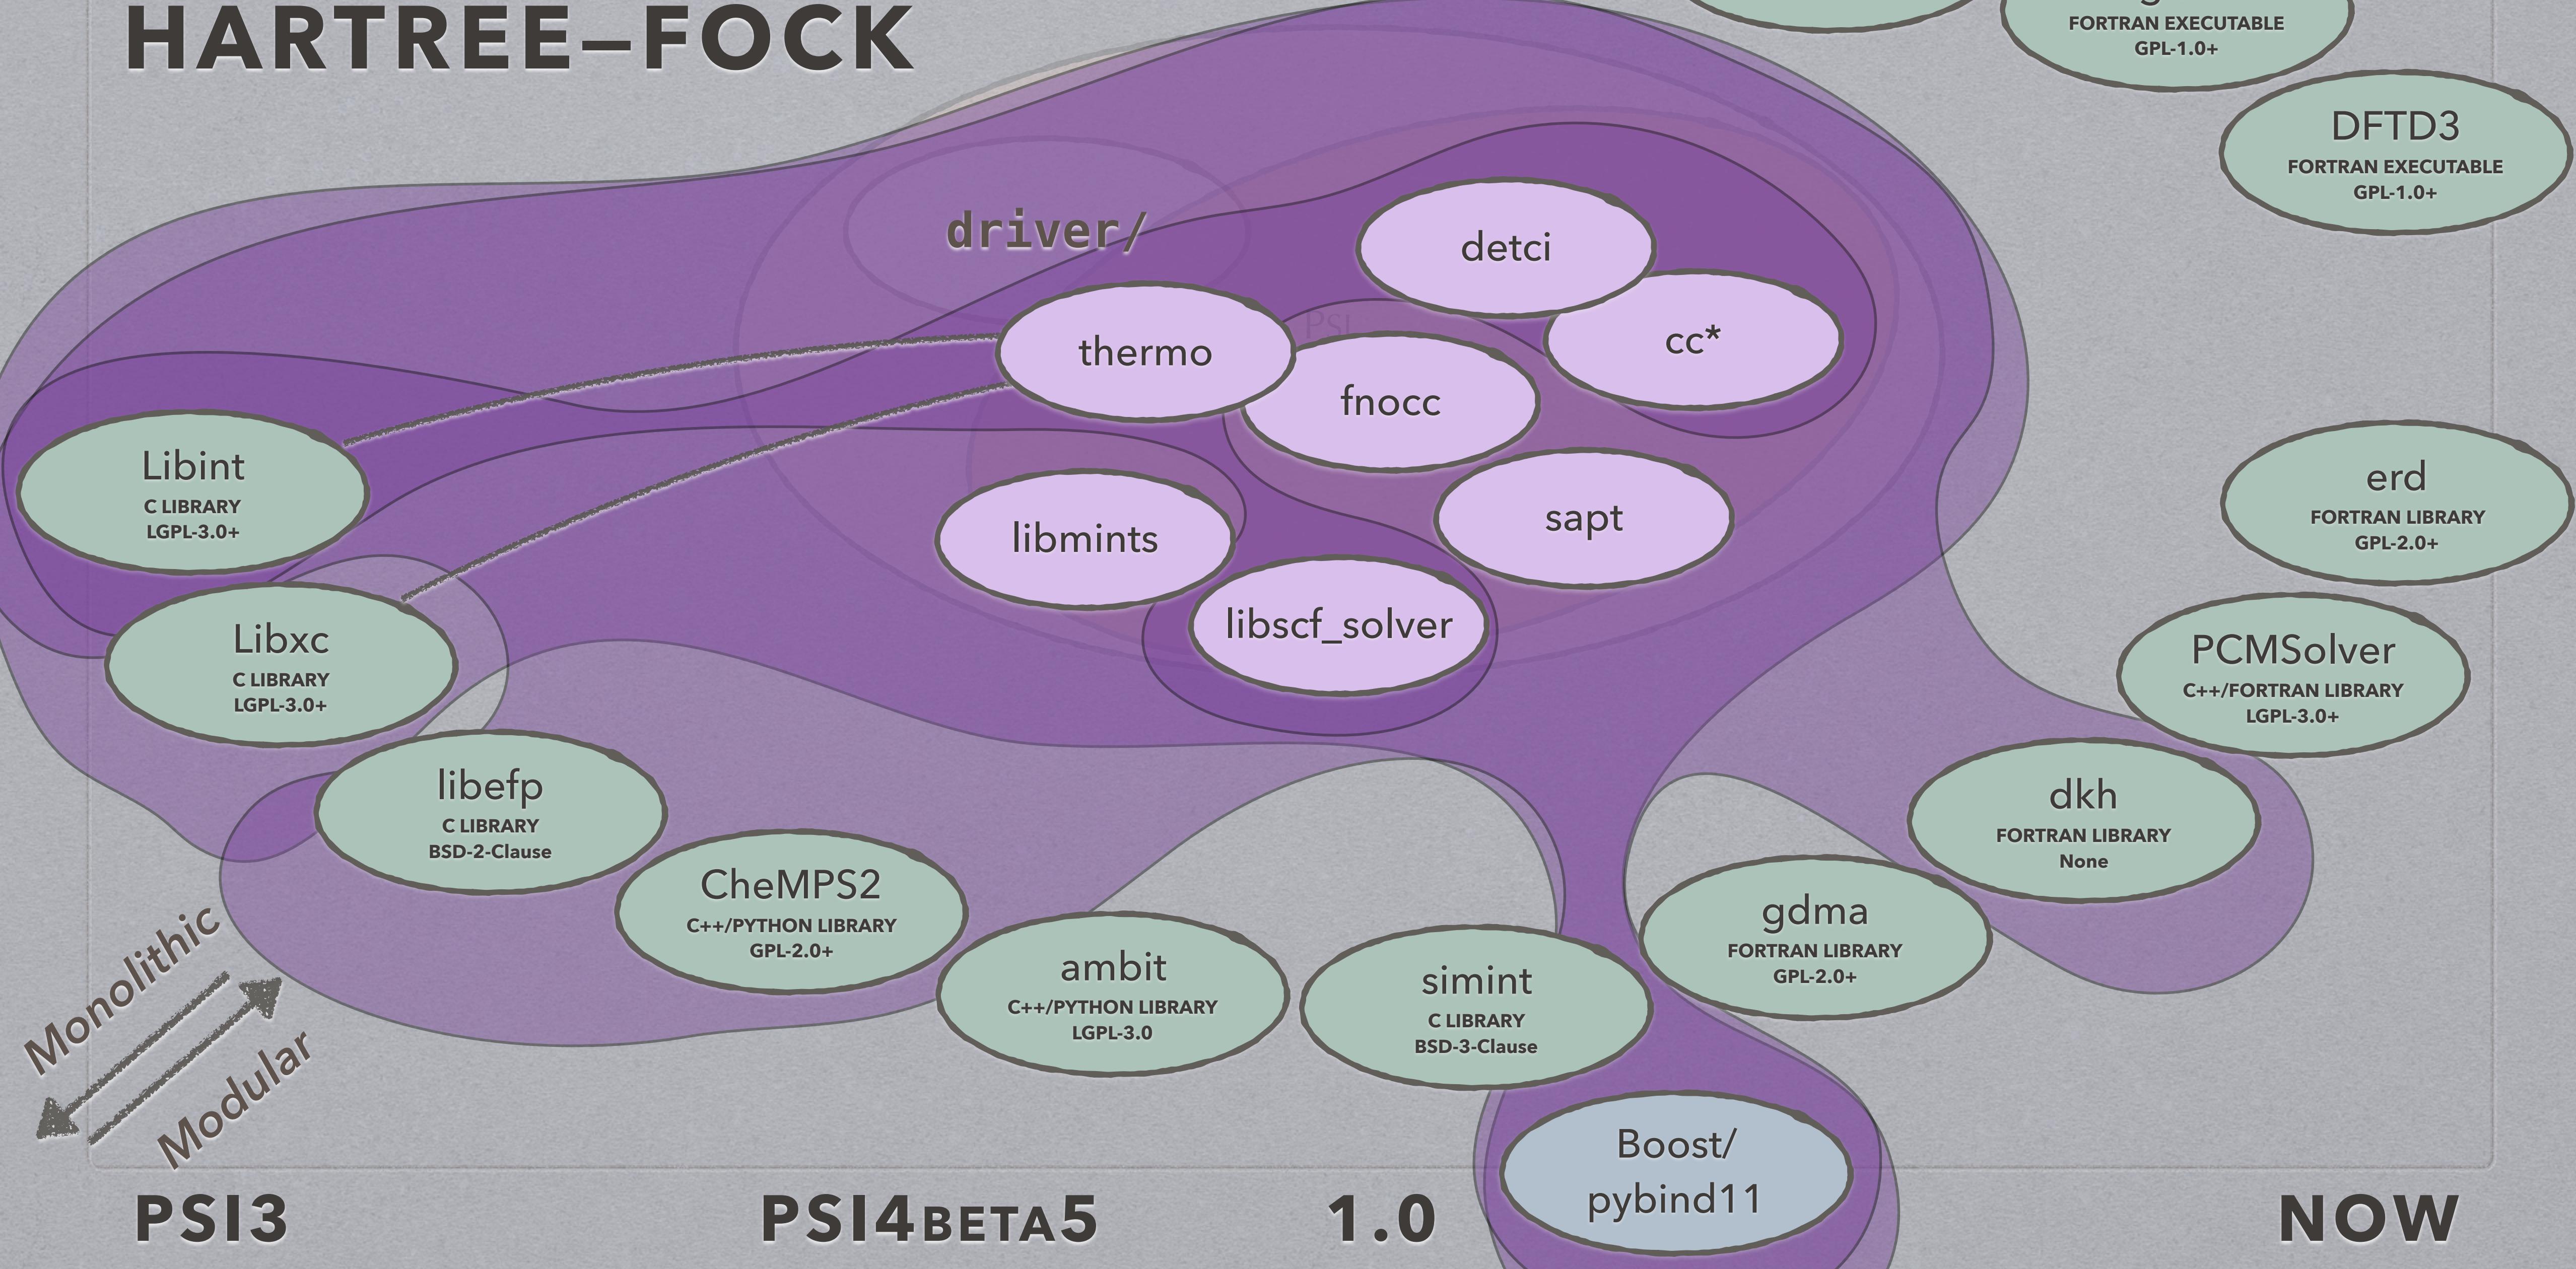
PSI3

PSI4 BETA5

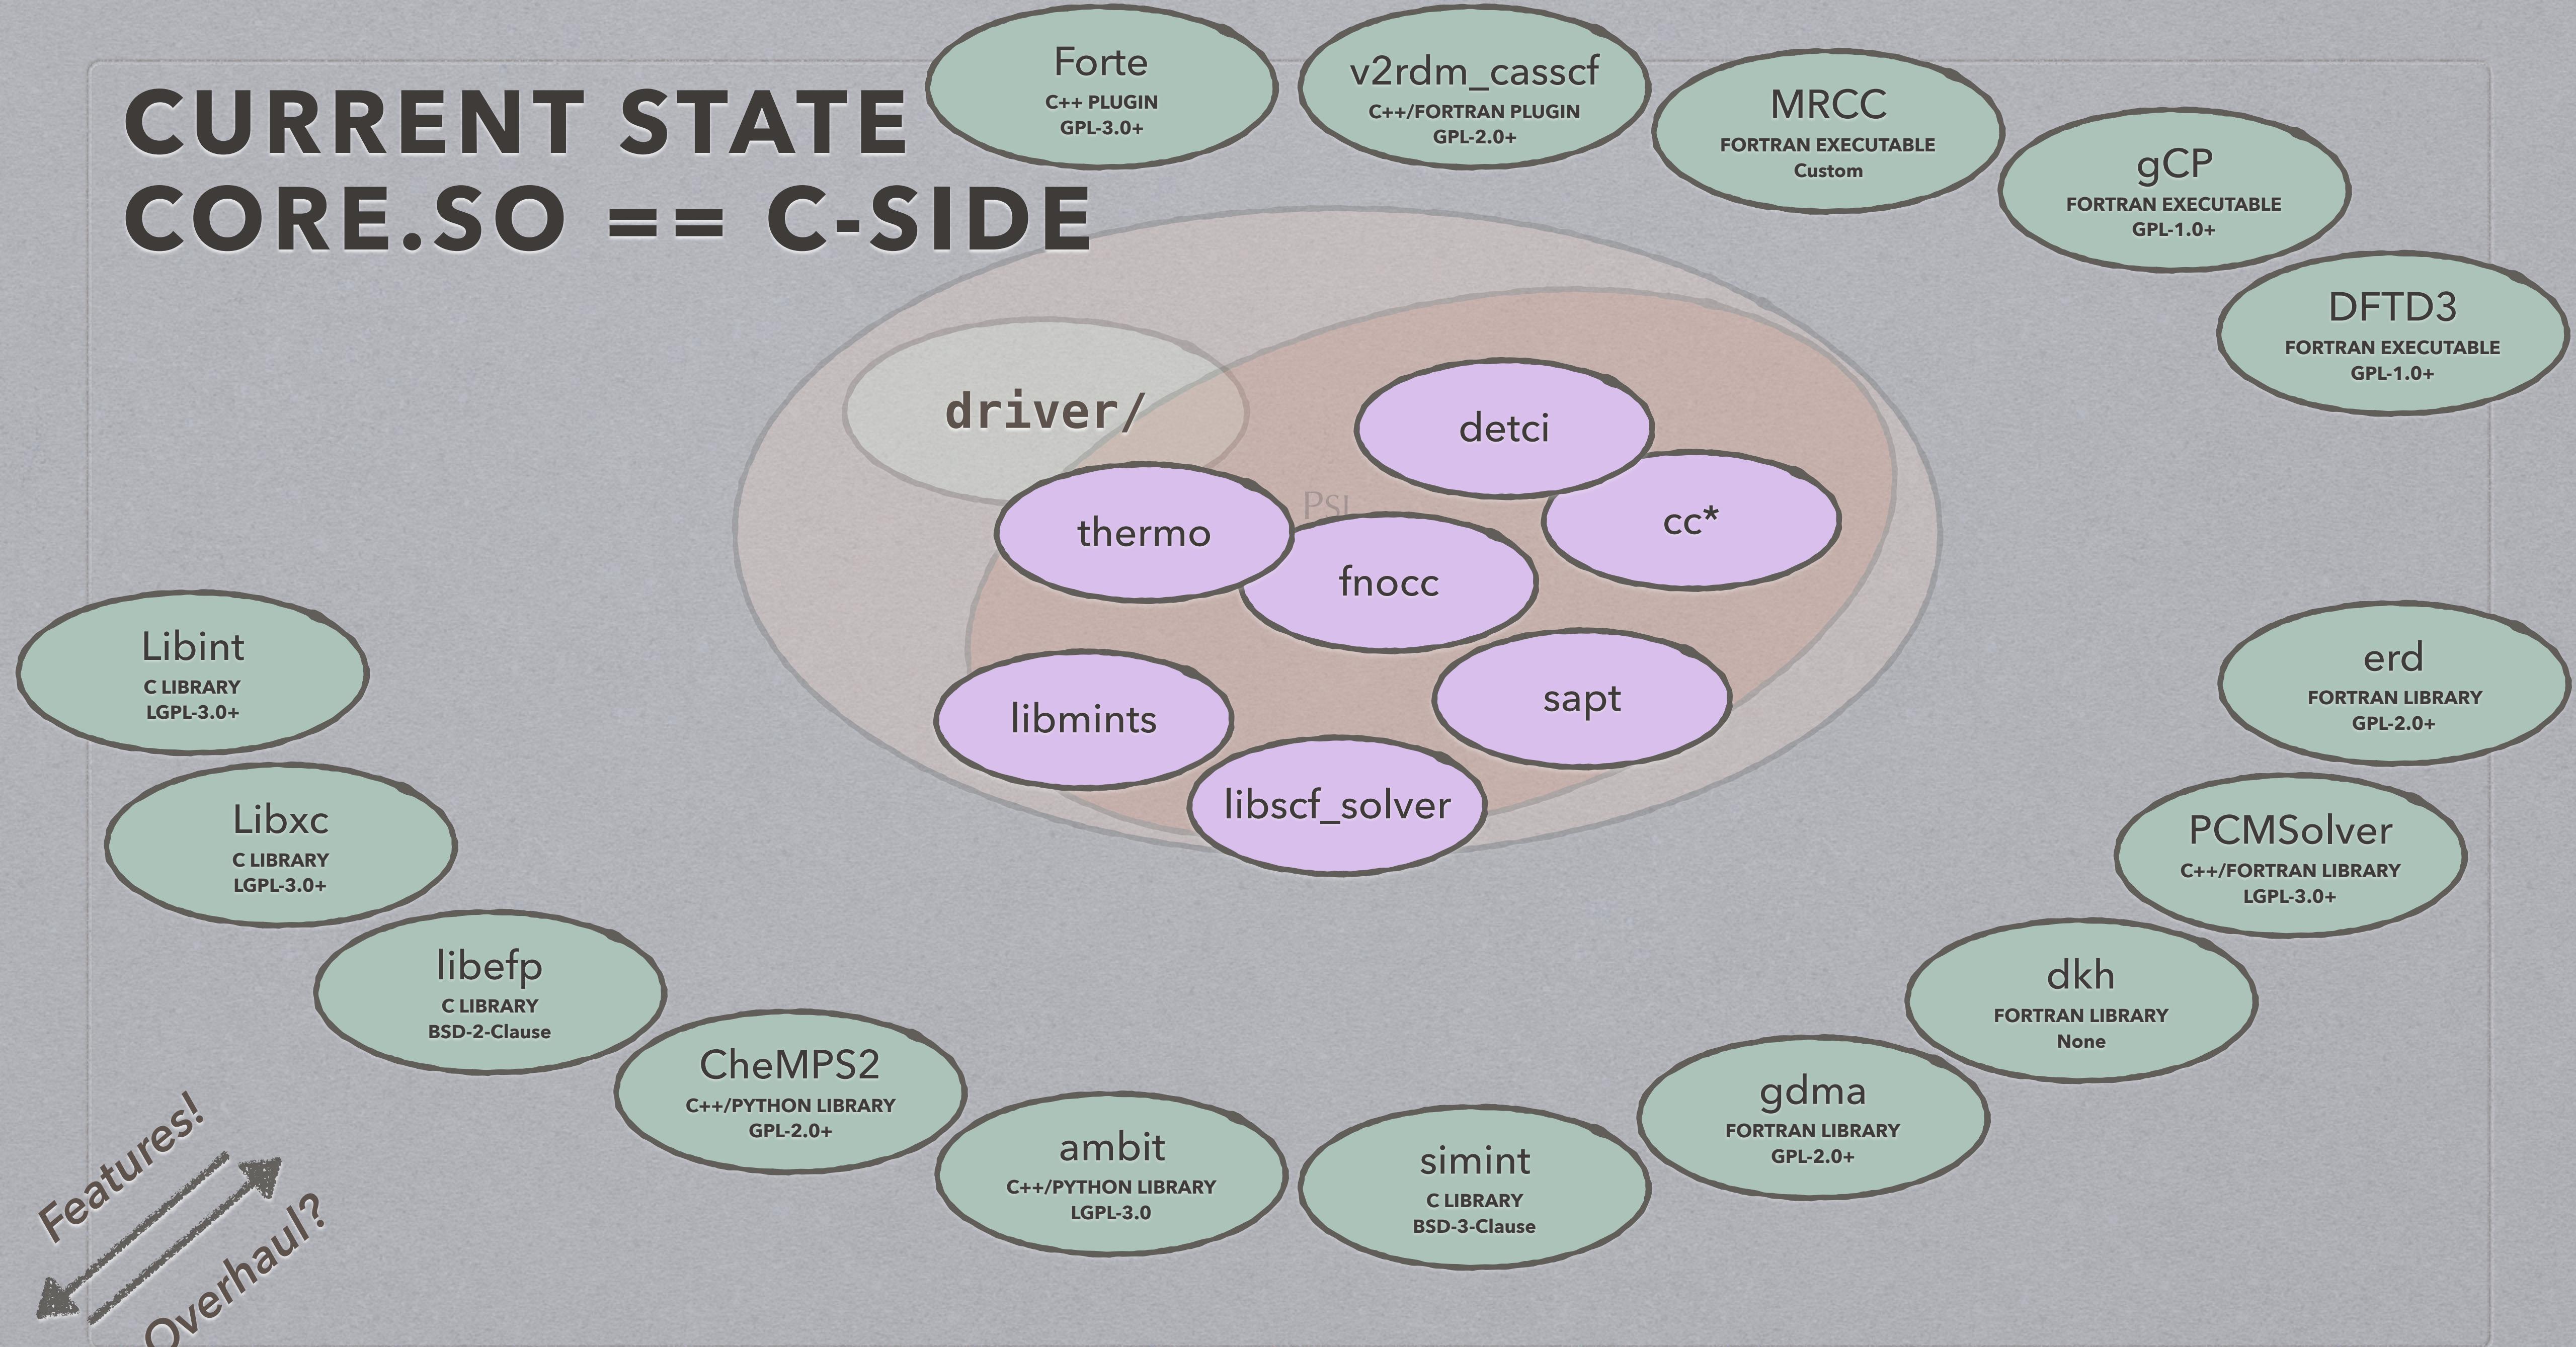
1.0

NOW

COMPILE TO RUN HARTREE-FOCK

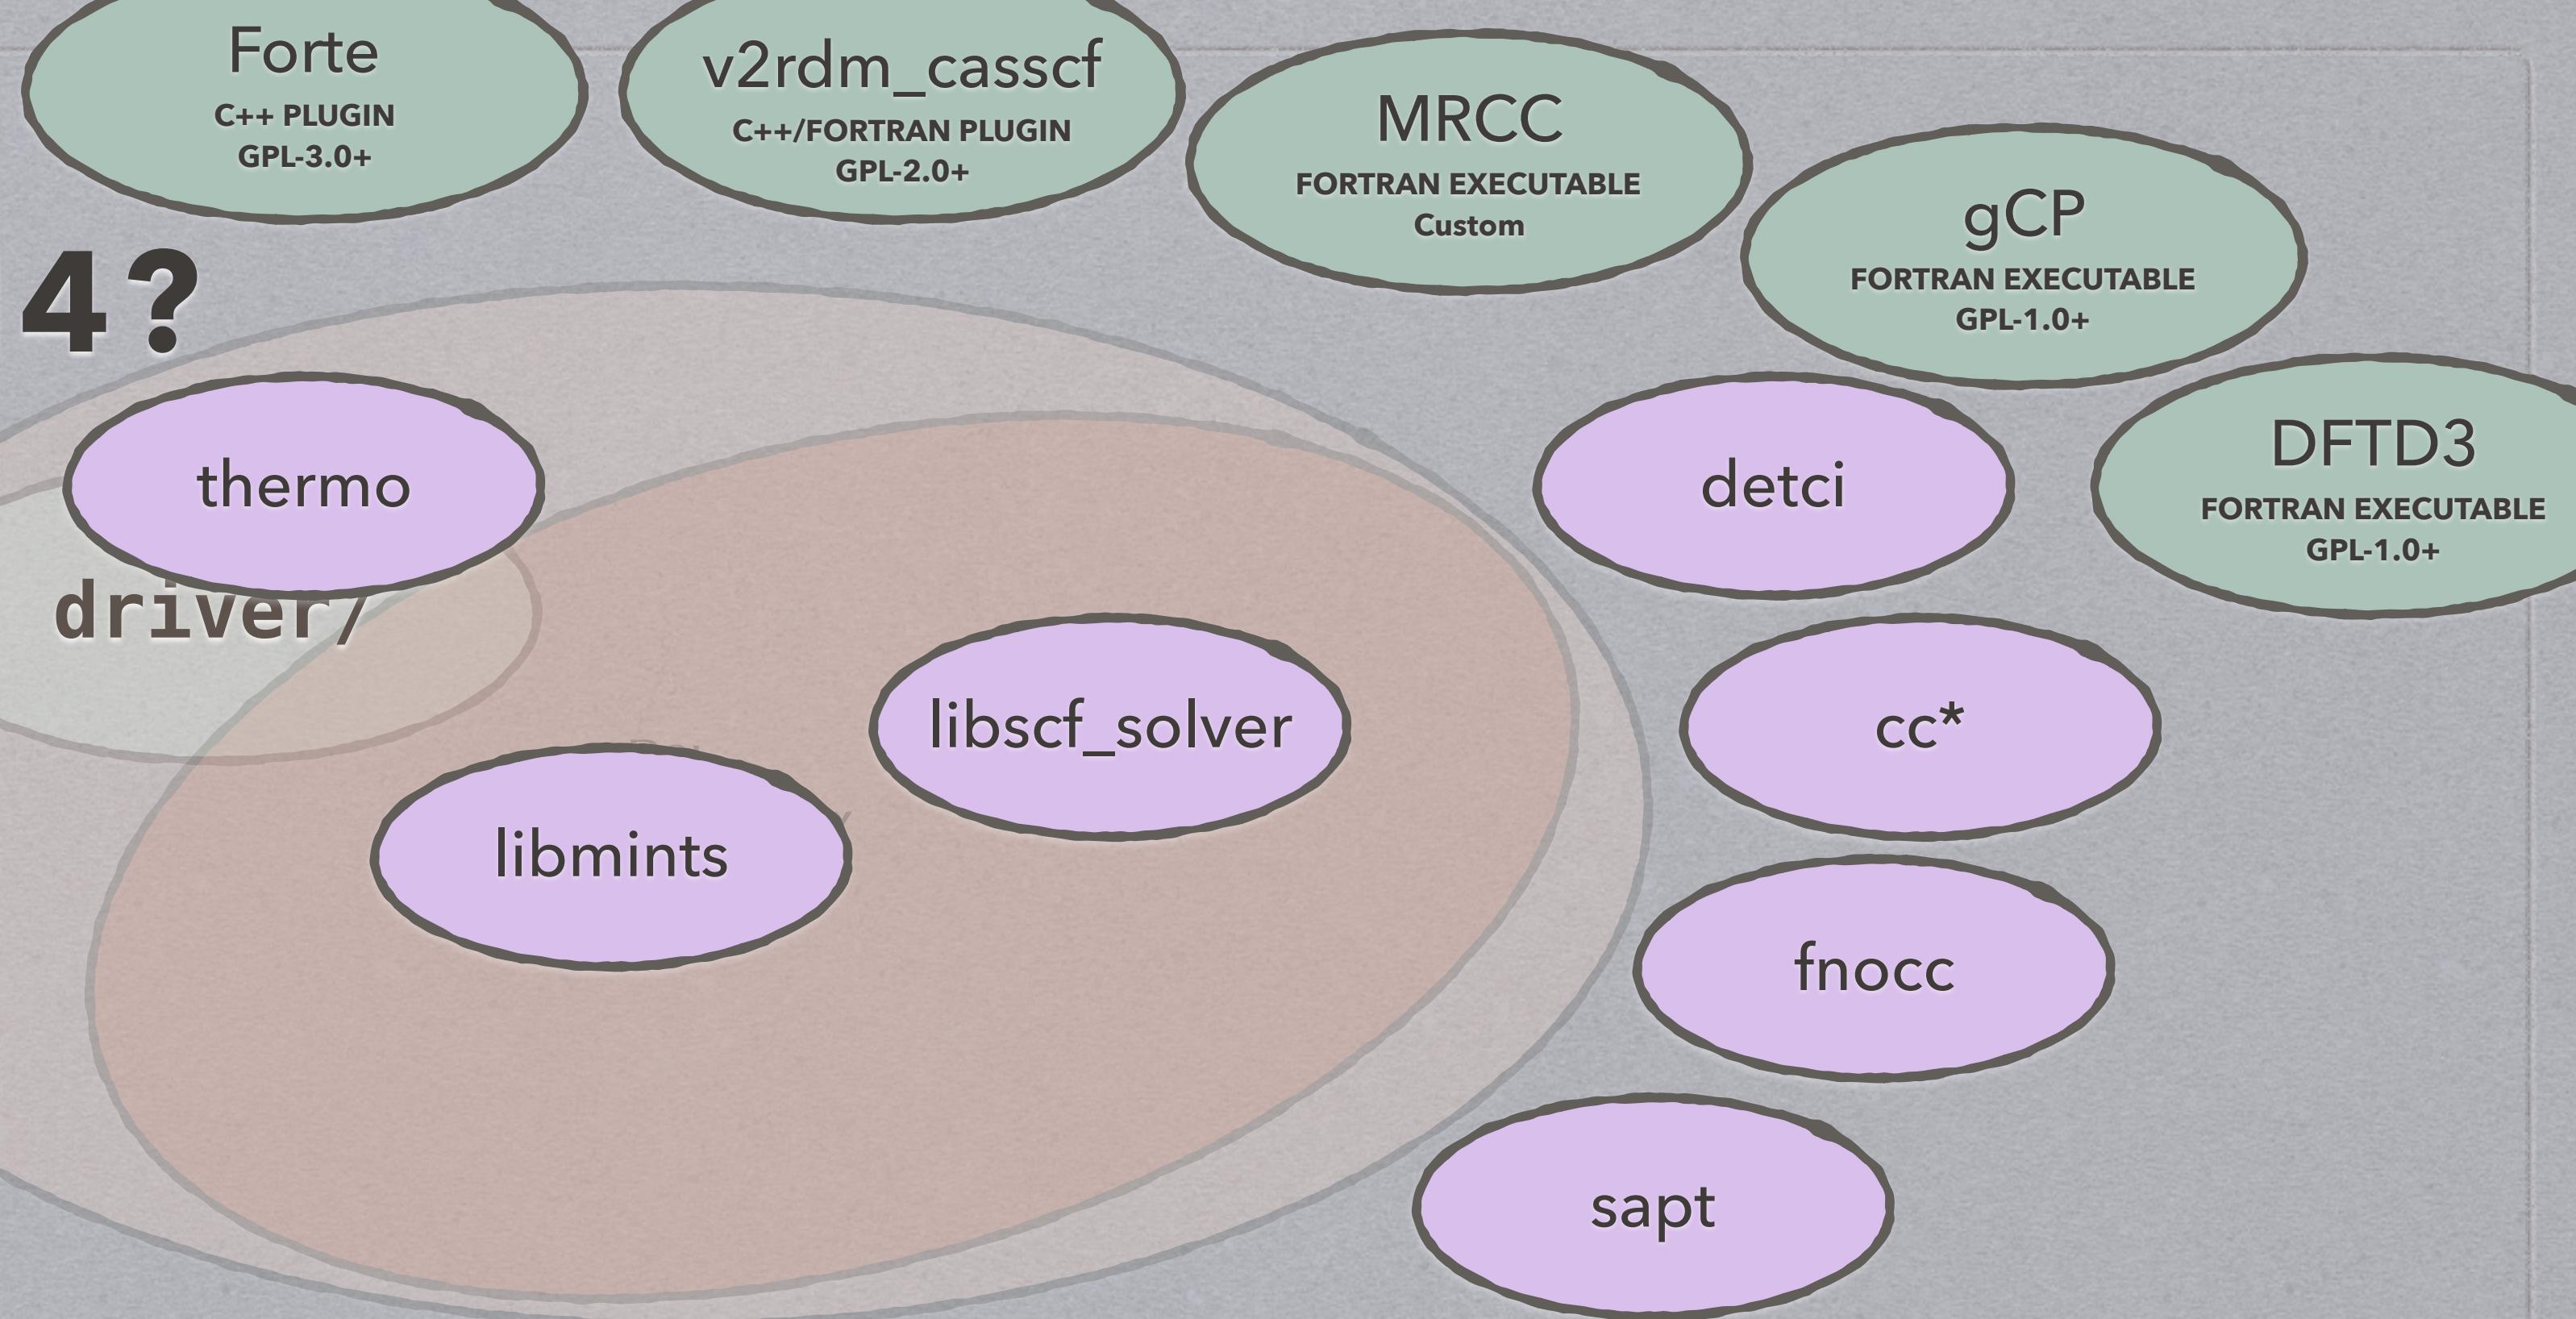


CURRENT STATE CORE.SO == C-SIDE



FUTURE CORE/ PERIPHERAL PSI4?

Monolithic
Modular

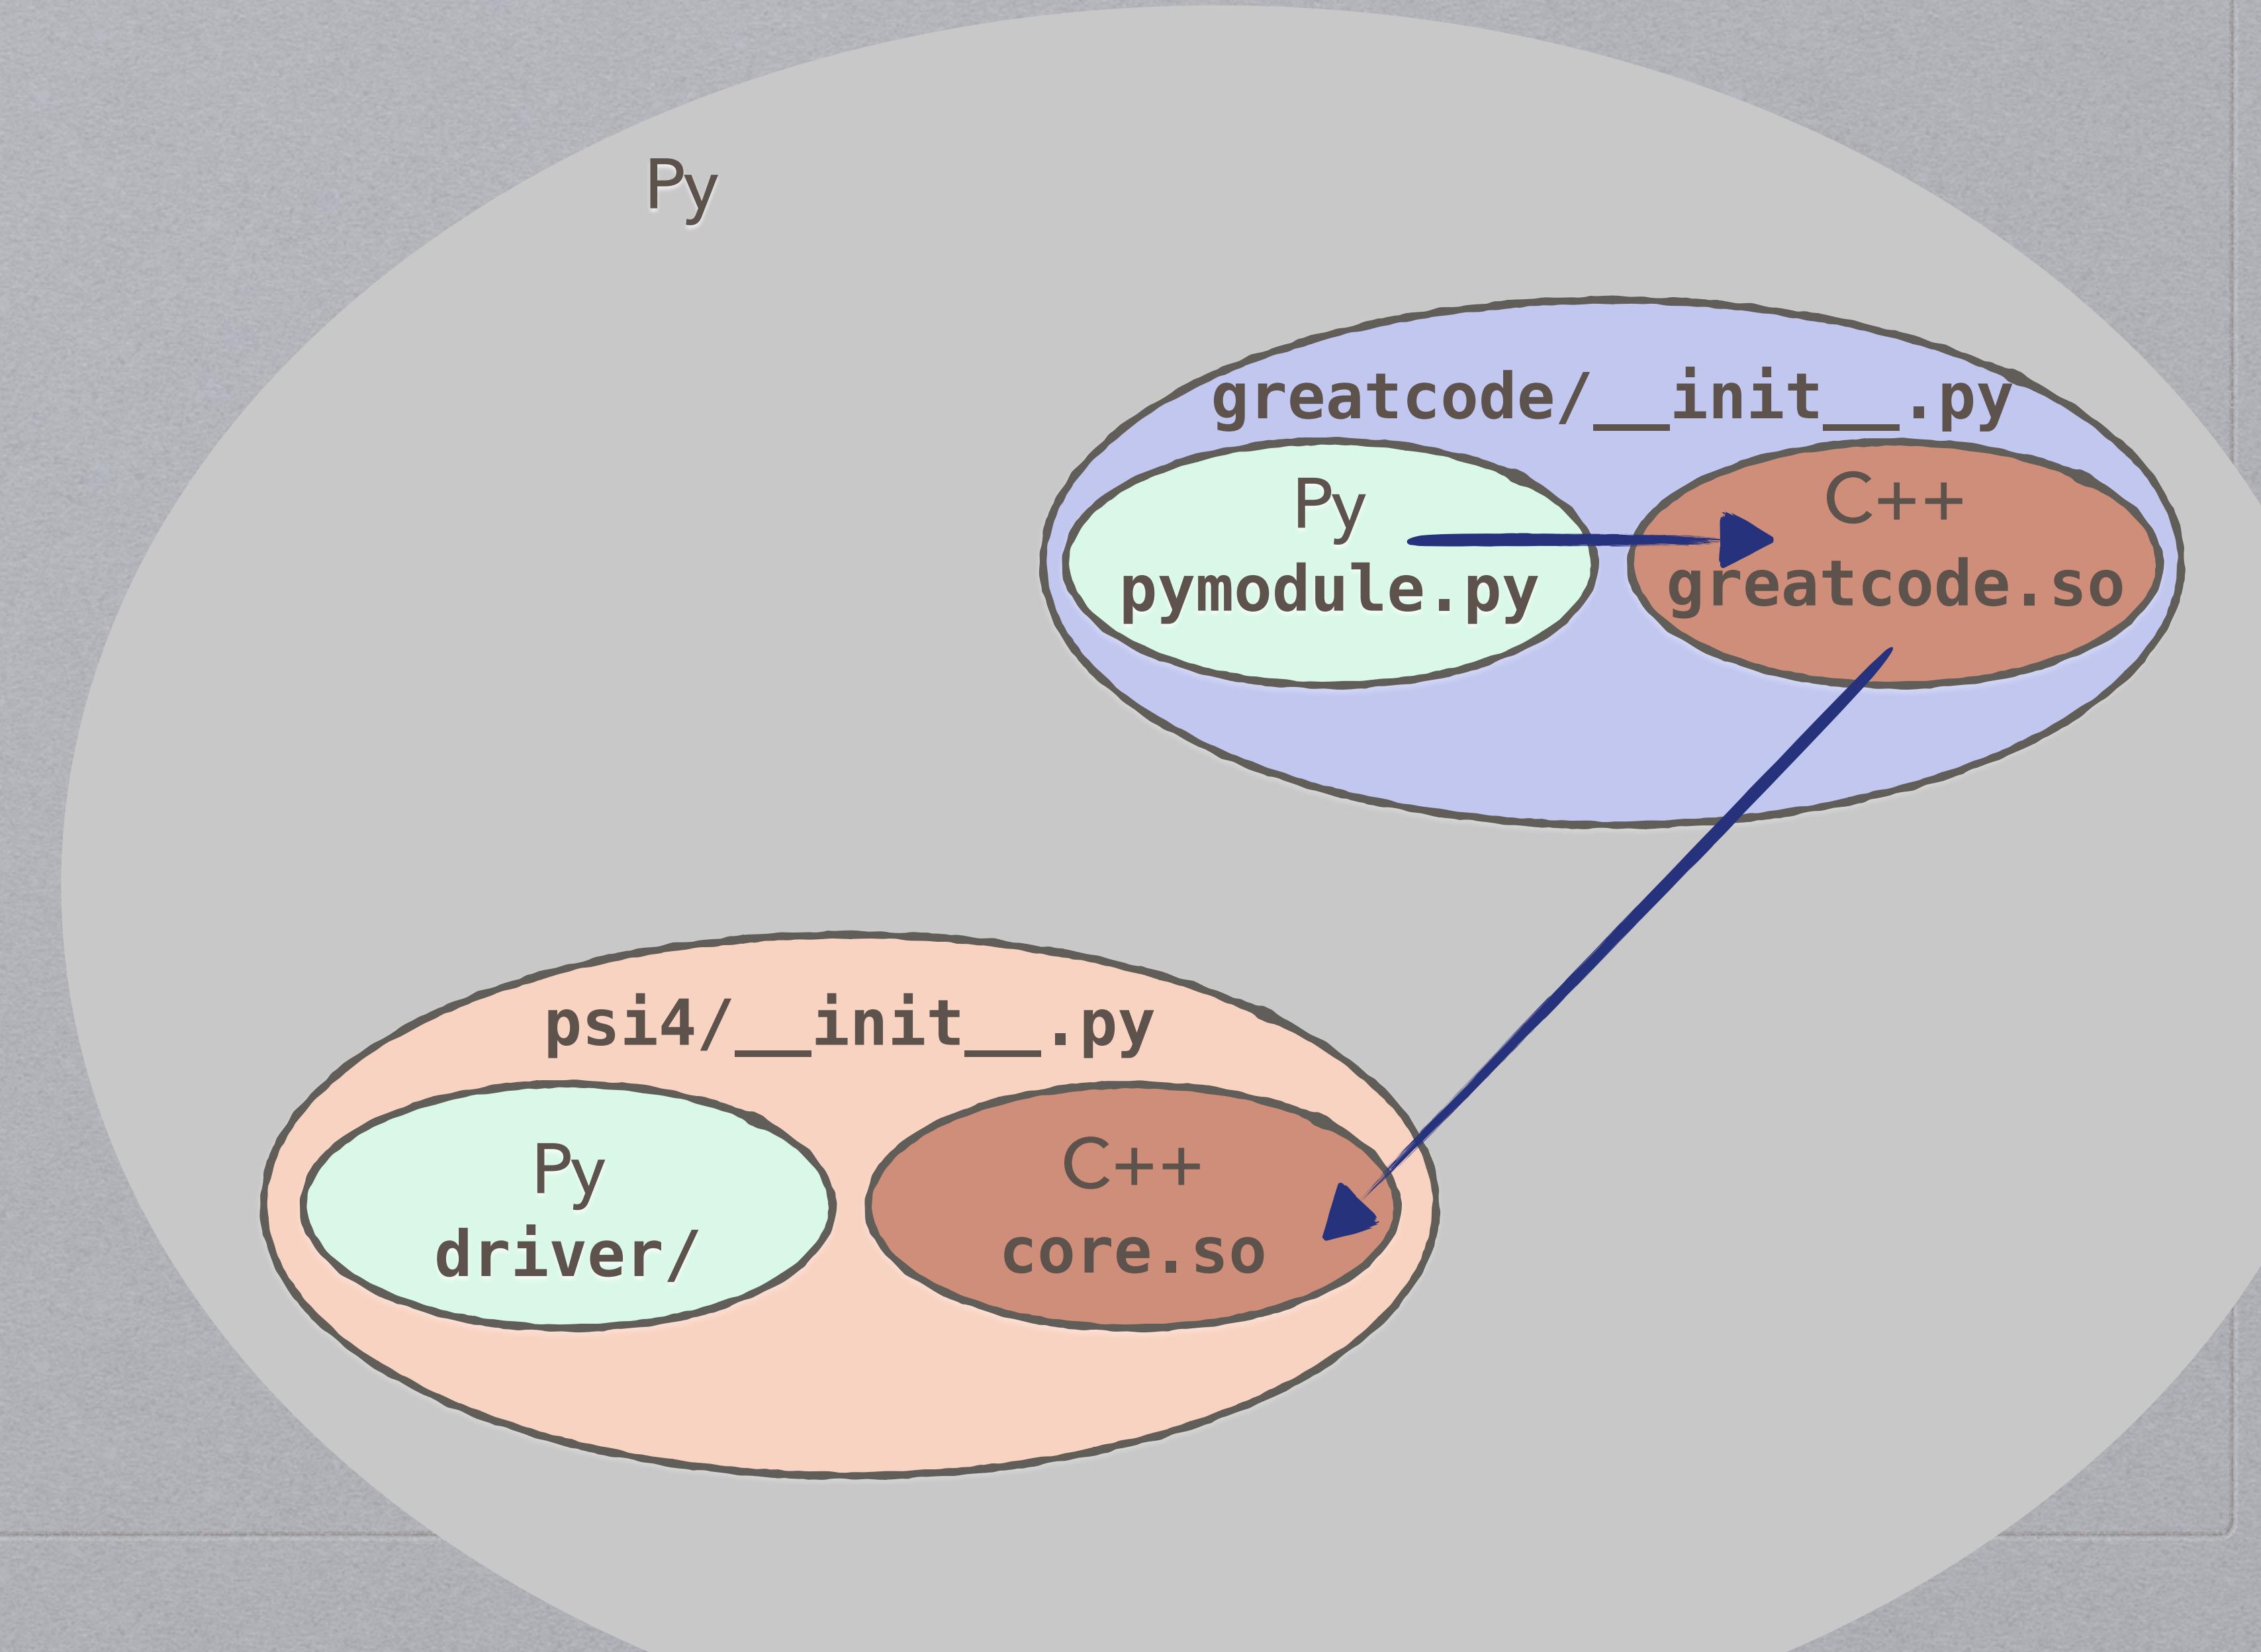


- Less distinction btwn methods *in* Psi and *out of* Psi in plugins, so new faculty can keep clear attribution of their research code.
- New contributors can grasp a finite level of commitment to working with psi4-core, rather than feeling uneasy obligation toward 1M loc.
- Existing projects can stay in-repo (if no longer under active dev) or move out to own repo (if you want dev branches and not compiling Psi a every CI).
- For users, everything gets distributed, so it looks just like before.

CHALLENGES OF STATEHOOD

- Once liberated, Psi4 no longer providing framework for
 - build system, including dependencies
 - oddly, no one wants to manage a CMake system, much less a superbuild
 - documentation build
 - theme, linking, autodoc all a pain
 - all psi4's docs tools now available outside psi4 repo
 - versioning
 - if packages depend on you, I highly recommend the self-versioning scheme like psi4 has, so every commit is traceable
 - if a commit that's valuable downstream isn't released in short working order (couple weeks), then it needs to be otherwise addressable
 - options
 - uggh, but handled by current plugins (not nice for pure-Py to forcibly have to compile just to get options)
 - testing
 - a few inputs with Makefiles are better than nothing, but far better is a testing scheme that works for *installed* software. that is, pytest, not ctest. we can maintain current psithon test suite and run it through pytest. Just need to agree that pytest is authoritative.
 - downstream deps may not even build us, so no ctest and if they use a stable psi, won't even notice broken wrt us
- All of this calls for a more comprehensive plugin template, namely, a cookie cutter

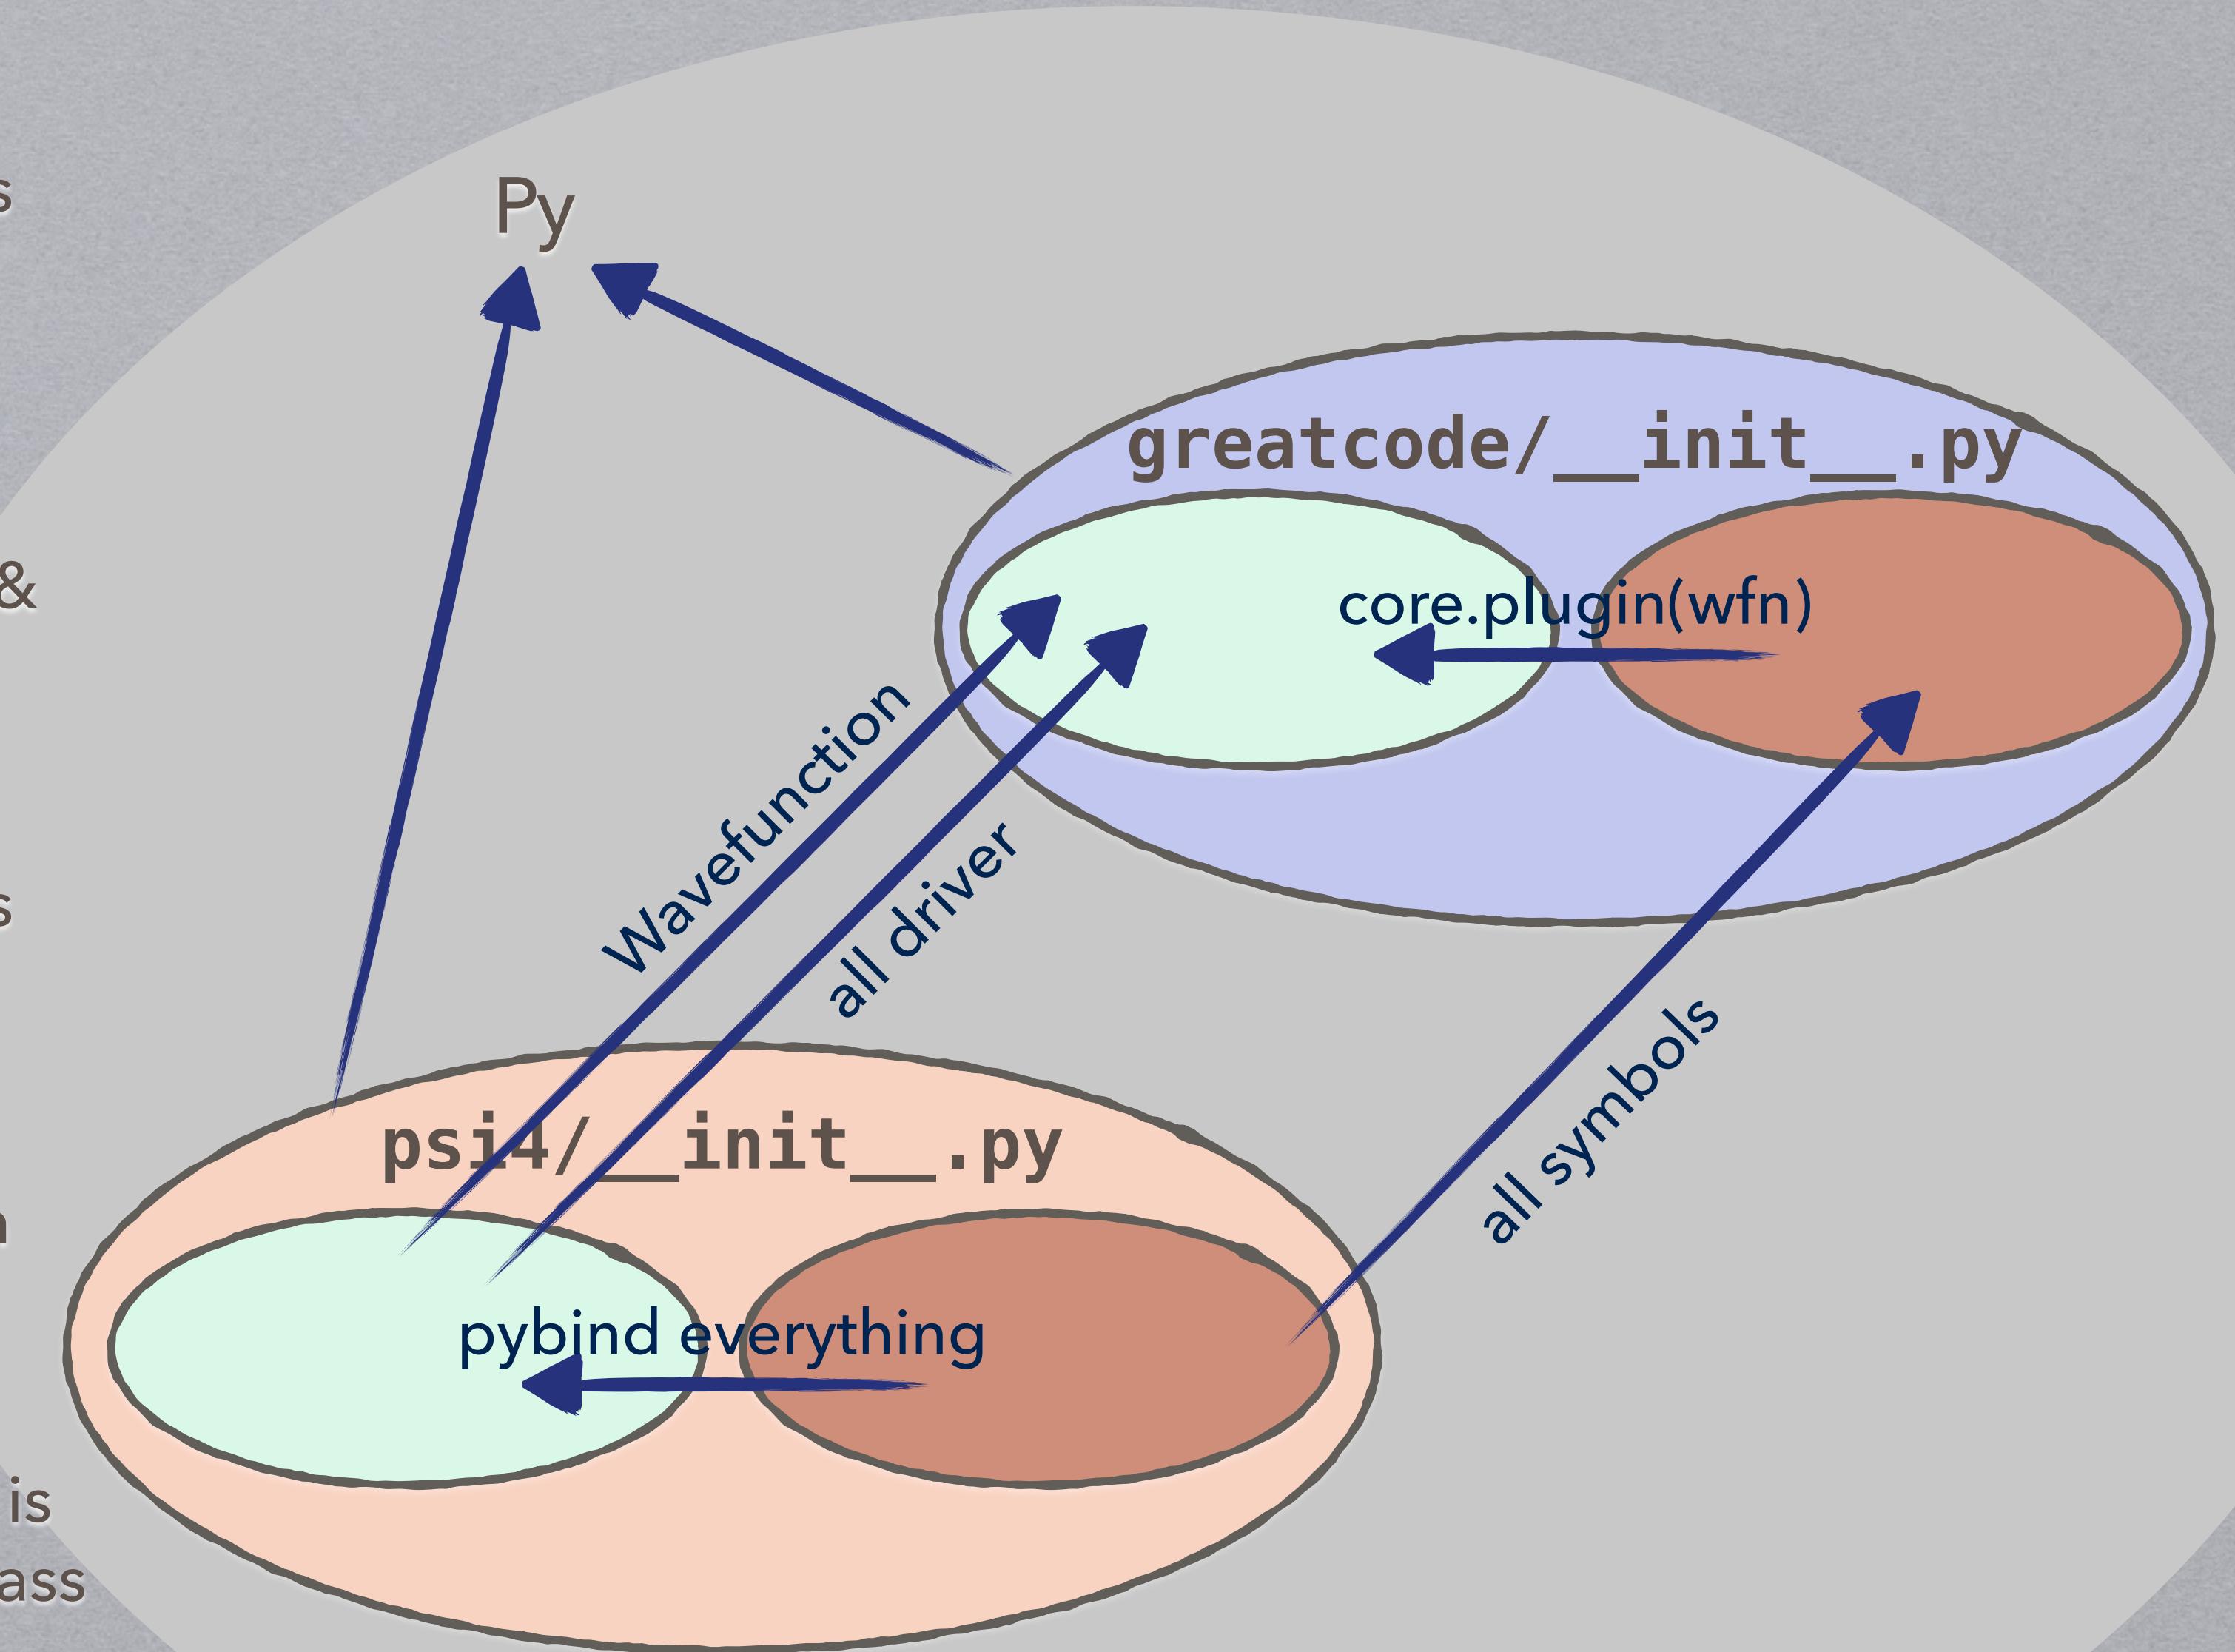
PLUGIN FASHIONABILITY



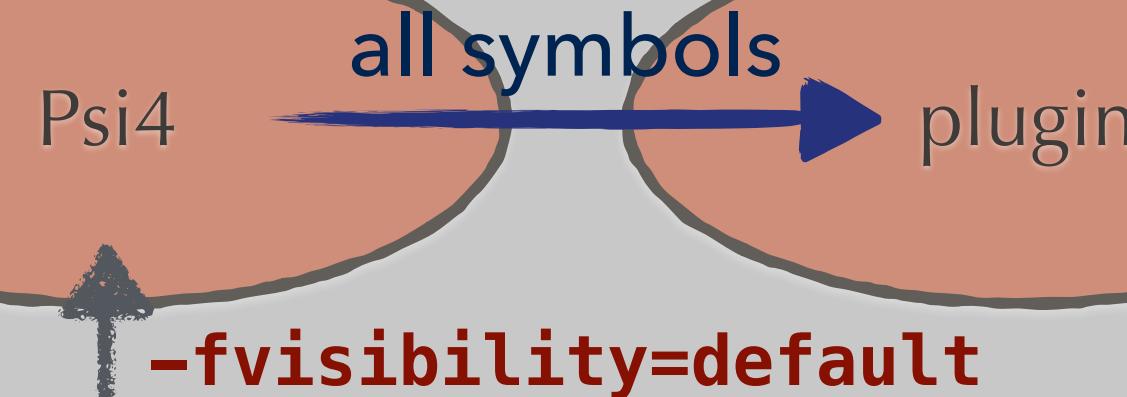
PLUGIN FASHIONABILITY

- a Psi method
 - C-side has access to all Psi C-libraries
 - Py-side has access to all Psi Py-/C-libraries
 - thus can shift Py/C boundary transparently to user as prototyping & efficiency demand
- a plugin method
 - C-side has access to all Psi C-libraries
 - Py-side has access to all Psi Py-/C-libraries
 - very limited Py/C interface w/i plugin
 - plugins are excellent workspace for developing a Psi C-library (not that the compile time-penalty for in-tree lib dev is high anymore) but decidedly second-class for method prototyping or sustaining

arrow head can access arrow tail

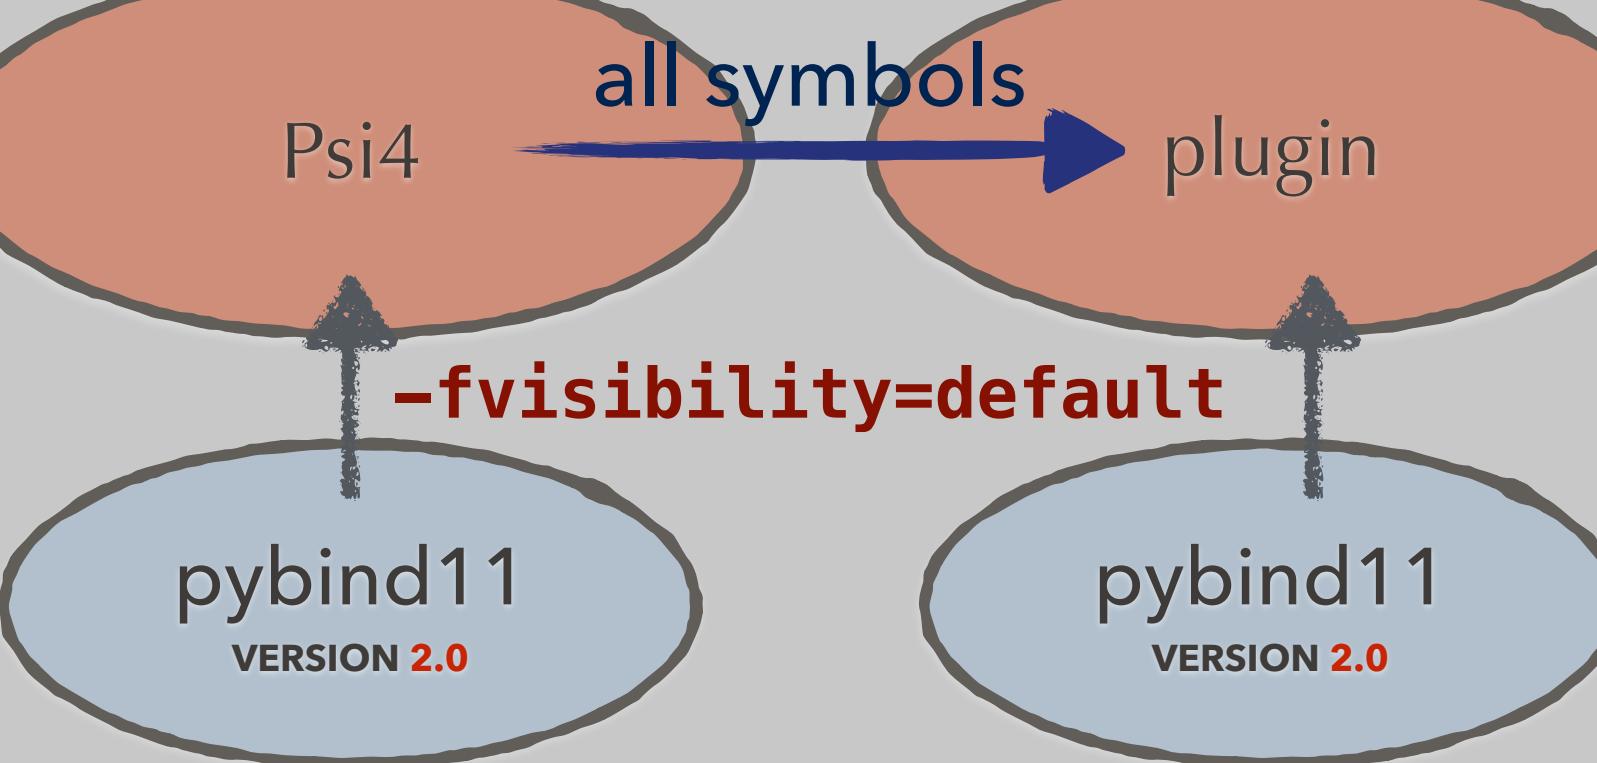


FVISIBILITY CRISIS

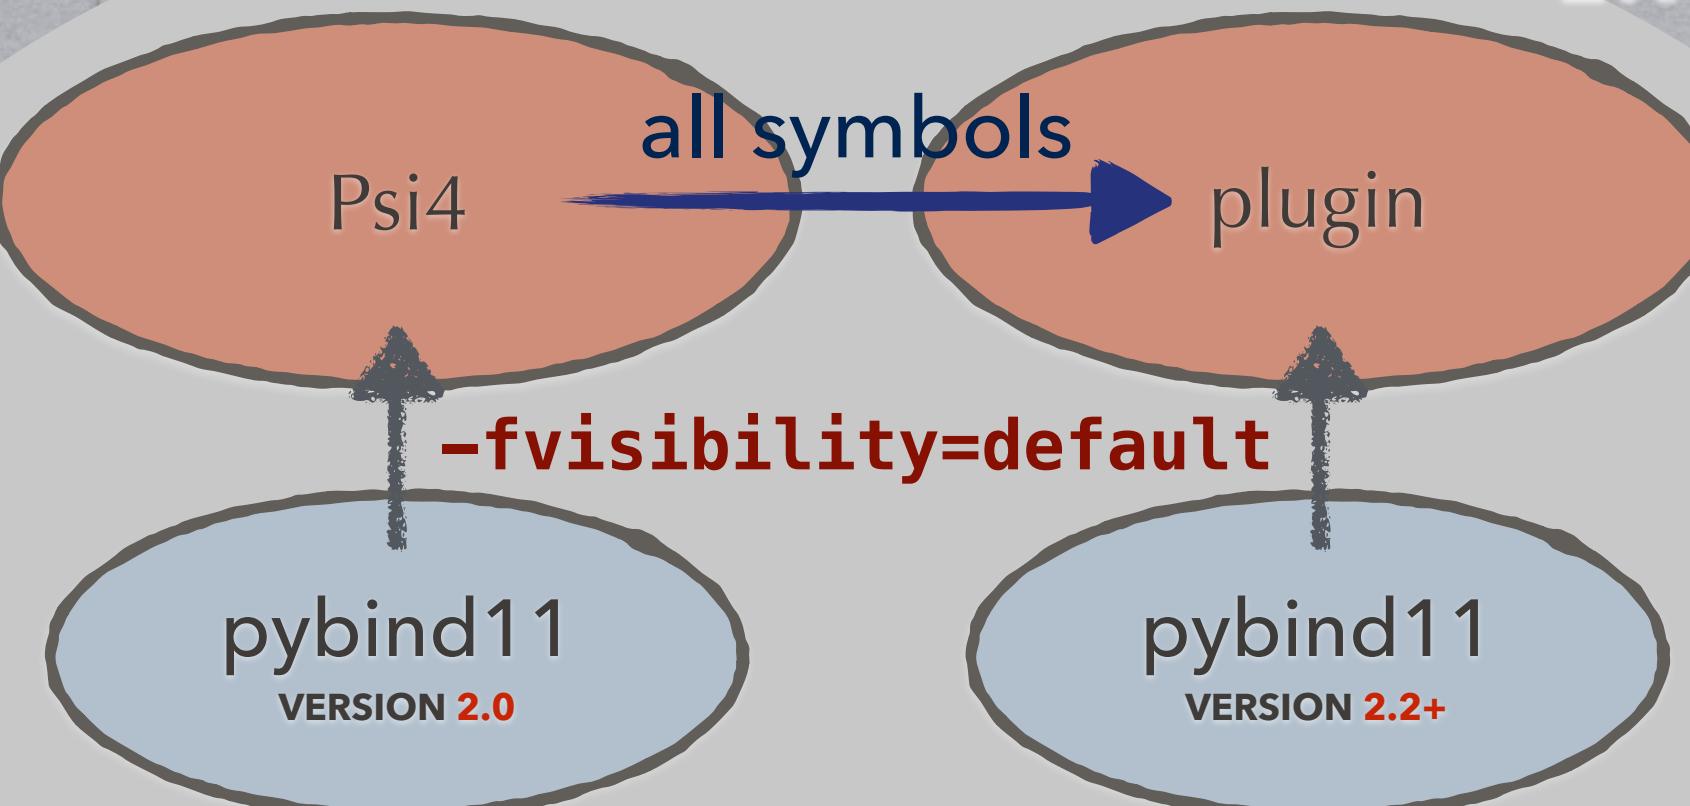


VERSION 2.0

NOW



ENFORCE YEAR-OLD VERSION

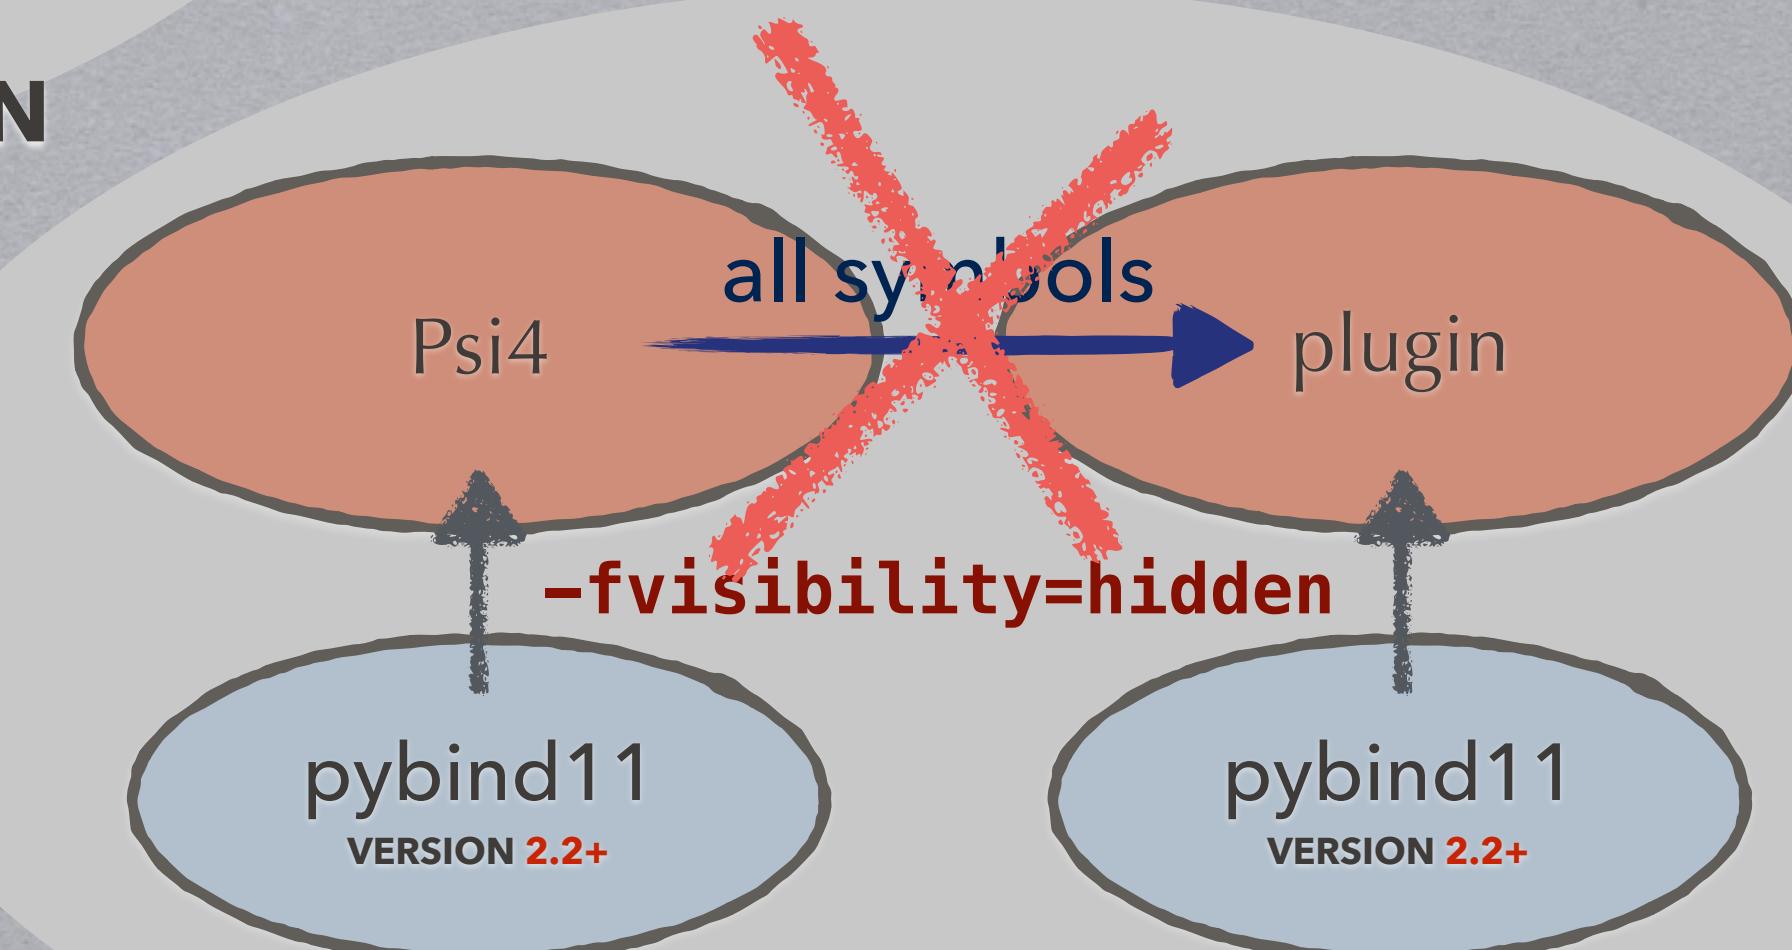


MIGHT WORK; ENFORCE PSI SOLE 2.0

- Presently can't build Psi4/Pybind 2.2+ & dependencies together. directives can't compete with Linux header-finding tech. To control pybind version, must be out of sight; even an install next to a dependency can interfere.

TACKLE SYMBOL CURATION

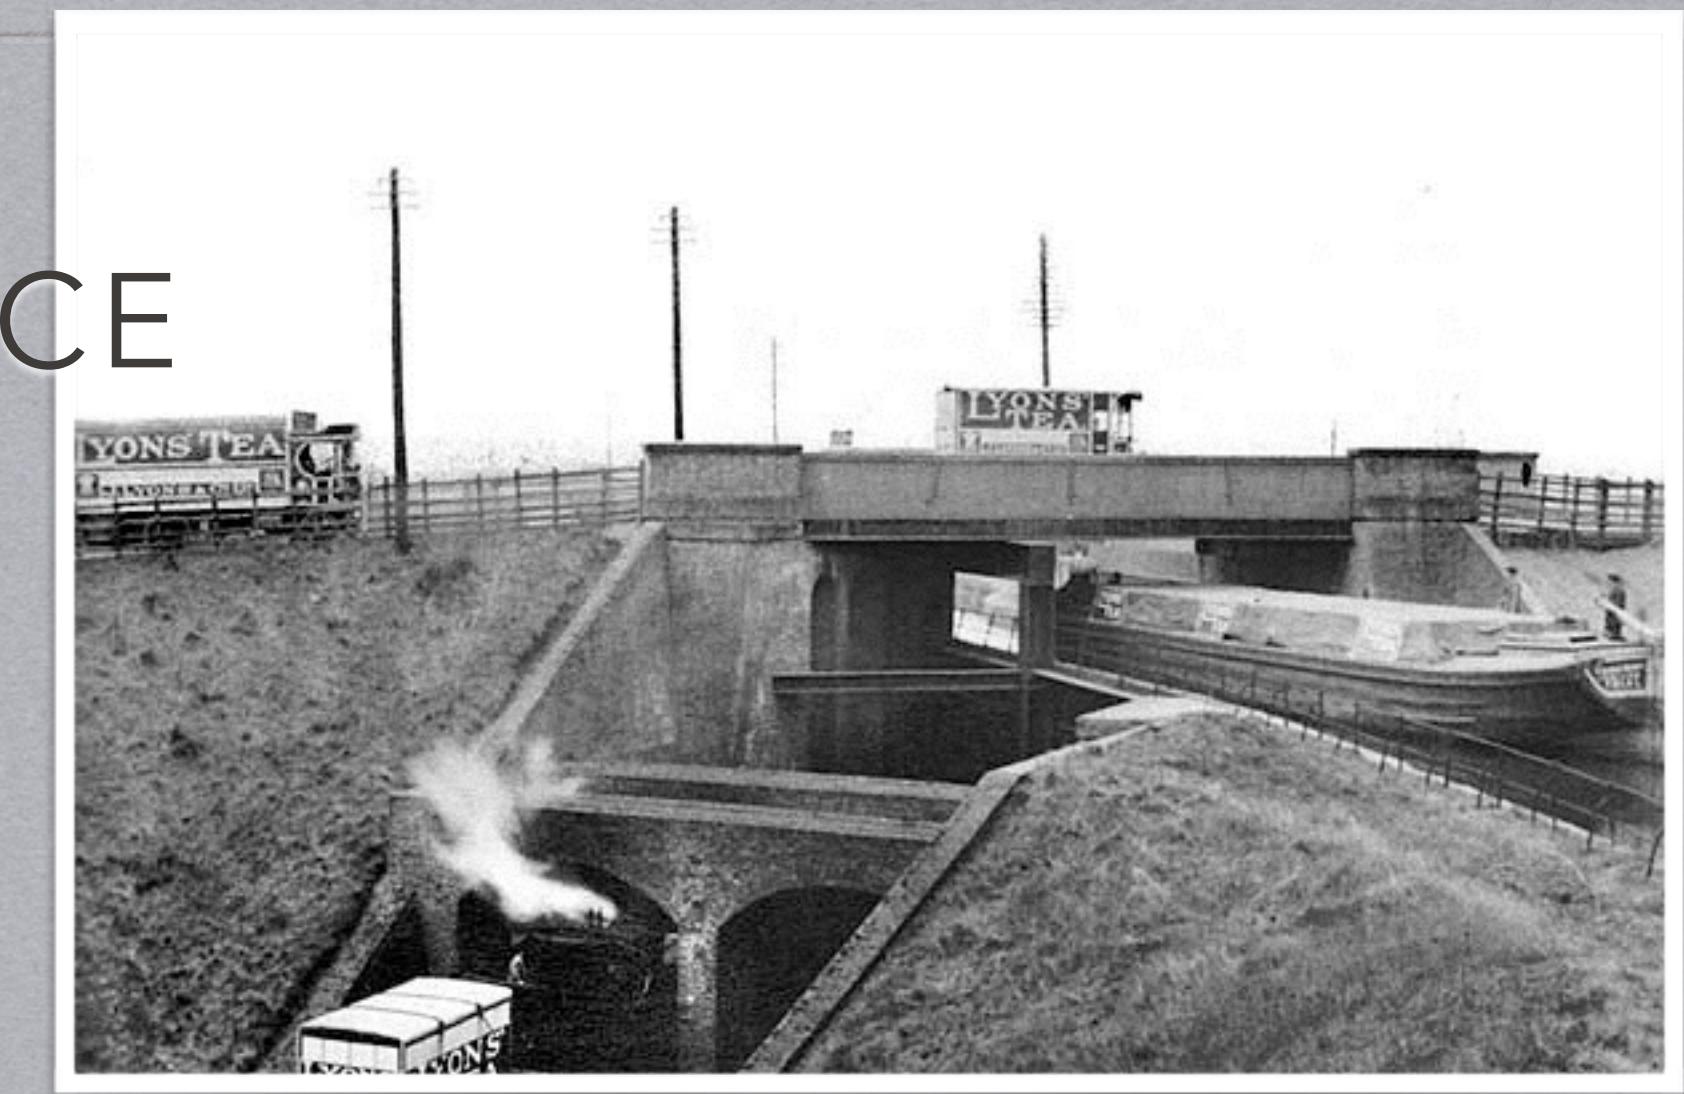
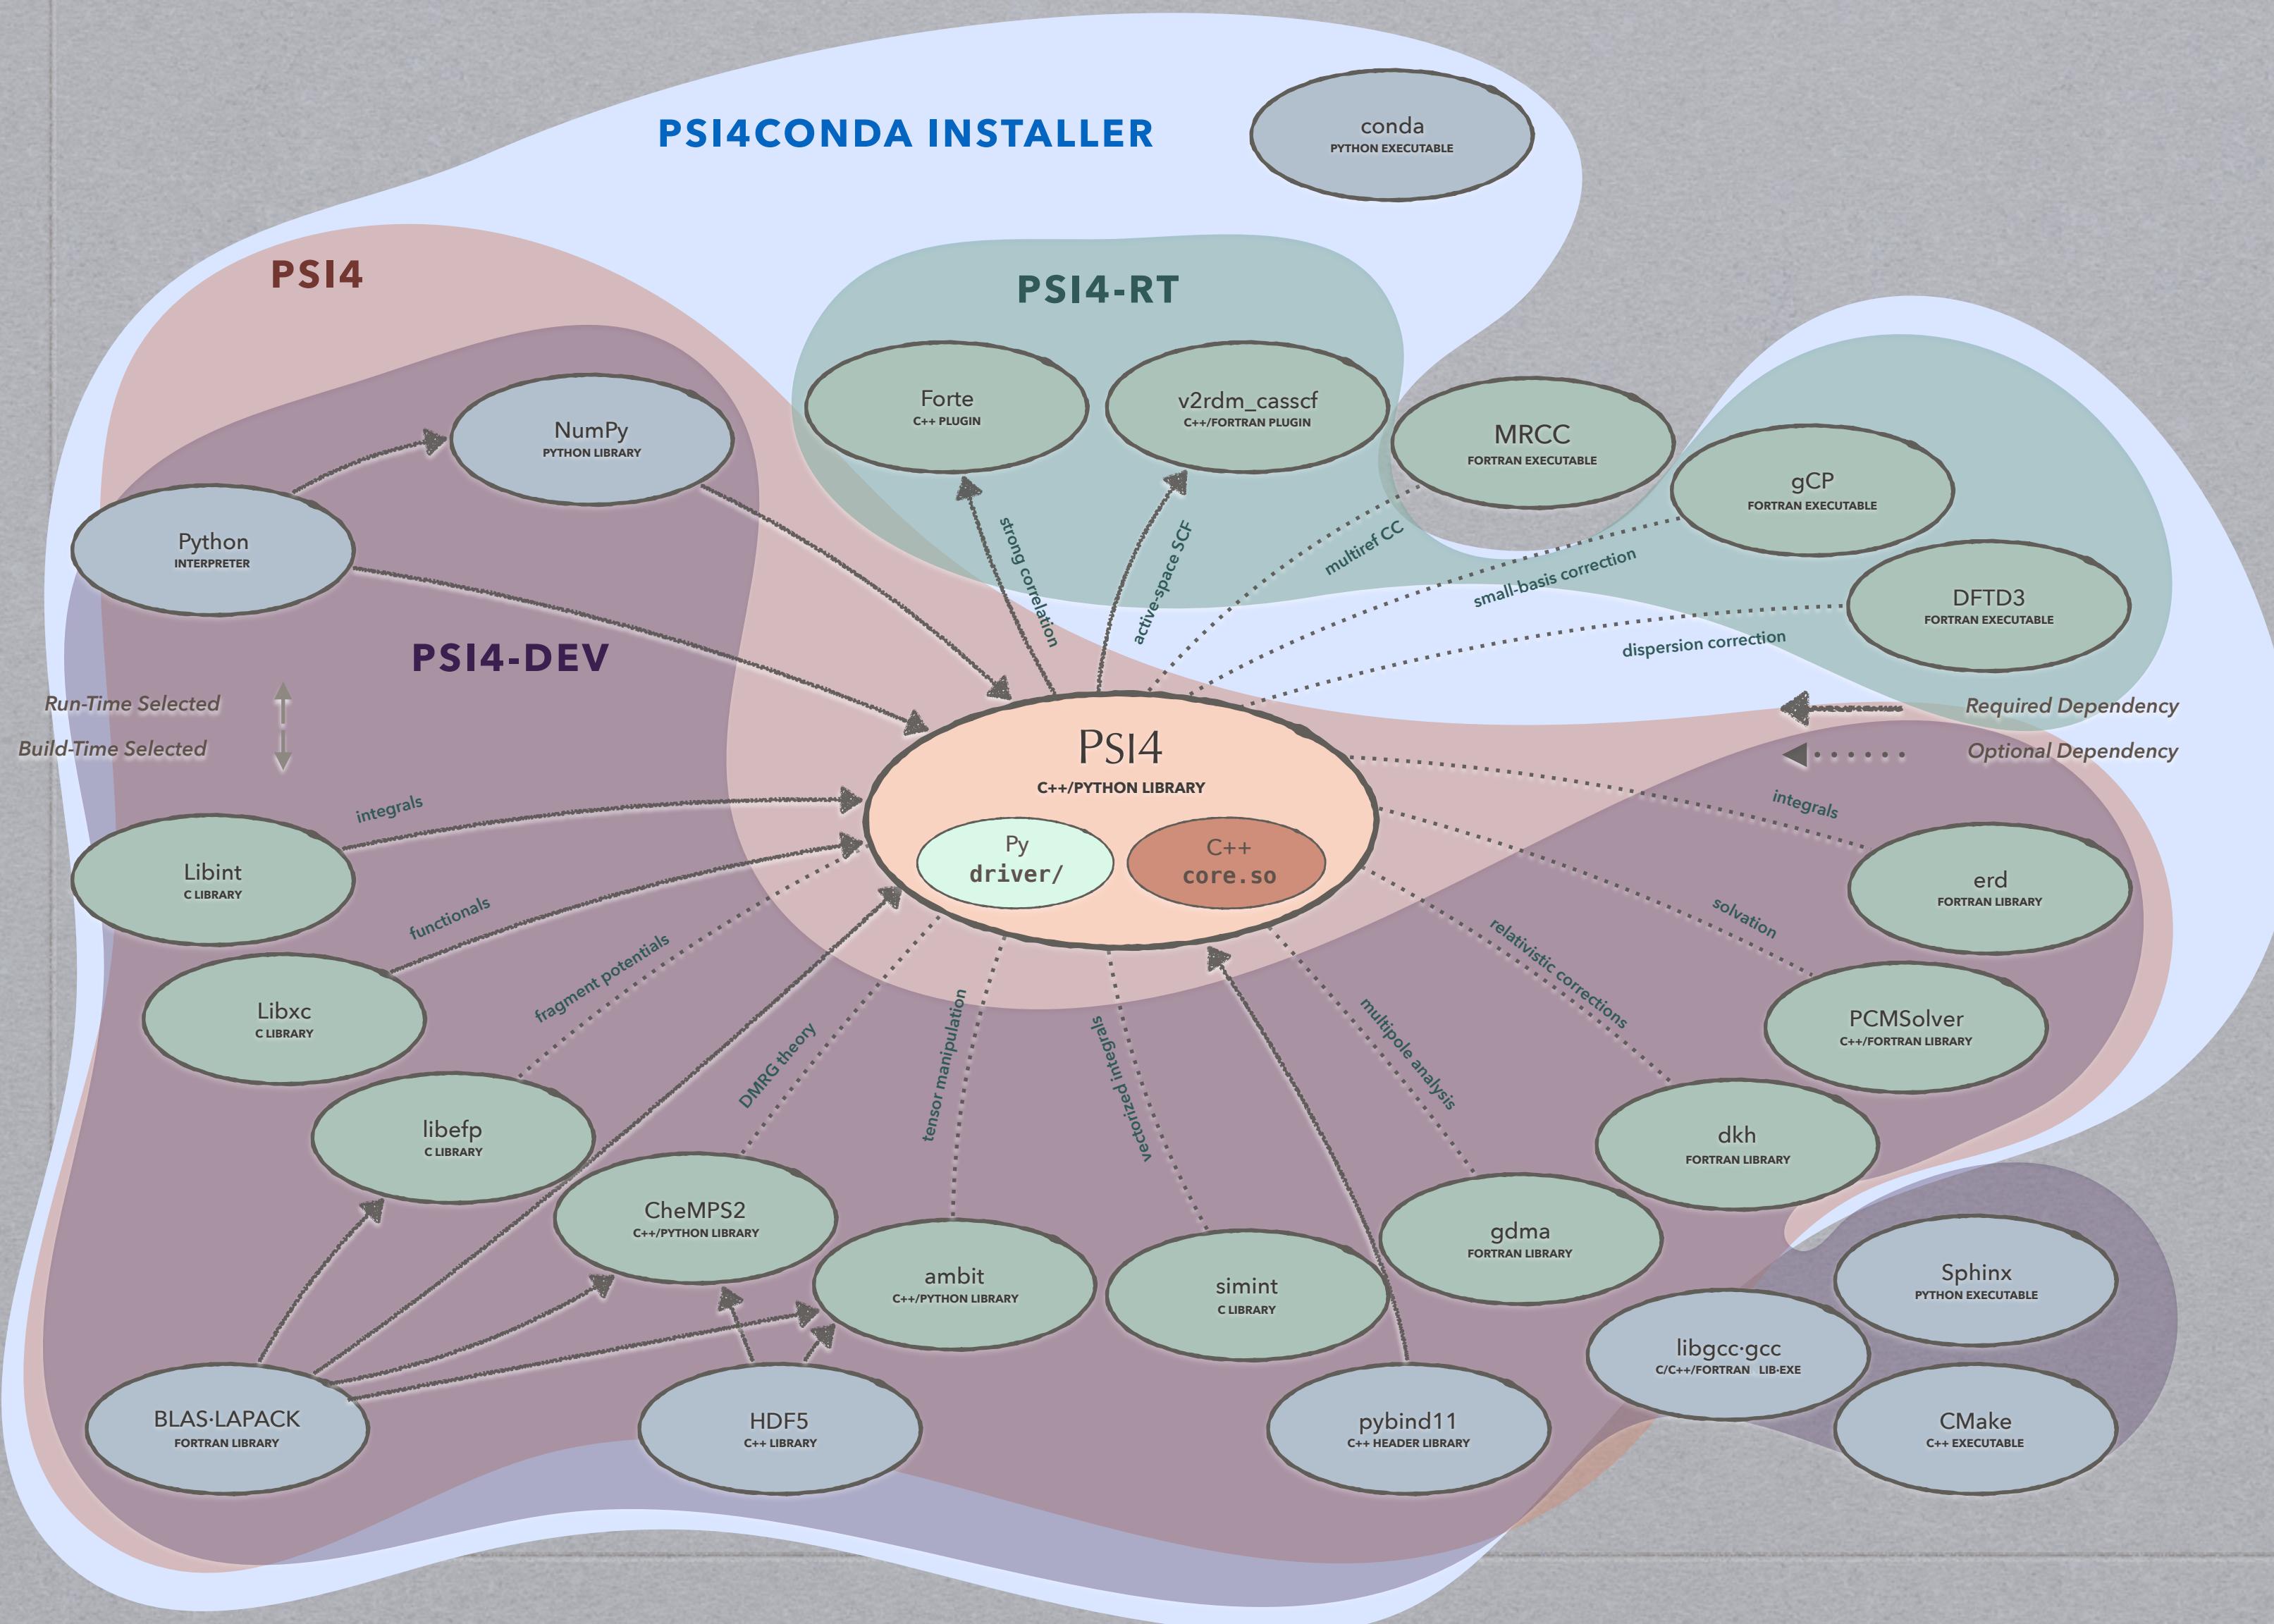
- so can pybind more deps
 - forte
 - pylibefp
 - only psi4-core?



CAN'T ACCESS SYMBOLS

CMAKE & CONDA FOR ALL

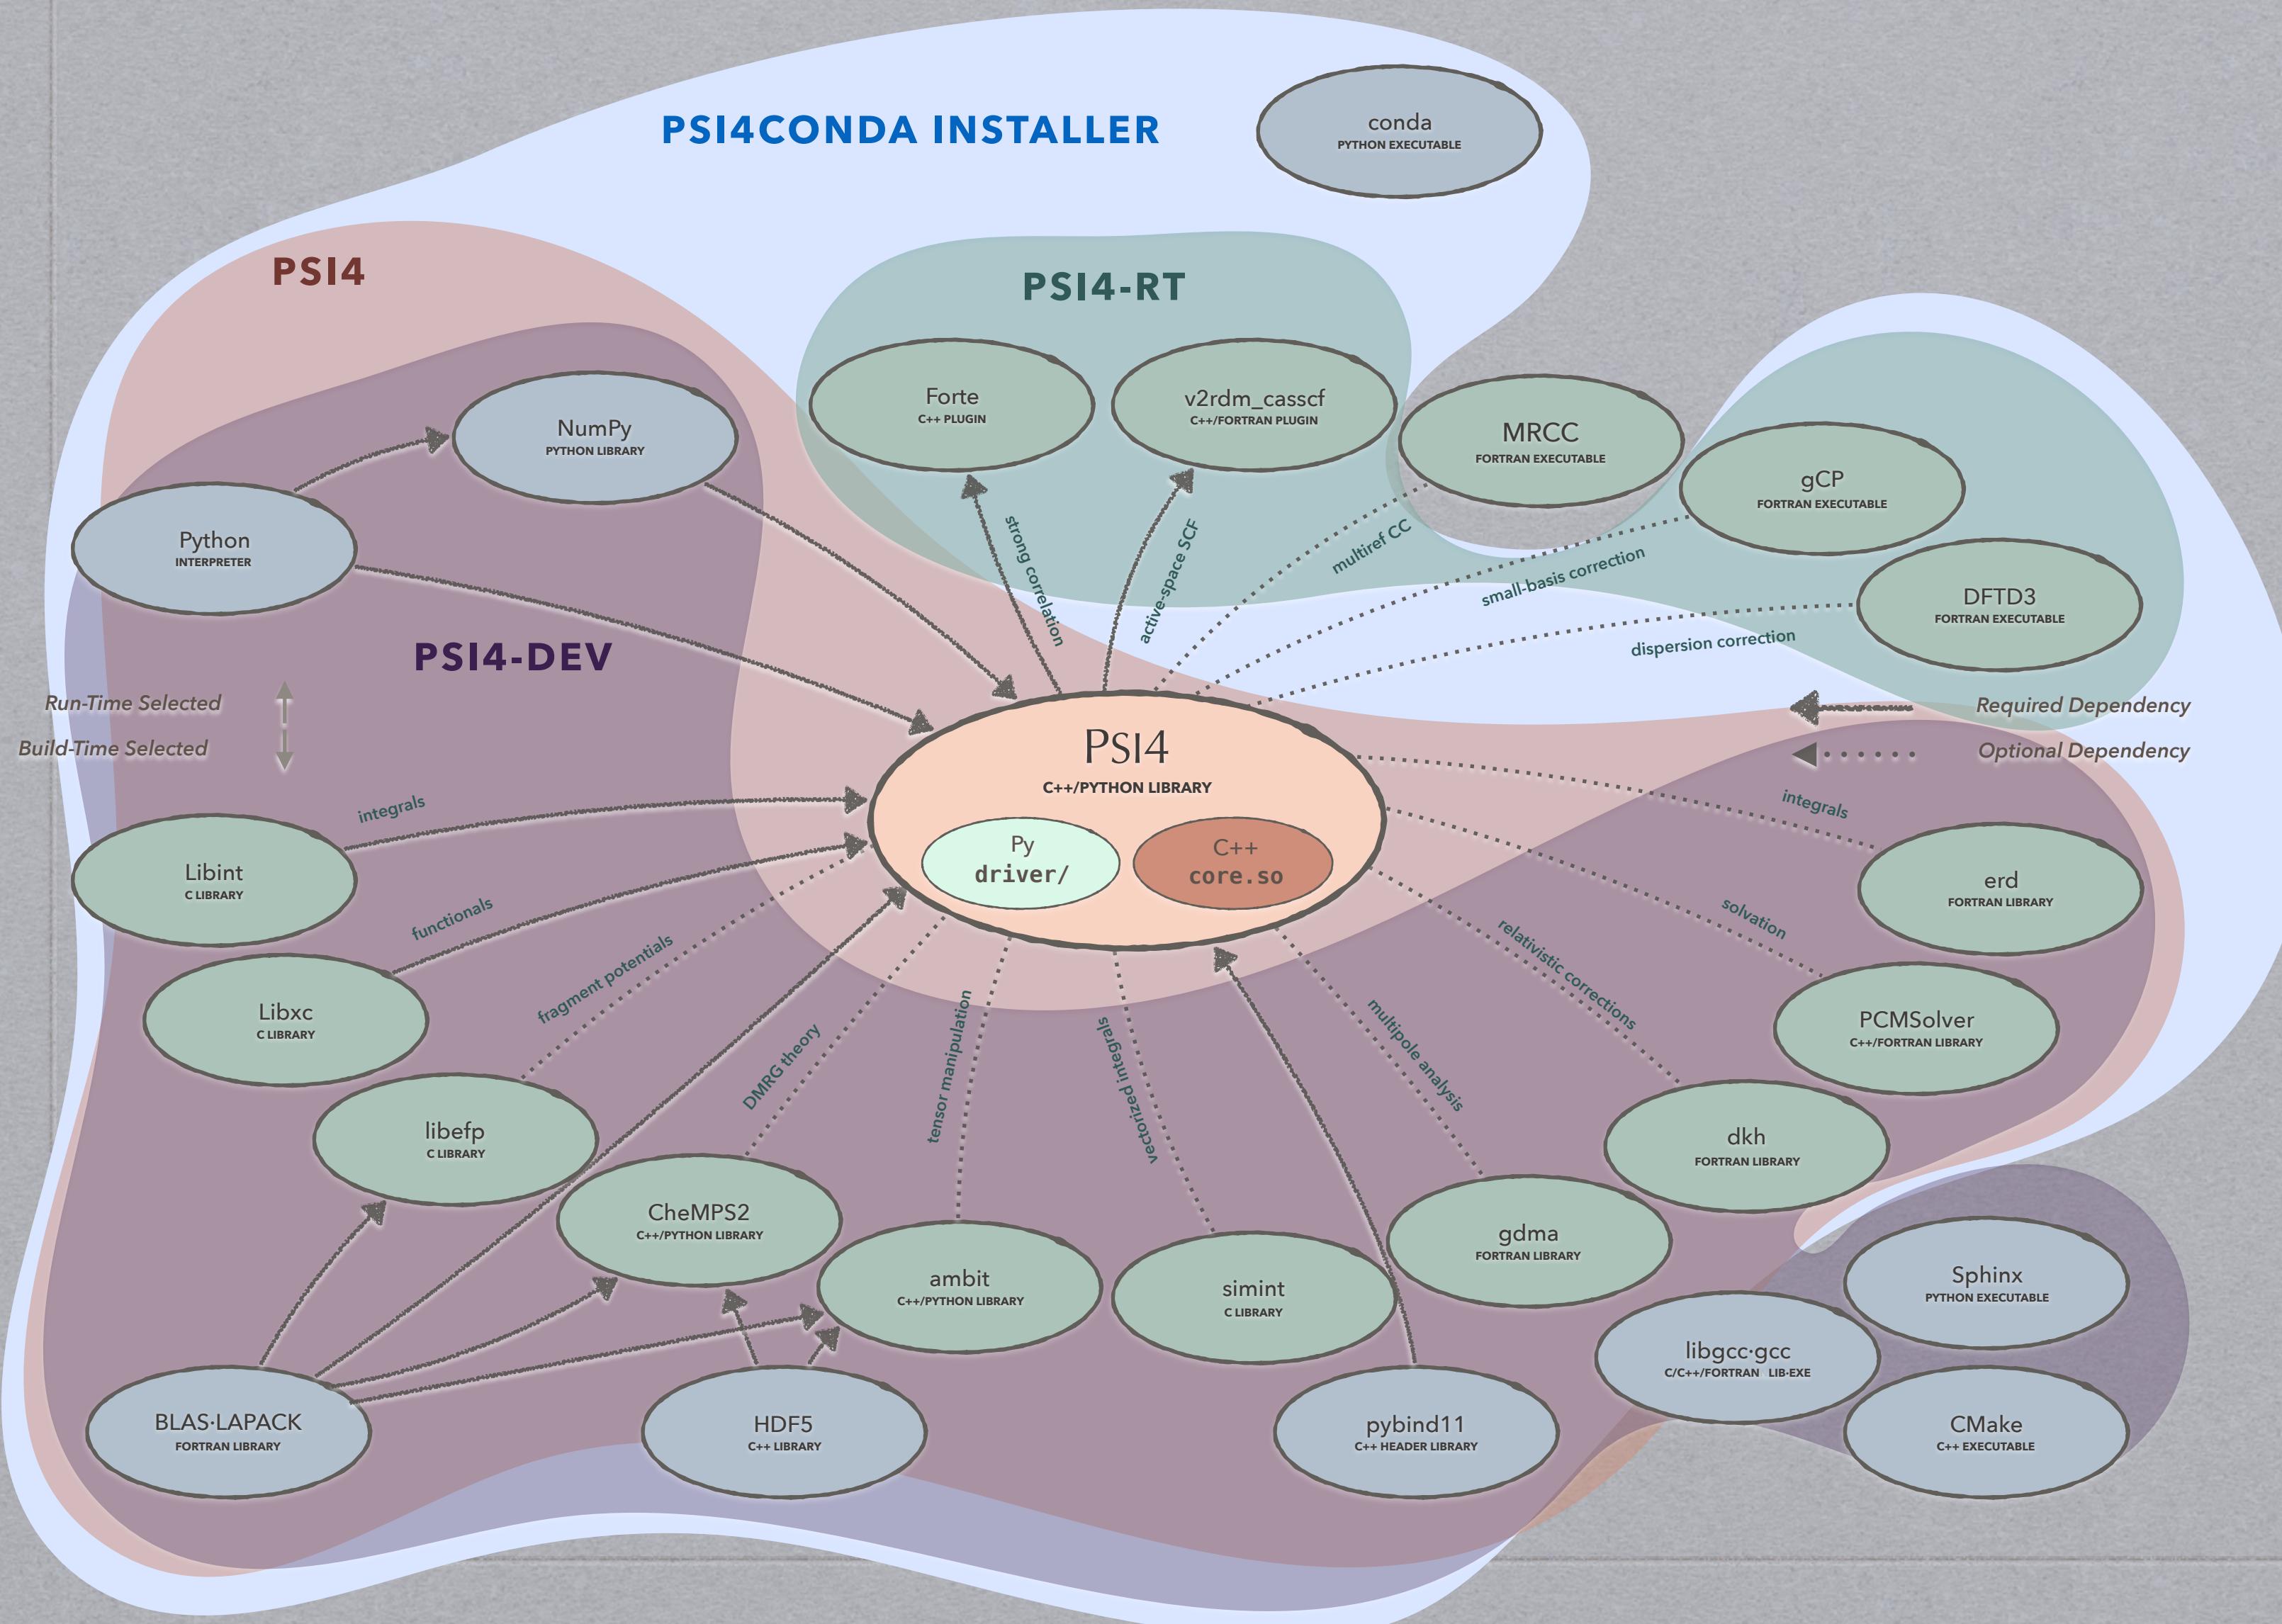
TRICKY OPTIMIZING FOR ALL USERS AT ONCE



**THREE BRIDGES, LONDON
(ROAD OVER CANAL OVER RAIL)**

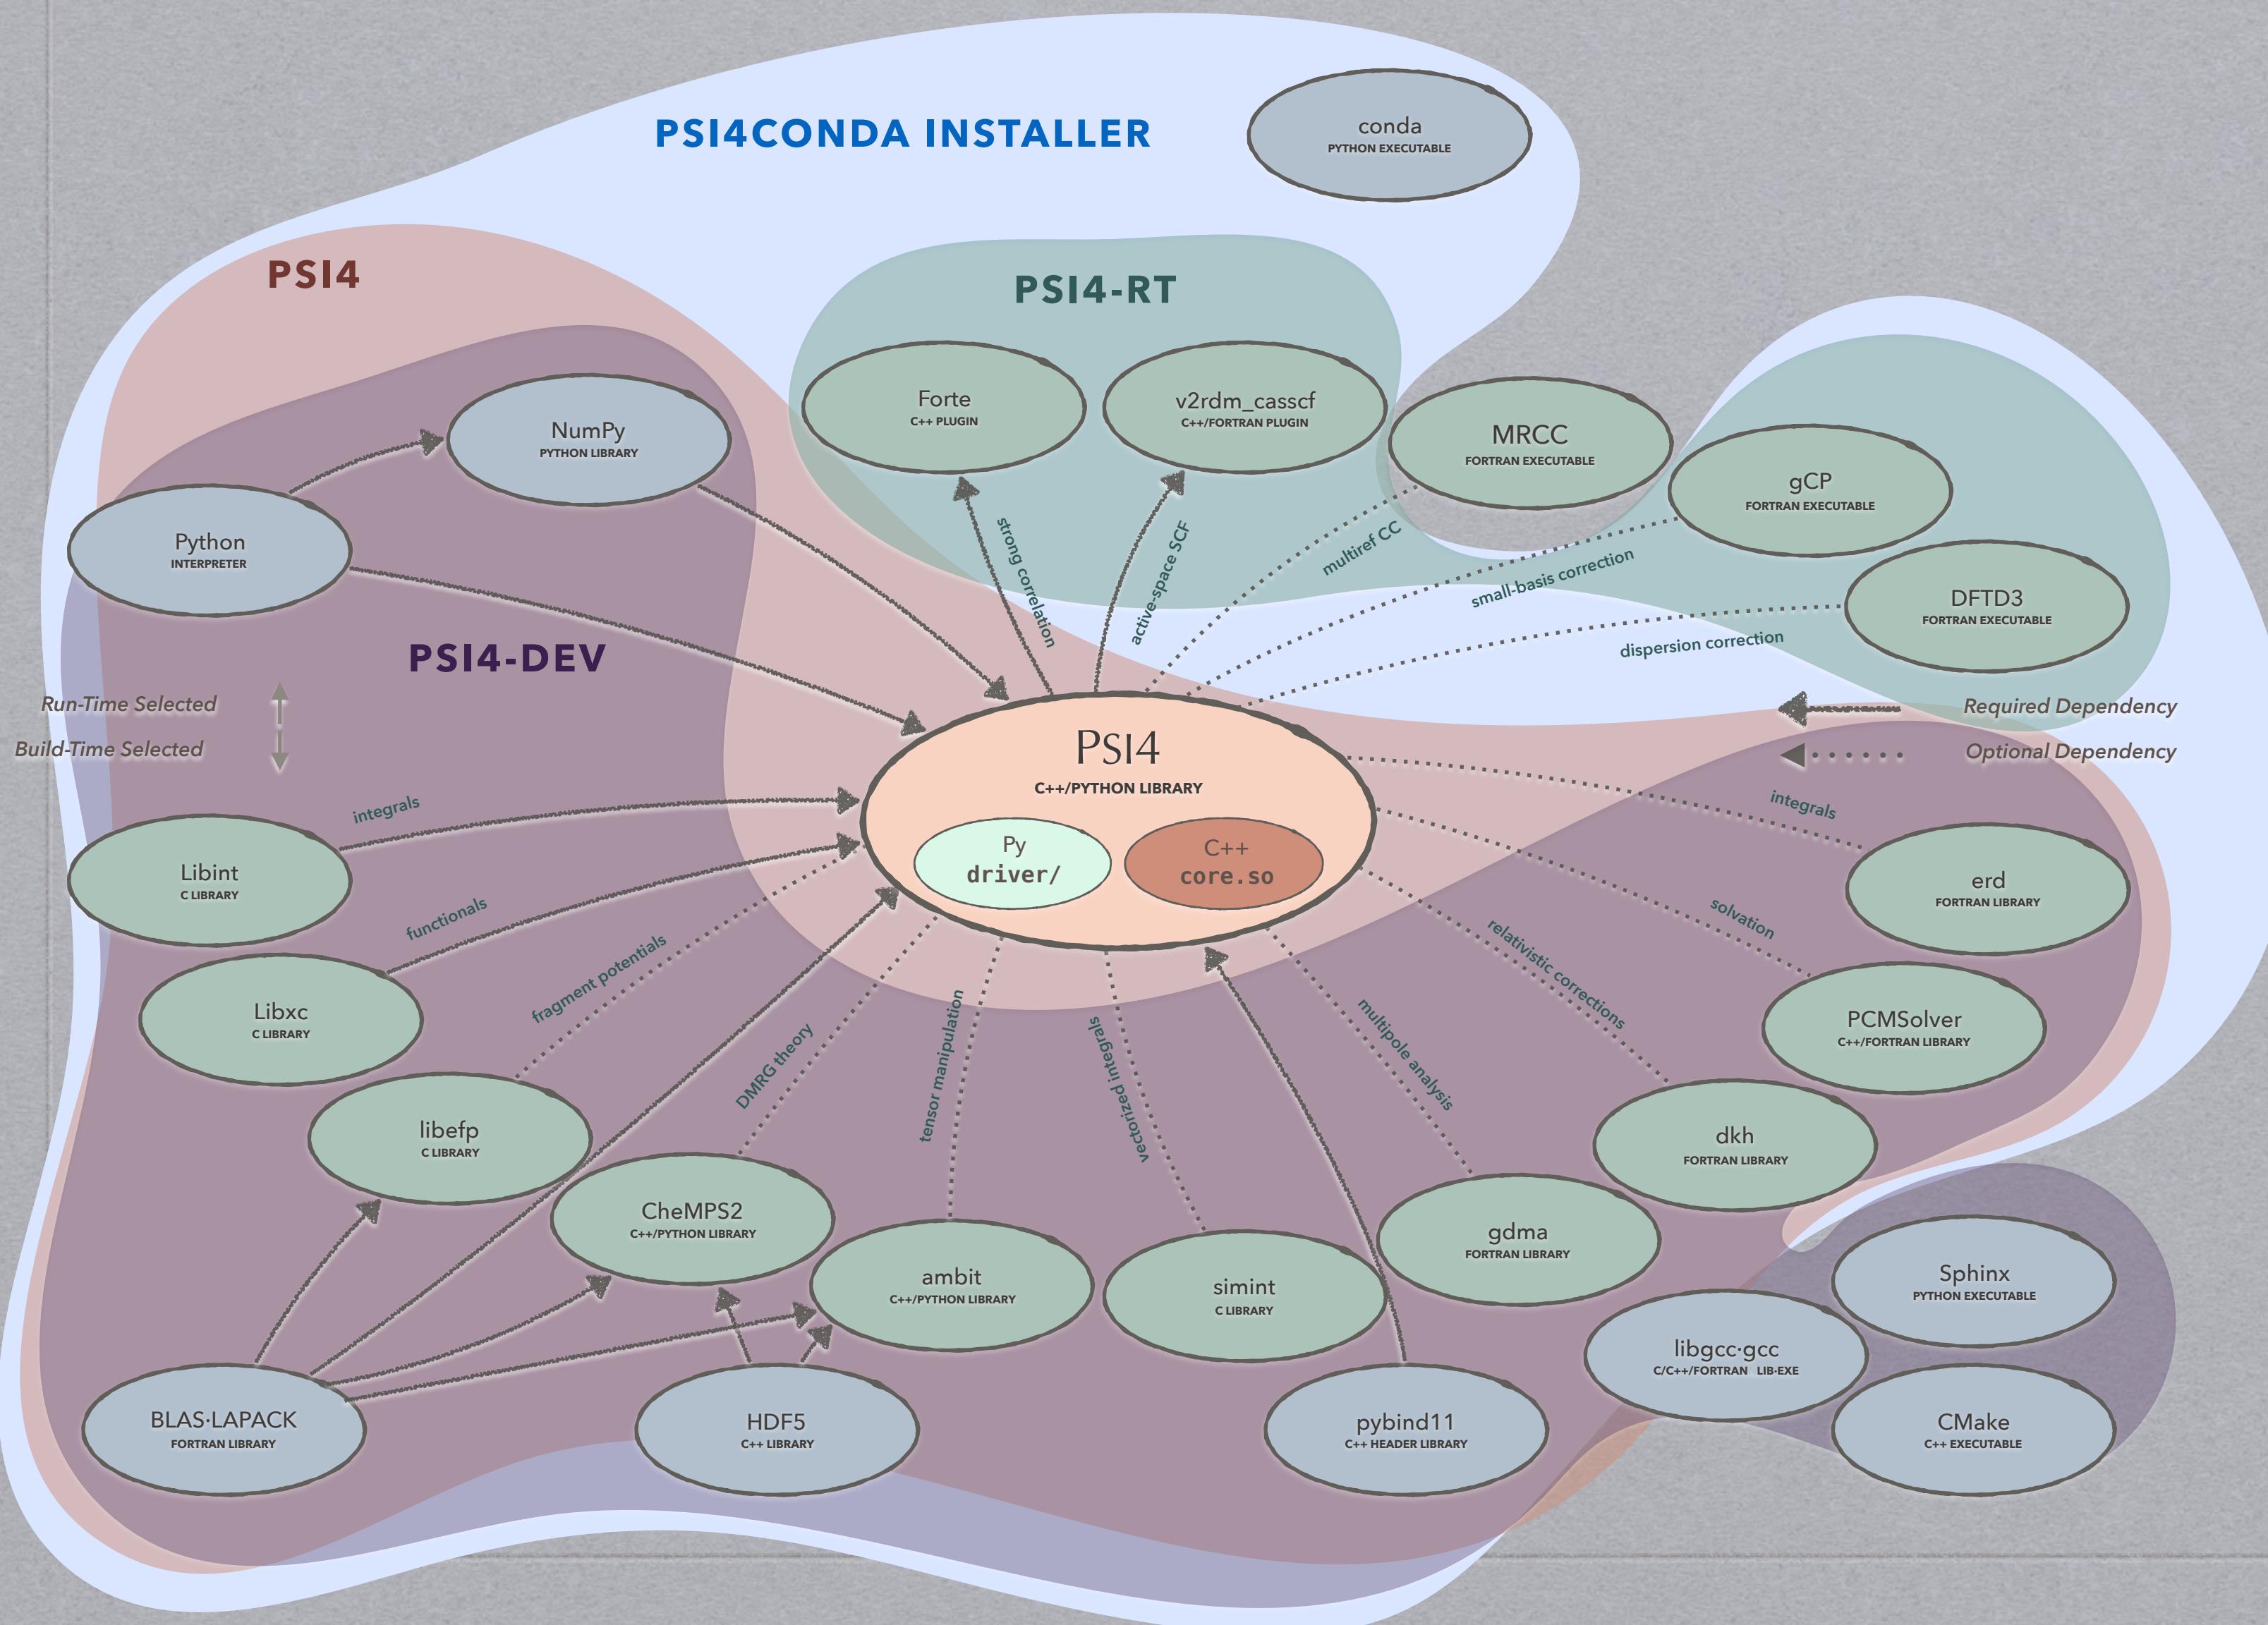
CMAKE & CONDA FOR ALL

TRICKY OPTIMIZING FOR ALL USERS AT ONCE



CMAKE & CONDA FOR ALL

TRICKY OPTIMIZING FOR ALL USERS AT ONCE



- Questions for you wrt psi4-dev:
 - I've copied over a headers pkg so you can get "mkl.h" via "-c psi4", but I'd rather not do too much of that.
 - Can you handle "conda create -n p4dev -c intel -c psi4/label/dev -c psi4" so you get mkl and numpy from Intel?
 - How much docs stuff in psi4-dev? Psi4 will tell you the long path with several packages & channels to get the docs deps. Do you really want them in psi4-dev? If so, can you handle "-c conda-forge -c astropy"?
 - Any problems with psi4-path-advisor? Remember its "-h"

WHAT, PRAY, SHOULD I CONDA INSTALL?

1

```
conda create -n p4env
```

What do you want to do?

run PSI4

build PSI4 plugin

build PSI4

run PSI4 in Travis CI

+`= psi4 psi4-rt`

+`= psi4 psi4-dev`

+`= psi4-dev`

+`= ci-psi4 psi4`

2.7

3.5

3.6

+`= python=2.7`

+`= python=3.5`

+`= python=3.6`

3

What computing platform are you targeting?

5

Do you want the last stable/tagged release or a development/nightly build?

latest release

nightly build

+`= -c psi4`

+`= -c psi4/label/dev -c psi4`

6

Do you want to build the docs?

yes (secondary cmd below)

```
conda install sphinx-psi-theme cloud_sptheme astropy-helpers graphviz python-graphviz -c conda-forge -c astropy -c psi4
```

4dLW

What does the following return?
`ldd --version`

< 2.7 (e.g., 2.6, 2.2.4, 1.09)

≥ 2.7 (e.g., 2.7, 2.14, 2.26)

Sorry, c. 2007 fundamental library. You will not get a PSI4 conda package today.

Good, PSI4 conda packages will work for you.

4cLW

Execute the following:
`arch`

x86_64

Anything else

Good, PSI4 conda packages will work for you.

Sorry, 32-bit. You will not get a PSI4 conda package today.

4aM

Mac
Execute the following:
`sysctl machdep.cpu | grep features | grep -o -e SSE4.1 -e SSE4.2 -e AVX1.0 -e AVX2`

(no output)

SSE4.1

SSE4.2 (among others)

AVX2 (among others)

Sorry, ancient hardware. You will not get a PSI4 conda package today.

You'll have to use the special sse41 version, but PSI4 conda packages will work for you.

You CAN'T use Intel Distribution for Python conda packages.

+`= sse41`

+`= sse41 -c intel`

+`= -c intel`

Windows

Do you have Bash on Ubuntu on Windows (<https://msdn.microsoft.com/en-us/commandline/wsl/about>)?

4aW

4bLW

No output

Output includes:
`sse2`

Output includes:
`sse4_2`

Output includes:
`avx` or `avx2`

Sorry, you will not get a PSI4 conda package today.

Good, PSI4 conda packages will work for you.

Good, PSI4 conda packages will work for you. Also, you can use Intel Distribution for Python conda packages.

Good, PSI4 conda packages are optimized for your instruction set. Also, you can use Intel Distribution for Python conda packages.

+`= -c intel`

+`= -c intel`

Do you get the following:
`OMP: Error #100: Fatal system error detected.`

Execute the following:
`export KMP_AFFINITY=disabled`

DOWN WITH GLOBALS

PSI4 BETA5

[US] | <https://github.com/psi4/psi4archive/blob/4.0b5/include/psi4-dec.h>

Psi4 Docs   PDFdb  

```
public:
    class Environment
    {
        std::map<std::string, std::string> environment_;
        unsigned long int memory_;
        int nthread_;

        boost::shared_ptr<Molecule> molecule_;
        SharedMatrix gradient_;
        boost::shared_ptr<Vector> frequencies_;
        boost::shared_ptr<Wavefunction> wavefunction_;
        boost::shared_ptr<PointGroup> parent_symmetry_;
```

PSI4 1.0

[US] | <https://github.com/psi4/psi4archive/blob/1.0.x/include/process.h>

Psi4 Docs   PDFdb  

```
public:
    class Environment
    {
        std::map<std::string, std::string> environment_;
        unsigned long int memory_;
        int nthread_;

        boost::shared_ptr<Molecule> molecule_;
        SharedMatrix gradient_;
        boost::shared_ptr<efp::EFP> efp_;
        SharedMatrix efp_torque_;
        boost::shared_ptr<Vector> frequencies_;
        boost::shared_ptr<PointGroup> parent_symmetry_;

        boost::shared_ptr<Molecule> legacy_molecule_;
        boost::shared_ptr<Wavefunction> legacy_wavefunction_;
```

NOW

[] | <https://github.com/psi4/psi4/blob/master/psi4/src/psi4/libpsi4util/process.h>

Docs   PDFdb  

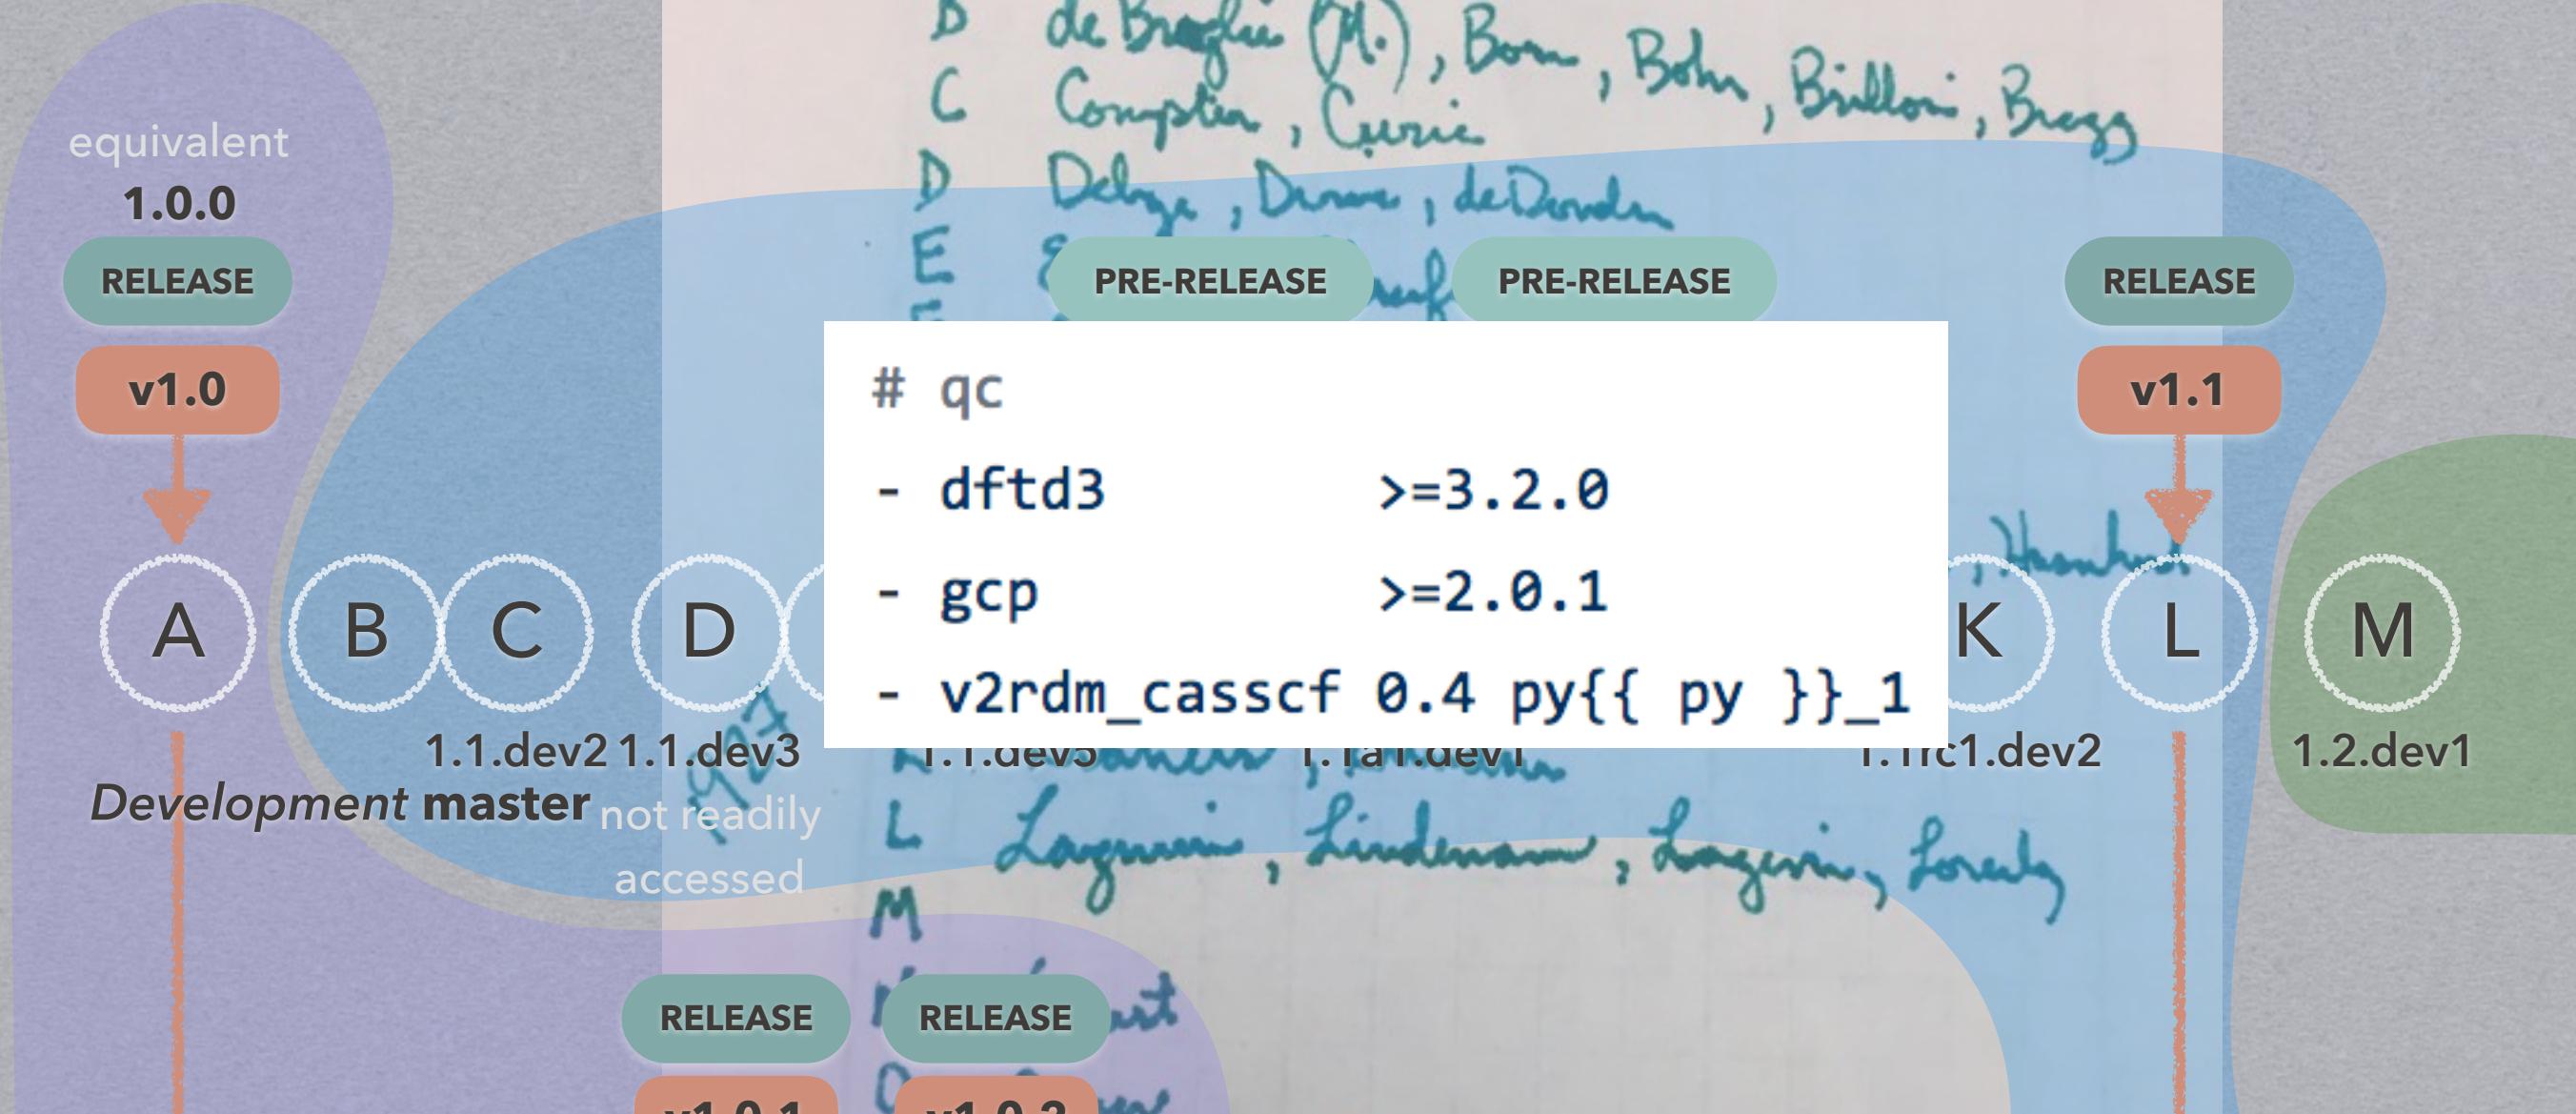
```
class Environment {
    std::map<std::string, std::string> environment_;
    size_t memory_;
    int nthread_;
    std::string datadir_;

    std::shared_ptr<Molecule> molecule_;
    SharedMatrix gradient_;
    std::shared_ptr<efp::EFP> efp_;
    SharedMatrix efp_torque_;
    std::shared_ptr<Vector> frequencies_;
    std::shared_ptr<PointGroup> parent_symmetry_;

    std::shared_ptr<Molecule> legacy_molecule_;
    std::shared_ptr<Wavefunction> legacy_wavefunction_;
```

VERSIONING

PRAY I DON'T ALTER IT ANY FURTHER



- Last year's scheme is working great! ...
 - needs a checklist and a calm mind to bump versions, but that's ok
 - ... for computers
 - ... for advertisers (1.1 hoopla)
 - ... not so much for people or downstream devs
 - have to keep rebuilding for ABI changes
 - have to remember how many commits past tag their feature of interest is

```
using_psi4_libxc = pytest.mark.skipif(is_psi4_new_enough("1.2a1.dev100") is False,  
                                       reason="Psi4 does not include DFT rewrite to use Libxc. Update to development head")
```

```
using_psi4_efpmints = pytest.mark.skipif(is_psi4_new_enough("1.2a1.dev507") is False,  
                                         reason="Psi4 does not include EFP integrals in mints. Update to development head")
```

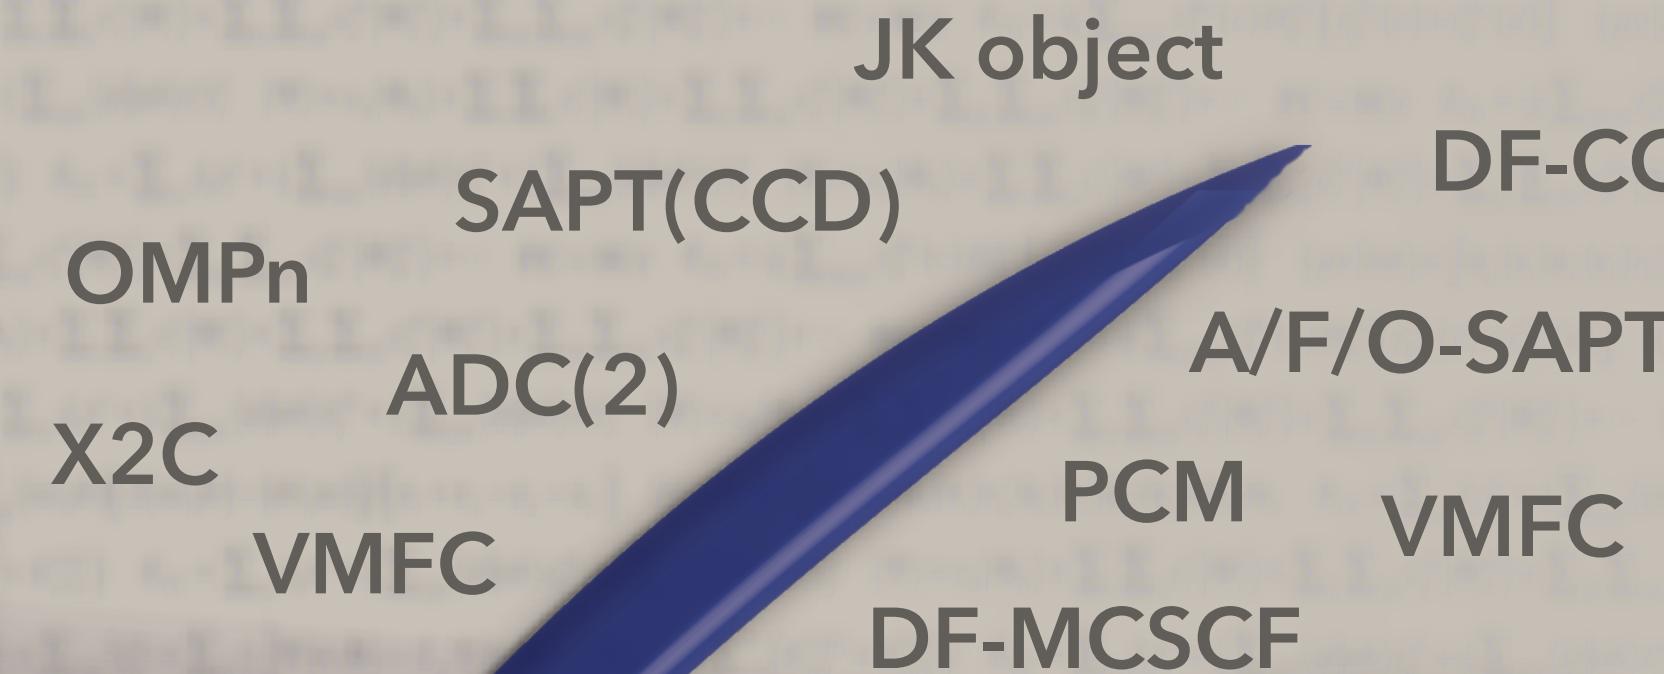
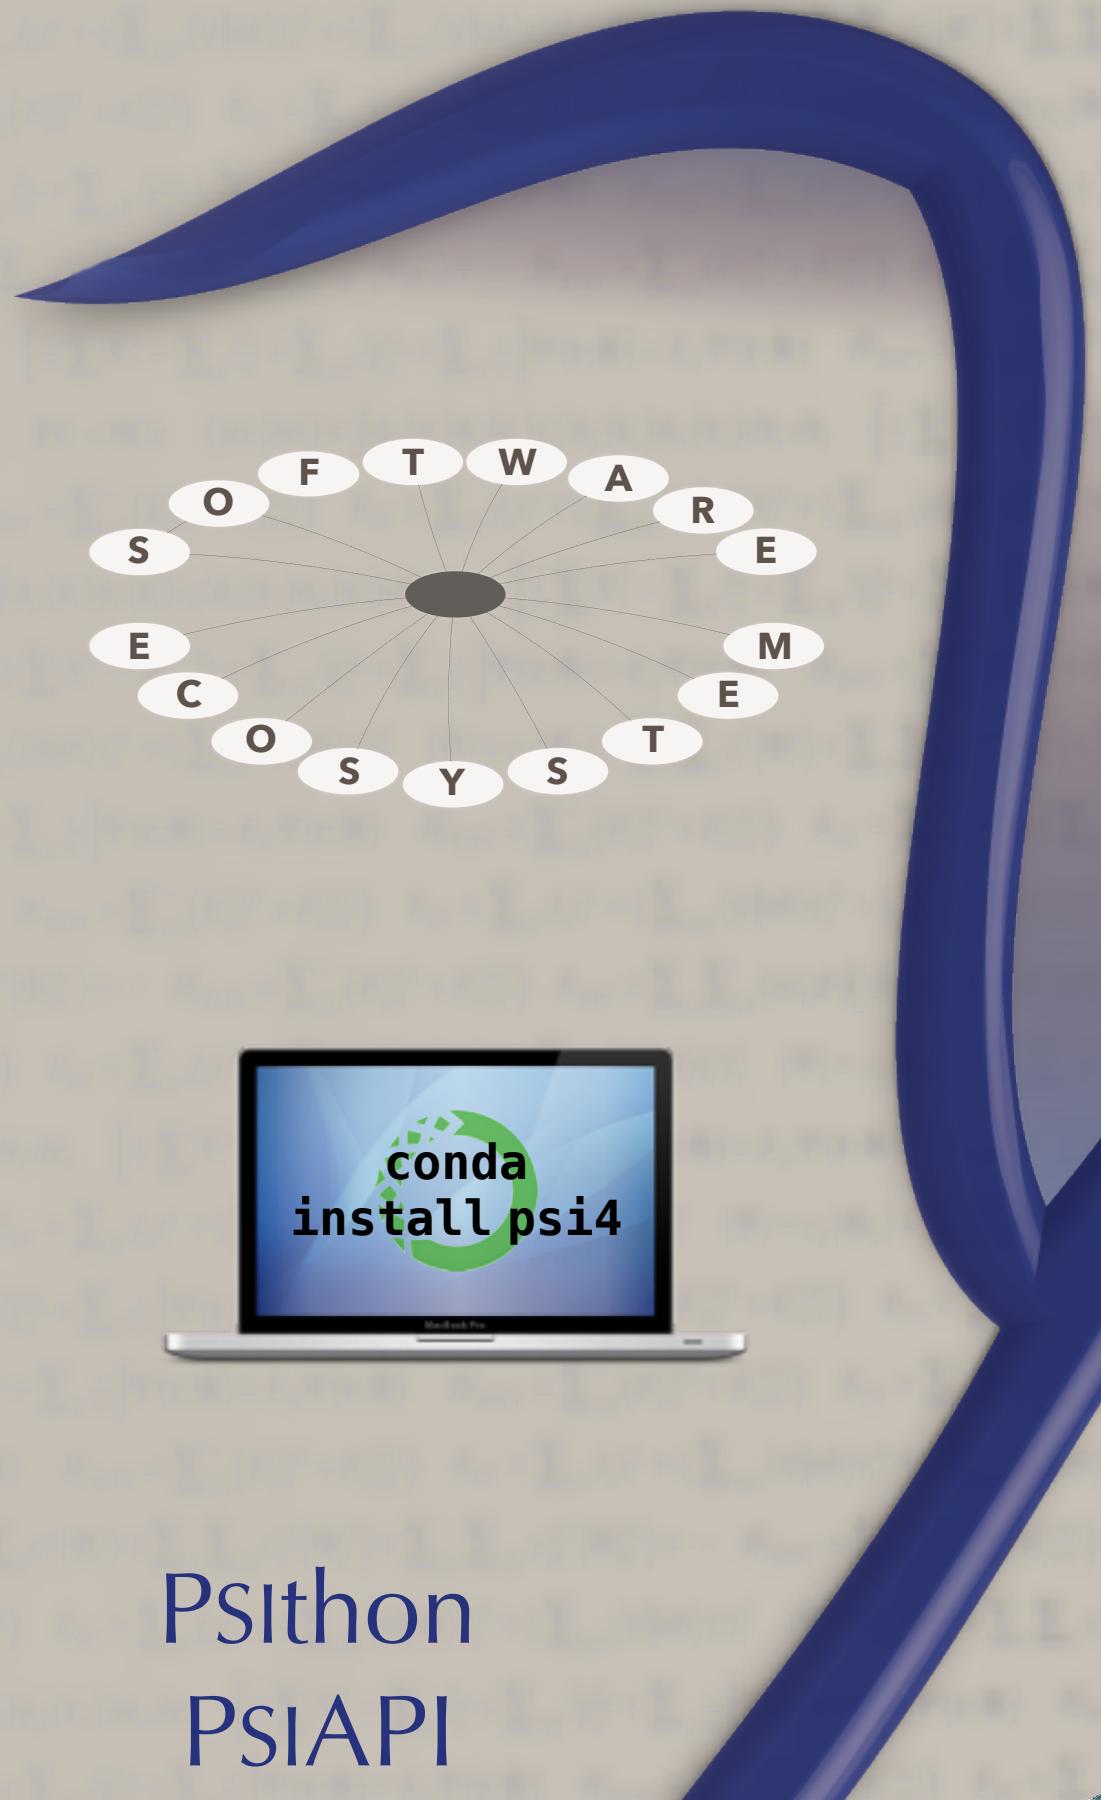
```
using_psi4_sadpy = pytest.mark.skipif(is_psi4_new_enough("1.2a1.dev600") is False,  
                                       reason="Psi4 does not include SADGuess bindings. Update to development head")
```

vertising
from API-
Psi4 1.13)

SCIPY 2018

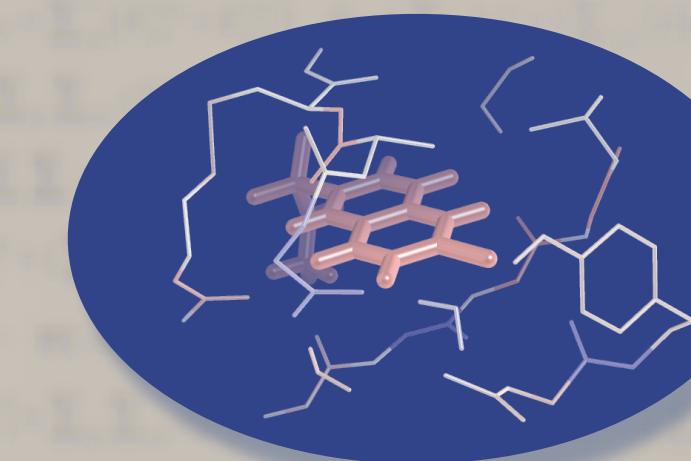
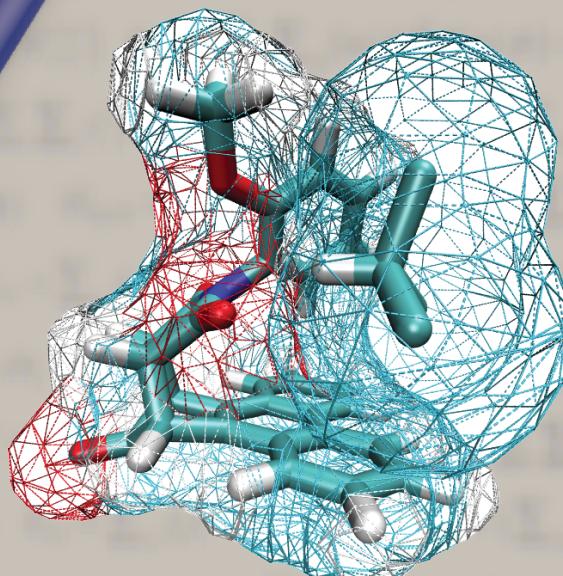
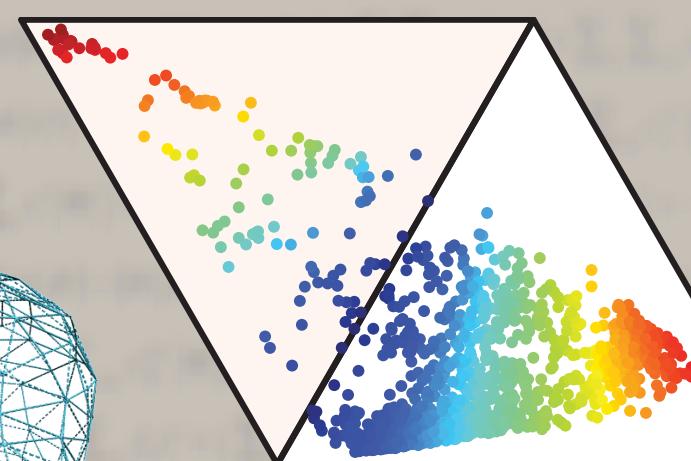
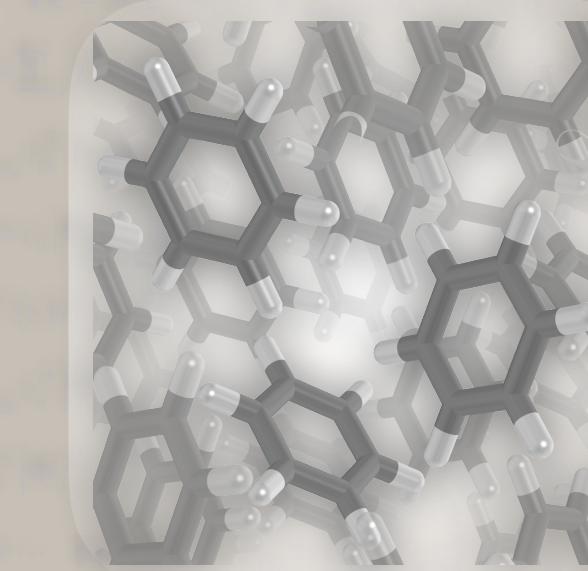
- Austin, TX
- July 11-13
- MolSSI & others organizing a Chemistry & Materials mini symposium
- talk slots allocated proportionately to abstracts submitted, so think about what python aspects of your science are useful to others
- meet developers of your favorite software tools

PSI4 OPEN-SOURCE QUANTUM CHEMISTRY



JK object
OMPn
X2C
VMFC
SAPT(CCD)
ADC(2)
DF-CC
A/F/O-SAPT
PCM
VMFC
DF-MCSCF

PSI4NUMPY



PSIthon
PSIAPI
PSIfound

PSI4EDUCATION